

Developing an automated iterative near-term forecasting system for an ecological study

*Ethan P. White^{1,2,3}, Glenda M. Yenni¹, Shawn D. Taylor⁴, Erica M.
Christensen¹, Ellen K. Bledsoe⁴, Juniper L. Simonis¹, S. K. Morgan
Ernest^{1,3}*

¹ Department of Wildlife Ecology and Conservation, University of Florida, Gainesville,
FL, United States

² Informatics Institute, University of Florida, Gainesville, FL, United States

³ Biodiversity Institute, University of Florida, Gainesville, FL, United States

⁴ School of Natural Resources and Environment, University of Florida Gainesville, FL,
United States

Abstract

1. Most forecasts for the future state of ecological systems are conducted once and never updated or assessed. As a result, many available ecological forecasts are not based on the most up-to-date data, and the scientific progress of ecological forecasting models is slowed by a lack of feedback on how well the forecasts perform.
2. Iterative near-term ecological forecasting involves repeated daily to annual scale

forecasts of an ecological system as new data becomes available and regular assessment of the resulting forecasts. We demonstrate how automated iterative near-term forecasting systems for ecology can be constructed by building one to conduct monthly forecasts of rodent abundances at the Portal Project, a long-term study with over 40 years of monthly data. This system automates most aspects of the six stages of converting raw data into new forecasts: data collection, data sharing, data manipulation, modeling and forecasting, archiving, and presentation of the forecasts.

3. The forecasting system uses R code for working with data, fitting models, making forecasts, and archiving and presenting these forecasts. The resulting pipeline is automated using continuous integration (a software development tool) to run the entire pipeline once a week. The cyberinfrastructure is designed for long-term maintainability and to allow the easy addition of new models. Constructing this forecasting system required a team with expertise ranging from field site experience to software development.
4. Automated near-term iterative forecasting systems will allow the science of ecological forecasting to advance more rapidly and provide the most up-to-date forecasts possible for conservation and management. These forecasting systems will also accelerate basic science by allowing new models of natural systems to be quickly implemented and compared to existing models. Using existing technology, and teams with diverse skill sets, it is possible for ecologists to build these systems and use them to advance our understanding of natural systems.

Key-words: forecasting, prediction, mammals, iterative forecasting, Portal Project

43 **Introduction**

44 Forecasting the future state of ecological systems is important for management,
45 conservation, and evaluation of how well models capture the processes governing
46 ecological systems (Clark et al., 2001; Tallis & Kareiva, 2006; Díaz et al., 2015; Dietze,
47 2017). In 2001, Clark et al. (2001) called for a more central role of forecasting in
48 ecology. Since then, an increasing number of ecological forecasts are being published
49 that focus on a societally important questions from daily to decadal time scales (Dietze
50 et al., 2018). At daily scales, ecological forecasts predict the occurrence of
51 environmental issues like toxic algal blooms (Stumpf et al., 2009). At quarterly scales
52 forecasts are used to predict the stocks of fisheries and make decisions about quotas for
53 fishing (NOAA, 2016). At decadal time scales, ecological forecasts are used to predict
54 how biodiversity will change as it responds to anthropogenic influences (Harris et al.,
55 2018). These forecasting examples highlight the important role that ecological forecasts
56 play in recasting ecological knowledge in societally relevant ways and also improve our
57 understanding of ecological systems by testing the ability of our models to predict how
58 systems will change in the future (Dietze et al., 2018, Harris et al. (2018)).

59 While some of the examples given above (e.g., fisheries stock estimates) are regularly
60 repeated, most ecological forecasts are made once, published, and never assessed or
61 updated (Dietze et al., 2018). This lack of both regular assessment and active updating
62 has limited the progress of ecological forecasting and hindered our ability to make
63 useful and reliable predictions. The lack of active assessment results in limited
64 information on how much confidence to place in forecasts and makes it difficult to
65 determine on which forecasting methods to build. Without regular updates, forecasts
66 lack the most current data, and the longer a forecast remains out of date, the less
67 accurate it becomes (Petchey et al., 2015; Dietze et al., 2018). More regular updating
68 and assessment will advance ecological forecasting as a field by accelerating the
69 identification of the best models for individual forecasts and improving our

70 understanding of how to best design forecasting approaches for ecology in general. This
71 approach has helped accelerate forecasting ability in other fields including meteorology
72 (Kalnay, 2003, McGill (2012), Bauer et al. (2015)). For ecological forecasting to
73 mature as a field, we need to change how we produce and interact with forecasts,
74 creating a more dynamic interplay between model development, prediction generation,
75 and incorporation of new data and information (Dietze et al., 2018).

76 With the goal of making ecological forecasting more dynamic and responsive, Dietze et
77 al (2018) recently called for an increase in iterative near-term forecasting. Iterative
78 near-term forecasting is defined as making predictions for the near future and repeatedly
79 updating those predictions through a cycle of evaluation, integration of new data, and
80 generation of new forecasts. Because forecasts are made ‘near-term’—daily to annual
81 time scales instead of multi-decadal—predictions can be assessed more quickly and
82 frequently, leading to more rapid model improvements (Tredennick et al., 2016; Dietze
83 et al., 2018). Since forecasts are made repeatedly through time, new data can be
84 continuously integrated with each iteration (Dietze et al., 2018). By quickly identifying
85 how models are failing, facilitating rapid testing of improved models, and incorporating
86 the most up-to-date data available, iterative near-term forecasting has the potential to
87 promote rapid improvement in the state of ecological forecasting. In addition to
88 yielding improved information for guiding policy and management (Clark et al., 2001;
89 Luo et al., 2011; Petchey et al., 2015), this iterative approach will help improve our
90 basic understanding of ecological systems (Dietze et al., 2018). For example, alternative
91 mechanistic models can be compared to determine which model provides the best
92 forecasts, thus providing insights into the importance of different ecological processes
93 (Dietze et al., 2018). Iterative near-term forecasting provides the more dynamic
94 interplay between models, predictions, and data that has been identified as necessary for
95 improving ecological forecasting and our understanding of ecological systems more
96 broadly.

97 Because iterative near-term forecasting requires a dynamic integration of models,
98 predictions, and data, Dietze et al (2018) highlight approaches to data management,
99 model construction and evaluation, and cyberinfrastructure that are necessary to
100 effectively implement this type of forecasting (Box 1). Data needs to be released quickly
101 under open licenses (Vargas et al., 2017; Dietze et al., 2018) and structured so that it can
102 be used easily by a variety of researchers and in multiple modeling approaches (Borer et
103 al., 2009; Strasser et al., 2011). Models need to be able to deal with uncertainty, in both
104 the predictors and the predictions, to properly convey uncertainty in the resulting
105 forecasts (Diniz-Filho et al., 2009). Multiple models should be developed, both to assess
106 which models are performing best (Dietze et al., 2018) and to facilitate combining
107 models to form ensemble predictions which tend to perform better than single models
108 (Araujo & New, 2007; Diniz-Filho et al., 2009). Ensuring that data and models are
109 regularly updated and new forecasts are made requires cyberinfrastructure to automate
110 data processing, model fitting, prediction, model evaluation, forecast visualization, and
111 archiving. In combination, these approaches should allow forecasts to be easily rerun
112 and evaluated as new data becomes available (Box 1; Dietze et al., 2018).

113 While iterative near-term forecasting is an important next step in the evolution of
114 ecological forecasting, the requirements outlined by Dietze et al (Box 1) are not trivial
115 to implement (e.g., making quality data available in near real-time and automatically
116 rerunning forecasts in reproducible ways), and few of their recommendations are in
117 widespread use in ecology today (Stodden & Miguez, 2014; e.g., Wilson et al., 2014;
118 Yenni et al., 2018). We explored what it would entail to operationalize Dietze et al's
119 recommendations by constructing our own iterative near-term forecasting pipeline for
120 an on-going, long-term ecological study that collects high-frequency data on desert
121 rodent abundances (J. H. Brown, 1998; S. K. M. Ernest et al., 2008). We constructed an
122 automated forecasting pipeline with the goal of being able to forecast rodent
123 abundances and evaluate our predictions on a monthly basis. In this paper, we discuss

124 our approach for creating this iterative near-term forecasting pipeline, the challenges we
125 encountered, the tools we used, and the lessons we learned so that others can create
126 their own iterative forecasting systems. For those interested in implementing iterative
127 forecasting, either on their own or as part of a team, this paper will provide a roadmap
128 for how to build such a system and what skills will be helpful to do so. For readers
129 looking for an introduction to automation and continuous integration in an ecological
130 context, we recommend our paper on data management for continuously collected data,
131 which includes a tutorial on how to set up some of the aspects of automation described
132 in this paper (Yenni et al. (2018)).

133 **System Background**

134 Iterative forecasting is most effective with frequently collected data, since it provides
135 more opportunities for updating model results and assessing (and potentially improving)
136 model performance (Box 1; Dietze et al., 2018). The Portal Project is a long-term
137 ecological study situated in the Chihuahuan Desert (2 km north and 6.5 km east of
138 Portal, Arizona, US). Researchers have been continuously collecting data at the site
139 since 1977, including data on the abundance of rodent and plant species (monthly and
140 twice yearly, respectively) and climatic factors such as air temperature and precipitation
141 (daily) (J. H. Brown, 1998; S. K. M. Ernest et al., 2009, 2016). The site consists of 24
142 50m x 50m experimental plots. Each plot contains 49 permanently marked trapping
143 stations laid out in a 7 x 7 grid, and all plots are trapped with Sherman live traps for one
144 night each month. For all rodents caught during a trapping session, information on
145 species identity, size, and reproductive condition is collected, and new individuals are
146 given identification tags. This information on rodent populations is high-frequency, uses
147 consistent trapping methodology, and has an extended time-series (470 monthly samples
148 and counting), making this study an ideal case for near-term iterative forecasting.

149 **Implementing an automated iterative forecasting system**

150 Implementation of iterative forecasting requires the regular rebuilding of models with
151 new raw data as it becomes available and the presentation of those forecasts in usable
152 forms; in our case, this occurs monthly. Rebuilding models in an efficient and
153 maintainable way relies on developing an automated pipeline to handle the six stages of
154 converting raw data into new forecasts: data collection, data sharing, data manipulation,
155 modeling and forecasting, archiving, and presentation of the forecasts (Figure 1a). To
156 implement the pipeline outlined in Figure 1a, we used a “continuous analysis”
157 framework (*sensu* Beaulieu-Jones & Greene, 2017) that automatically processes the
158 most up-to-date data, refits the models, makes new forecasts, archives the forecasts, and
159 updates a website with analysis of current and previous forecasts. In this section we
160 describe our approach to streamlining and automating the multiple components of the
161 forecasting pipeline and the tools and infrastructure we employed to execute each
162 component.

163 **Continuous Analysis Framework**

164 A core aspect of iterative near-term forecasting is the regular rerunning of the
165 forecasting pipeline. We employed “continuous analysis” (*sensu* Beaulieu-Jones &
166 Greene, 2017) to drive the automation of both the full pipeline and a number of its
167 individual components. Continuous analysis uses a set of tools originally designed for
168 software development called “continuous integration” (CI). CI combines computing
169 environments for running code with monitoring systems to identify changes in data or
170 code. Essentially, CI is a computer helper who watches the pipeline and, when it sees a
171 change in the code or data, runs all the computer scripts needed to ensure that the
172 forecasting pipeline runs from beginning to end. This is useful for iterative near-term
173 forecasting because it does not rely on humans to create new forecasts whenever new

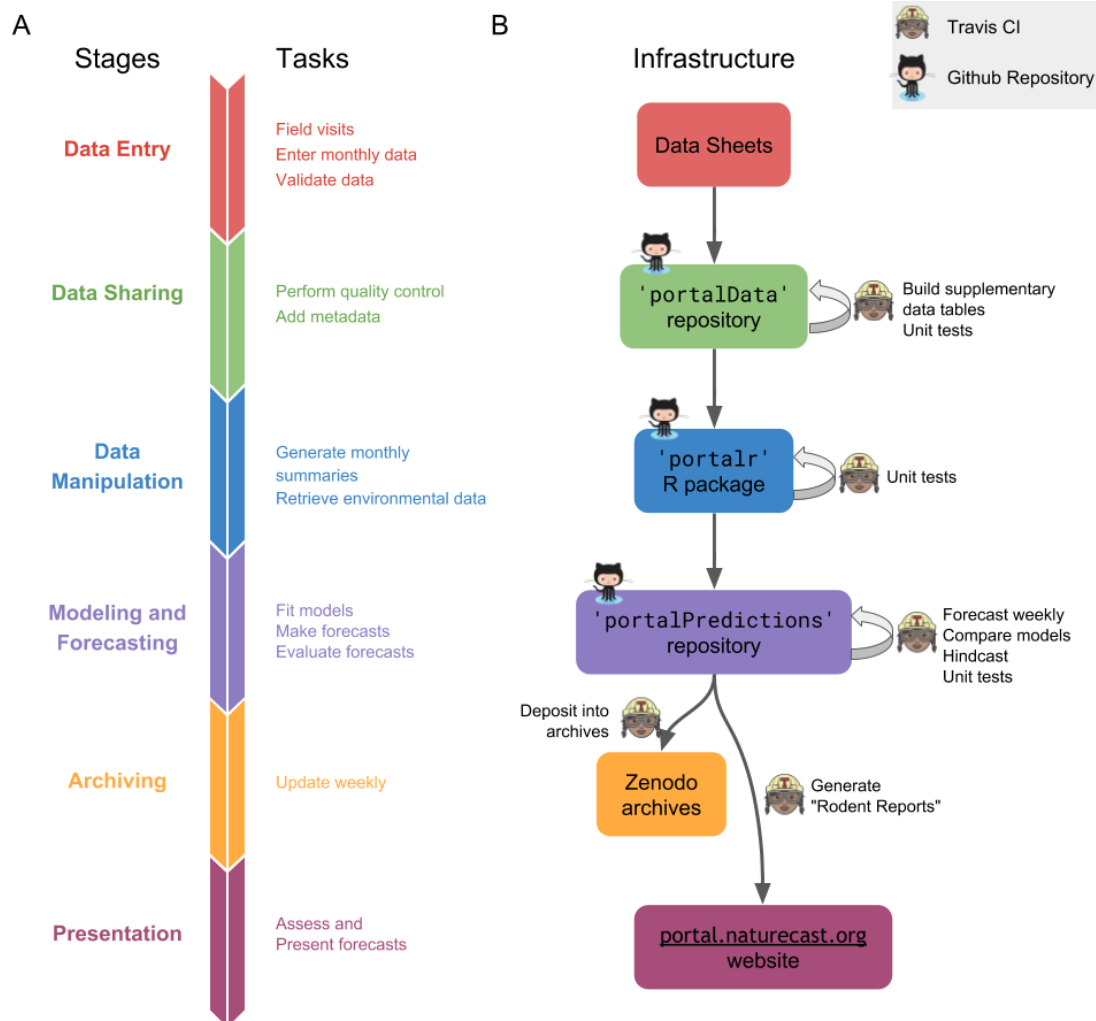


Figure 1: a) Stages of the forecasting pipeline. To go from raw data to forecast presentation involves a number of stages, each of which requires unique tasks, tools and infrastructure. The stages are interdependent, with outputs from one stage forming the inputs for the subsequent stage. Tasks in all stages are run using code written in R. b) Continuous integration system. Each box denotes the core infrastructure used for each stage of the forecasting pipeline. Continuous integration (denoted by the Travis icon, a woman wearing safety glasses and hardhat) triggers the code involved in events that link the stages of the pipeline, such as using the output from the forecasting stage (purple box) to create an updated website (rose box). Travis also runs tasks within a stage, such as testing code and adding weather data (icons on arrows originating and ending on the same box). The code for driving different stages of this pipeline is stored on GitHub (denoted by the GitHub icon, an “octocat”).

174 models or data are added. These tools are common in the area of software development,
175 where they are used to automate software testing and integrate work by multiple
176 developers working on the same code base. However, these tools can be used for any
177 computational task that needs to be regularly repeated or run after changes to code or
178 data (Beaulieu-Jones & Greene, 2017). Our forecasting pipeline currently runs on a
179 publicly available continuous integration service (Travis CI; <https://travis-ci.org/>) that is
180 free for open source projects (up to a limited amount of computing time). This
181 continuous integration integrates directly with GitHub (<https://github.com>), the online
182 repository where we store the associated code and data. Because of the widespread use
183 of CI in software development, alternative services that can run code on local or
184 cloud-based computational infrastructure also exist (Beaulieu-Jones & Greene, 2017).
185 We use CI to quality check data, test code using “unit tests” (Wilson et al., 2014), build
186 models, make forecasts, and publicly present and archive the results (Figure 1b).

187 In addition to automatically running software pipelines, the other key component of
188 “continuous analysis” is making sure that the pipelines will continue to run even as
189 software dependencies change (Beaulieu-Jones & Greene, 2017). Many of us have
190 experienced the frustrations that can occur when software updates (e.g., changes in R
191 package versions) create errors in previously functional code. We experienced this issue
192 when the `tscount` package (Liboschik et al., 2015), used by one of our forecasting
193 models, was temporarily removed from CRAN (the R package repository) and could
194 not be installed in the usual way. This broke our forecasting pipeline, as we could no
195 longer run models that used that package. To make our pipeline robust to changes in
196 external software dependencies, we follow Beaulieu and Greene’s (2017)
197 recommendation to use software containers. Software containers are standalone
198 packages that contain copies of everything needed to run a given piece of software,
199 including the operating system. Once created, a software container is basically a time
200 capsule, containing all the software dependencies in the exact state used to develop and

run the software. If those dependencies change (or disappear) in the wider world, they still exist, unchanged, in the container. We use an existing platform, Docker (Merkel, 2014), to store an exact image of the complete software environment for running the forecasts. Docker also allows a specified set of packages to be used consistently across different computer and server environments. Using containers allows us to control transitions to new package versions, implementing them only after we have tested them and made any necessary changes to the data processing and analysis code. We use a container created by the Rocker project, which is a Docker image with many important R packages (i.e. tidyverse) pre-installed (Boettiger & Eddelbuettel, 2017). We add our code and dependencies to this existing Rocker image to create a software container for our forecasting pipeline. In combination, the automated running of the pipeline (continuous integration) and the guarantee it will not stop working unexpectedly due to software dependencies (via a software container) allows continuous analysis to serve as the glue that connects all stages of the forecasting pipeline.

Data Collection, Entry, and Processing

Iterative forecasting benefits from frequently updated data so that state changes can be quickly incorporated into new forecasts (Dietze et al., 2018). Both frequent data collection and rapid processing are important for providing timely forecasts. Since we collect data monthly, ensuring that the models have access to the newest data requires a data latency period of less than 1 month from collection to availability for modeling. To accomplish this, we automated components of the data processing and quality assurance/quality control (QA/QC) process to reduce the time needed to add new data to the database (Figure 1).

New data are double-entered into Microsoft Excel using the “data validation” feature. The two versions are then compared using an R script to control for errors in data entry. Quality control (QC) checks using the `testthat` R package (Wickham, 2011) are run

227 on the data to test for validity and consistency both within the new data and between the
228 new and archived data. The local use of the QC scripts to flag problematic data greatly
229 reduces the time spent error-checking and ensures that the quality of data is consistent.
230 The cleaned data are then uploaded to the GitHub-based PortalData repository
231 (<https://github.com/weecology/PortalData>). GitHub (<https://github.com/>) is a software
232 development tool for managing computer code development, but we have also found it
233 useful for data management. On GitHub, changes to data can be tracked through the Git
234 version control system which logs all changes made to any files in the repository, giving
235 us a record of exactly of when specific lines of data were changed or added. All updates
236 to data are processed through “pull requests,” which are notifications that someone has a
237 modified version of the data to contribute. QA/QC checks are automatically run on the
238 submitted data using continuous integration to ensure that no avoidable errors reach the
239 official version of the dataset.

240 We also automated the updating of supplementary data tables, including information on
241 weather and trapping history, that were previously updated manually. As soon as new
242 field data is merged into the repository, continuous integration updates all
243 supplementary files. Weather data is automatically fetched from our cellular-connected
244 weather station, cleaned, and appended to the weather data table. Supplementary data
245 tables related to trapping history are updated based on the data added to the main data
246 tables. Using CI for this ensures that all supplementary data tables are always
247 up-to-date with the core data.

248 **Data Sharing**

249 The Portal Project has a long history of making its data publicly available so that anyone
250 can use it for forecasting or other projects. Historically, the publication of the data was
251 conducted through data papers (S. K. M. Ernest et al., 2009, S. K. M. Ernest et al.
252 (2016)), the most common approach in ecology; this approach, however, caused years

253 of data latency. With the recent switch to posting data directly to a public GitHub
254 repository (Figure 1) with a CC0 waiver (i.e. no restrictions on data use;
255 <https://creativecommons.org/publicdomain/zero/1.0/>), data latency for everyone has
256 been reduced to less than one month, making meaningful iterative near-term forecasting
257 possible for not only our group but other interested parties, as well.

258 **Data Manipulation**

259 Once data is available, it must be processed into a form appropriate for modeling
260 (Figure 1). For many ecological datasets, this requires not only simple data
261 manipulation but also a good understanding of the data to facilitate appropriate
262 aggregation. Data manipulation steps are often conducted using custom one-off code to
263 convert the raw data into the desired form (Morris & White, 2013), but this approach
264 has several limitations. First, each researcher must develop and maintain their own data
265 manipulation code, which is inefficient and can result in different researchers producing
266 different versions of the data for the same task. Subtle differences in data processing
267 decisions have led to confusion when reproducing results for the Portal data in the past.
268 Second, this kind of code is rarely robust to changes in data structure and location.
269 Based on our experience developing and maintaining the Data Retriever (Morris &
270 White, 2013; Senyondo et al., 2017), these kinds of changes are common. Finally, this
271 kind of code is generally poorly tested, which can lead to errors based on mistakes in
272 data manipulation. To avoid these issues for the Portal Project data, the Portal team has
273 been developing an R package (portalr; <http://github.com/weecology/portalr>) for
274 acquiring the data and handling common data cleaning and aggregation tasks. As a
275 result, our modeling and forecasting code only needs to install this package and run the
276 data manipulation and summary functions to get the appropriate data (Figure 1b). The
277 package undergoes thorough automated unit testing to ensure that data manipulations
278 are achieving the desired results. Having data manipulation code maintained in a

279 separate package that focuses on consistently providing properly summarized forms of
280 the most recent data has made maintaining the forecasting code itself much more
281 straightforward.

282 **Modeling and Forecasting**

283 Iterative near-term forecasting involves regularly refitting a variety of different models
284 (Figure 1). Ideally, new models should be easy to incorporate to allow for iterative
285 improvements to the general modeling structure and approach. We use CI to refit the
286 models and make new forecasts each time the modeling code changes and when new
287 data become available (Figure 1b). We use a plugin infrastructure to allow new models
288 to be easily added to the system. This approach treats each model as an interchangeable
289 black box; all models have access to the same input data and generate the same structure
290 for model outputs (Figure 2). During each run of the forecasting code, all existing
291 models are run and the standardized outputs are combined into a single file to store the
292 results of the different models' forecasts. A weighted ensemble model is then added
293 with weights based on how well individual models fit the training data. This plugin
294 infrastructure makes it easy to add and compare very different types of models, from the
295 basic time-series approaches currently implemented to the more complex state-space
296 and machine learning models we hope to implement in the future. As long as a model
297 script can load the provided data and produce the appropriate output, it will be run and
298 its results incorporated into the rest of the forecasting system. This means that anyone
299 can add a new model to the existing system by: 1) creating their own copy of the project
300 (typically by forking the project on GitHub); 2) developing a new model; and 3)
301 submitting a pull request to our repository.

302 In addition to flexibility in what model structures can be supported, we also wanted to
303 support flexibility in what the models predict. Allowing models to make forecasts for
304 system properties ranging from individual species' population abundances to total

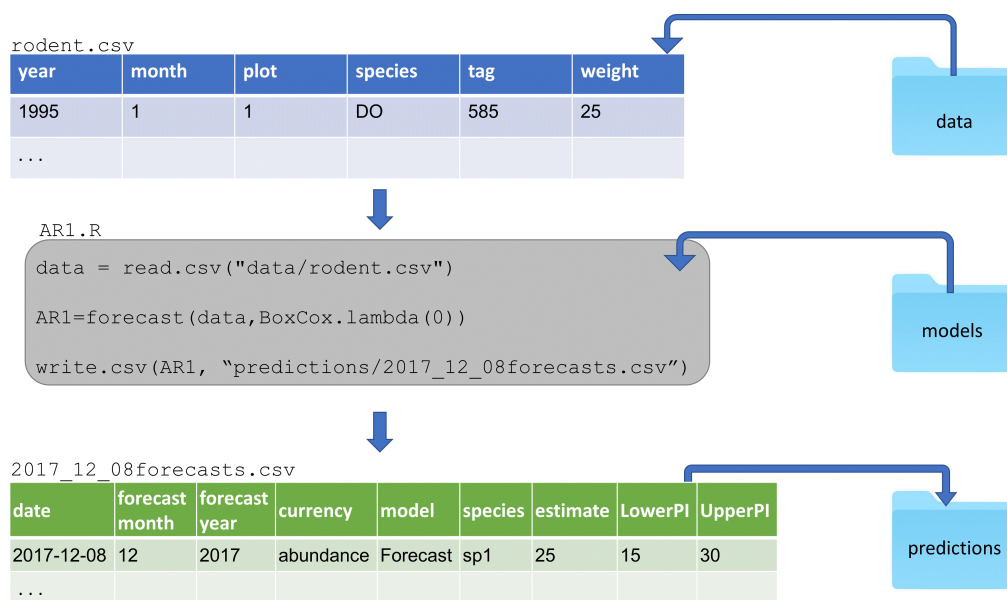


Figure 2: Demonstration of plugin infrastructure. All model scripts (represented here by the example AR1.R) are housed in a single folder. Each model script uses data provided by the core forecasting code (represented here by rodent.csv) and returns its forecast outputs in a predefined structure that is consistent across models (represented here by the example 2017_12_08forecasts.csv). Outputs from all models run on a particular date are combined into the same file (i.e. 2017_12_08forecasts.csv) to allow cross-model evaluations. Model output files are housed in a folder containing all forecast outputs from all previous dates to facilitate archiving and forecast assessment.

community biomass facilitates exploration of differences in forecastability across different aspects of ecological systems. We designed a forecast output format to support this. Each forecast output file contains the date being forecast, the collection date of the data used for fitting the models, the model name, the date the forecast was made, the state variable being forecast (e.g., rodent biomass, the abundance of a species), and the forecast value and associated uncertainty of that forecast (Figure 2). This allows us to store a variety of different forecasts in a common format and may serve as a useful starting point for developing a standard for storing ecological forecasts more generally. Forecasts are currently evaluated using root mean square error (RMSE) to evaluate point forecasts and coverage to evaluate uncertainty. We plan to add additional metrics in the future. In addition to evaluating the actual forecasts, we also use hindcasting (forecasting on already collected data; Jolliffe & Stephenson, 2003) to gain additional insight into the methods that work best for forecasting this system. For example, a model is fit using rodent observations up to June 2005, then used to make a forecast 12 months out to May 2006. The observations of that 12-month period can immediately be used to evaluate the model. Since hindcasting is conducted using data that has already been collected, it allows model comparisons to be conducted on large numbers of hindcasts and provides insight into which models make the best forecasts without needing to wait for new data to be collected (Harris et al., 2018). It can also be used to quickly evaluate new models instead of waiting for an adequate amount of data to accumulate. As the performance of different models is understood through evaluation of forecasts and hindcasts models can be refined or removed from the system or ensemble to iteratively improve the resulting forecasts.

Archiving

Publicly archiving forecasts before new data is collected allows the field to assess, compare, and build on forecasts made by different groups (McGill, 2012; Tredennick et

al., 2016; Dietze et al., 2018; Harris et al., 2018) (Figure 1). Archiving serves as a form of pre-registration for model predictions because the forecasts cannot be modified once the data to assess them has been collected. This helps facilitate an unbiased interpretation of model performance. To serve this role, archives should be publicly accessible and be a permanent record that cannot be changed or deleted. This second criterion means that GitHub is not sufficient for archival purposes because repositories can be changed or deleted (Bergman, 2012; White, 2015). We explored three major repositories for archiving forecasts: FigShare (<https://figshare.com/>), Zenodo (<https://zenodo.org/>), and Open Science Framework (<https://osf.io/>). While all three repositories allowed for easy manual submissions (i.e., a human uploading files after each forecast), automating this process was substantially more difficult. Various combinations of repositories, APIs (i.e., interfaces for automatically interacting with the archiving websites), and associated R packages had issues with: 1) integrating authorization with continuous integration; 2) automatically making archived files public; 3) adding new files to an existing location; or 4) automatically permanently archiving the files. Our eventual solution was to leverage the GitHub-Zenodo integration (<https://guides.github.com/activities/citable-code/>) and automatically push forecasts to a GitHub repository from the CI server and release them via the GitHub API. The GitHub-Zenodo integration is designed to automatically create versioned archives of GitHub repositories. We created a repository for storing forecasts (<https://github.com/weecology/forecasts>) and linked this repository with Zenodo (a one-time manual process). Each time a new forecast is created, our pipeline adds the new forecasts to the GitHub repository and uses the GitHub API to create a new “release” for that repository. This triggers the GitHub-Zenodo integration, which automatically archives the resulting forecasts under a top-level DOI that refers to all archived forecasts (<https://doi.org/10.5281/zenodo.839580>). Through this process, we automatically archive every forecast made with a documented time-stamp. In addition, we also archive the full state of the modeling and forecasting repository

359 (<https://doi.org/10.5281/zenodo.833438>). This ensures that every forecast is fully
360 reproducible since the exact code used to generate every forecast is preserved. Early
361 forecasts from this system are archived in the modeling and forecasting code archive,
362 not in the newer repository ‘forecasts’.

363 **Presentation**

364 Each month, we present our forecasts on a website that displays monthly rodent
365 forecasts, model evaluation metrics, monthly reports, and information about the study
366 site (Figure 3; <http://portal.naturecast.org>). The website includes a graphical
367 presentation of the most recent month’s forecasts (including uncertainty) and compares
368 the latest data to the previous forecasts. Information on the species and the field site are
369 also included. The site is built using Rmarkdown (Allaire et al., 2017), which naturally
370 integrates into the pipeline and is automatically updated after each forecast. The `knitr`
371 R package (Xie, 2015) compiles the code into HTML, which is then published using
372 Github Pages (<https://pages.github.com/>). The files for the website are stored in a
373 subdirectory of the forecasting repository. As a result, the website is also archived
374 automatically as part of archiving the forecast results.

375 **Discussion**

376 Following the recommendations of Dietze et al (2018), we developed an automated
377 iterative forecasting system (Figure 1) to support repeated forecasting of an ecological
378 system. Our forecasting system automatically acquires and processes the newest data,
379 refits the models, makes new forecasts, publicly archives those forecasts, and presents
380 both the current forecast and information on how previous forecasts performed. Every
381 week, the forecasting system generates a new set of forecasts with no human
382 intervention, except for the entry of new field data. Our forecasting system ensures that

Portal Forecast

Total Abundance Forecast

This is the forecast for next month's sampling of rodents at Portal.

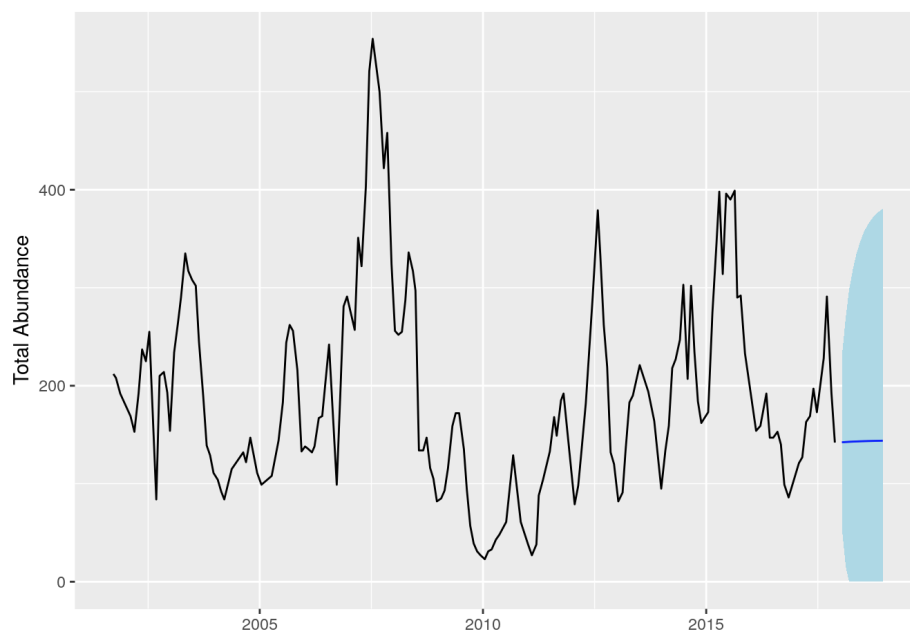


Figure 3: Screen capture of the homepage of the Portal Forecasting website (<http://portal.naturecast.org>). This site contains information on the most current forecasts, evaluation of forecast performance, and general information about the species being forecast.

383 forecasts based on the most recent data are always available and is designed to allow
384 rapid assessment of the performance of multiple forecasting models for a number of
385 different states of the system, including the abundances of individual species and
386 community-level variables such as total abundance. To create this iterative near-term
387 forecasting system, we used R to process data and conduct analyses and leveraged
388 existing tools and services (i.e. GitHub, Travis, Docker) for more complicated
389 cyberinfrastructure tasks. Thus, our approach to developing iterative near-term
390 forecasting infrastructure provides an example for how short-term ecological
391 forecasting systems can be developed.

392 We designed this forecasting system with the goal of making it relatively easy to build,
393 maintain, and extend. We used existing technology for both running the pipeline and
394 building individual components, which allowed us to build the system relatively cheaply
395 in terms of both time and money. This included the use of tools like Docker for
396 reproducibility, Travis CI continuous integration for automatically running the pipeline,
397 Rmarkdown and `knitr` for generating the website, and the already existing integration
398 between Github and Zenodo to archive the forecasts. By using this “continuous analysis”
399 approach (Beaulieu-Jones & Greene, 2017), where analyses are automatically rerun
400 when changes are made to data, models, or associated code, we have reduced the time
401 required by scientists to run and maintain the forecasting pipeline. To make the system
402 extensible so that new models could be easily incorporated, we used a plugin-based
403 infrastructure so that adding a new model to the system is as easy as adding a single file
404 to the ‘models’ folder in our repository (Figure 2). This should substantially lower the
405 barriers to other scientists contributing models to this forecasting effort. We also
406 automatically archive the resulting forecasts publicly so that the performance of these
407 forecasts can be assessed by both us and other researchers as new data is collected. This
408 serves as a form of pre-registration by providing a quantitative record of the forecast
409 before the data being predicted were collected.

410 While building this system was facilitated by the use of existing technological solutions,
411 there were still a number of challenges in making existing tools work for automated
412 iterative forecasting. Continuous integration is designed primarily for running
413 automated tests on software, not for running a coordinated forecasting pipeline. As a
414 result, extra effort was sometimes necessary to figure out how to get these systems to
415 work properly in non-standard situations, like running code that was not part of a
416 software package. In addition, hosted continuous integration solutions, like Travis,
417 provide only limited computational resources. As the number and complexity of the
418 models we fit has grown, we have had to continually invest effort in reducing our total
419 compute time so we can stay within these limits. Finally, we found no satisfactory
420 existing solution for archiving our results. All approaches we tried had limitations when
421 it came to automatically generating publicly-versioned archives of forecasts on a
422 repeated basis, and our eventual solution was difficult to configure to such a degree that
423 it will remain an impediment for most researchers. Overall, we found existing
424 technology to be sufficient to the task of creating an iterative forecasting pipeline, but it
425 required greater expertise and a greater investment of time than is ideal. Additional tool
426 development to reduce the effort required for scientists to set up their own short-term
427 forecasting systems would clearly be useful. Our efforts, however, show that it is
428 possible to use existing tools to develop initial iterative systems as a method for both
429 advancing scientific understanding and developing proof of concept forecasting systems.

430 Because of the breadth of expertise needed to set up our forecasting pipeline, our effort
431 required a team with diverse skills and perspectives, ranging from software
432 development to field site expertise. It is rare to find such breadth within a single
433 individual, and our system was developed as a collaboration between the lab collecting
434 the data and a computational ecology lab. When teams have a breadth of expertise,
435 communication can be challenging (Winowiecki et al., 2011). We found a shared base
436 of knowledge related to both the field research and computational skills was important

437 for the success of the group. The two labs are part of a joint interdisciplinary ecology
438 group that has a mission of breaking down barriers between field and
439 computational/theoretical ecologists (<http://weecology.org>). Everyone on the team had
440 received training in fundamental data management and computing skills through a
441 combination of university courses, Software and Data Carpentry workshops (Teal et al.,
442 2015), and lab training efforts. In addition, everyone was broadly familiar with the
443 study site and methods of data collection, and most team members had participated in
444 field work at the site on multiple occasions. This provided a shared set of knowledge
445 and vocabulary that actively facilitated interdisciplinary interactions. All members of
446 the team actively participated in the development of the forecasting pipeline. Given the
447 current state of tools for forecasting, forecasting teams require some experience in
448 working with continuous integration and APIs. This means either interdisciplinary
449 teams or additional training will often be required for creating these pipelines until tool
450 development improves. To improve the success of these diverse groups, we believe
451 efforts at providing ‘team science’ training to scientists interested in forecasting will be
452 beneficial for the success of iterative forecasting attempts for the foreseeable future
453 (Read et al., 2016).

454 We developed infrastructure for automatically making iterative forecasts with the goals
455 of making accurate forecasts for this well-studied system, learning what methods work
456 well for ecological forecasting more generally, and improving our understanding of the
457 processes driving ecological dynamics. The most obvious application of automated
458 iterative ecological forecasting is for speeding up development of forecasting models by
459 using the most recent data available and by quickly iterating to improve the models used
460 for forecasting. By learning what works best for forecasting in this and other ecological
461 systems, we will better understand what the best approaches are for ecological
462 forecasting more generally. By designing the pipeline so that it can forecast many
463 different aspects of the ecological community, we also hope to learn about what aspects

464 of ecology are more forecastable. Finally, automated forecasting infrastructures like this
465 one also provide a core foundation for faster scientific inquiry because new models can
466 quickly be applied to data and compared to existing models. The forecasting
467 infrastructure does the time-consuming work of data processing, data integration, and
468 model assessment, allowing new research to focus on the models being developed and
469 the inferences about the system that can be drawn from them (Dietze et al., 2018). We
470 plan to use this pipeline to drive future research into understanding the processes that
471 govern the dynamics of individual populations and the community as a whole. By
472 regularly running different models for population and community dynamics, a near-term
473 iterative pipeline such as ours should also make it possible to rapidly detect changes in
474 how the system is operating, which should allow the rapid identification of ecological
475 transitions or even possibly allow them to be prevented (Pace et al., 2017). By building
476 an automated iterative near-term forecasting infrastructure, we can improve our ability
477 to forecast natural systems, understand the biology driving ecological dynamics, and
478 detect or even predict changes in system state that are important for conservation and
479 management.

480 **Acknowledgements**

481 We thank Henry Senyondo for help with continuous integration and Hao Ye for
482 discussions and feedback on the manuscript. We thank all of the graduate students,
483 postdocs, and volunteers who have collected the Portal Project over the last 40 years
484 and the developers of the software and tools that made this project possible. This
485 research was supported by the National Science Foundation through grant 1622425 to
486 S.K.M. Ernest and by the Gordon and Betty Moore Foundation's Data-Driven
487 Discovery Initiative through grant GBMF4563 to E.P. White.

488 **Data Accessibility**

489 The data used in this study is from the Portal Project and is openly available (CC0) on
490 GitHub (<https://github.com/weecology/PortalData>). Code for reproducing all analyses is
491 available on GitHub (<https://github.com/weecology/portalPredictions>) and archived on
492 Zenodo (White et al., 2018b). Forecasts made by this system are all archived to Zenodo
493 (White et al., 2018a).

494 **Box 1. Key practices for automated iterative near-term** 495 **ecological forecasting**

496 A list of some of the key practices developed by Dietze et al (2018) for facilitating
497 iterative near-term ecological forecasting and discussion of why these practices are
498 important.

499 **Data**

500 **1. Frequent data collection**

501 Frequent data collection allows models to be regularly updated and forecasts to be
502 frequently evaluated (Dietze et al., 2018). Depending on the system being studied, this
503 frequency could range from sub-daily to annual, but typically the more frequently the
504 data is collected the better.

505 **2. Rapid data release under open licenses**

506 Data should be released as quickly as possible (low latency) under open licenses so that
507 forecasts can be made frequently and data can be accessed by a community of
508 forecasters (Vargas et al., 2017; Dietze et al., 2018).

509 **3. Best practices in data structure**

510 To reduce the time and effort needed to incorporate data into models, best practices in
511 data structure should be employed for managing and storing collected data to ensure it
512 is easy to integrate into other systems (interoperability) (Borer et al., 2009; Strasser et
513 al., 2011; White et al., 2013).

514 **Models**

515 **4. Focus on uncertainty**

516 Understanding the uncertainty of forecasts is crucial to interpreting and understanding
517 their utility. Models used for forecasting should be probabilistic to properly quantify
518 uncertainty and to convey how this uncertainty increases through time. Evaluation of
519 forecast models should include assessment of how accurately they quantify uncertainty
520 as well as point estimates (Hooten & Hobbs, 2015).

521 **5. Compare forecasts to simple baselines**

522 Understanding how much information is present in a forecast requires comparing its
523 accuracy to simple baselines to see if the models yield improvements over the naive
524 expectation that the system is static (Harris et al., 2018).

525 **6. Compare and combine multiple modeling approaches**

526 To quickly learn about the best approaches to forecasting different aspects of ecology,
527 multiple modeling approaches should be compared (Harris et al., 2018). Different
528 modeling approaches should also be combined into ensemble models, which often
529 outperform single models for prediction (Weigel et al., 2008).

530 **Cyberinfrastructure**

531 In addition to improvements in data and models, iterative near-term forecasting requires
532 improved infrastructure and approaches to support continuous model development and
533 iterative forecasting (Dietze et al., 2018).

534 **7. Best practices in software development**

535 Best practices should be followed in the development of scientific software and
536 modeling to make it easier to maintain, integrate into pipelines, and build on by other
537 researchers. Key best practices include open licenses, good documentation, version
538 control, and cross-platform support (Wilson et al., 2014; Hampton et al., 2015).

539 **8. Support easy inclusion of new models**

540 To facilitate the comparison and ensembling of different modeling approaches, code for
541 fitting models and making forecasts should be easily extensible, to allow models
542 developed by different groups to be integrated into a single framework (Dietze et al.,
543 2018).

544 **9. Automated end-to-end reproducibility**

545 Each forecast iteration involves acquiring new data, refitting the models, and making
546 new forecasts. This should be done automatically without requiring human intervention.
547 Therefore, the process of making forecasts should emphasize end-to-end reproducibility,
548 including data, models, and evaluation (Stodden & Miguez, 2014), to allow the
549 forecasts to be easily rerun as new data becomes available (Dietze et al., 2018).

550 **10. Publicly archive forecasts**

551 Forecasts should be openly archived to demonstrate that the forecasts were made
552 without knowledge of the outcomes and to allow the community to assess and compare
553 the performance of different forecasting approaches both now and in the future (McGill,
554 2012; Tredennick et al., 2016; Dietze et al., 2018; Harris et al., 2018). Ideally, the

555 forecasts and evaluation of their performance should be automatically posted publicly in
556 a manner that is understandable by both scientists and the broader stakeholder
557 community.

558 **Box 2. Glossary of terms**

559 **CI.** ‘Continuous Integration.’ The practice of continuously building and testing a code
560 base as it is developed. **Data latency.** The time it takes for data to be available for use.
561 **Docker.** An open-source Linux program for containerization (see software container).
562 **git.** An open-source version control system. **GitHub.** A web-based host for git projects.
563 Other options for a similar service include GitLab or Bitbucket. **PortalData.** The git
564 repository for the Portal data, found on GitHub. **portalPredictions.** The git repository
565 for the forecasts made using Portal data, found on GitHub. **portalr.** An R package for
566 using the Portal data. **QA/QC.** ‘Quality Assurance.’ Testing the quality of a product.
567 ‘Quality Control.’ The process of ensuring the quality of a product. **Rocker.** A project
568 making it easy to use Docker containers in the R environment. **Software container.**
569 Allows a developer to package up an application with all of the parts it needs to run
570 reliably. **testthat.** R package used to set up automated testing for QA/QC. **Travis.** A
571 continuous integration service that integrates easily with GitHub and R. Examples of
572 similar programs are Jenkins or CodeShip. **Unit test.** A component of quality control in
573 which each smallest testable part of software is formally tested. **Zenodo.** An open data
574 archive that integrates easily with GitHub.

575 **References**

576 Allaire, J., Cheng, J., Xie, Y., McPherson, J., Chang, W., Allen, J., . . . Arslan, R. (2017).
577 *Rmarkdown: Dynamic documents for r*. Retrieved from

578 <https://CRAN.R-project.org/package=rmarkdown>

579 Araujo, M. B., & New, M. (2007). Ensemble forecasts of species distributions. *Trends*
 580 *in Ecology and Evolution*, 22(1), 42–47.

581 Bauer, P., Thorpe, A., & Brunet, G. (2015). The quiet revolution of numerical weather
 582 prediction. *Nature*, 525(7567), 47.

583 Beaulieu-Jones, B. K., & Greene, C. S. (2017). Reproducibility of computational
 584 workflows is automated using continuous analysis. *Nature Biotechnology*, 35(4),
 585 342–346.

586 Bergman, C. (2012). On the preservation of published bioinformatics code on github.
 587 Retrieved from [https://caseybergman.wordpress.com/2012/11/08/](https://caseybergman.wordpress.com/2012/11/08/on-the-preservation-of-published-bioinformatics-code-on-github/)
 588 [on-the-preservation-of-published-bioinformatics-code-on-github/](https://caseybergman.wordpress.com/2012/11/08/on-the-preservation-of-published-bioinformatics-code-on-github/)

589 Boettiger, C., & Eddelbuettel, D. (2017). An introduction to rocker: Docker containers
 590 for r. *arXiv Preprint arXiv:1710.03675*.

591 Borer, E. T., Seabloom, E. W., Jones, M. B., & Schildhauer, M. (2009). Some simple
 592 guidelines for effective data management. *The Bulletin of the Ecological Society of*
 593 *America*, 90(2), 205–214.

594 Brown, J. H. (1998). The desert granivory experiments at portal. *Experimental Ecology*.
 595 *Oxford University Press, Oxford, UK*, 71–95.

596 Clark, J. S., Carpenter, S. R., Barber, M., Collins, S., Dobson, A., Foley, J. A., . . . others.
 597 (2001). Ecological forecasts: An emerging imperative. *Science*, 293(5530), 657–660.

598 Díaz, S., Demissew, S., Carabias, J., Joly, C., Lonsdale, M., Ash, N., . . . others. (2015).
 599 The ipbes conceptual framework—connecting nature and people. *Current Opinion in*

600 *Environmental Sustainability*, 14, 1–16.

601 Dietze, M. C. (2017). *Ecological forecasting*. Princeton University Press.

602 Dietze, M. C., Fox, A., Beck-Johnson, L. M., Betancourt, J. L., Hooten, M. B.,
603 Jarnevich, C. S., ... White, E. P. (2018). Iterative near-term ecological forecasting:
604 Needs, opportunities, and challenges. *Proceedings of the National Academy of Sciences*.
605 doi:10.1073/pnas.1710231115

606 Diniz-Filho, J. A. F., Bini, L. M., Rangel, T. F., Loyola, R. D., Hof, C., Nogues-Bravo,
607 D., & Araujo, M. B. (2009). Partitioning and mapping uncertainties in ensembles of
608 forecasts of species turnover under climate change. *Ecography*, 32(6), 897–906.

609 Ernest, S. K. M., Brown, J. H., Thibault, K. M., White, E. P., & Goheen, J. R. (2008).
610 Zero sum, the niche, and metacommunities: Long-term dynamics of community
611 assembly. *The American Naturalist*, 172(6), E257–E269.

612 Ernest, S. K. M., Valone, T. J., & Brown, J. H. (2009). Long-term monitoring and
613 experimental manipulation of a chihuahuan desert ecosystem near portal, arizona, usa.
614 *Ecology*, 90(6), 1708–1708.

615 Ernest, S. K. M., Yenni, G. M., Allington, G., Christensen, E. M., Geluso, K., Goheen, J.
616 R., ... Valone, T. J. (2016). Long-term monitoring and experimental manipulation of a
617 chihuahuan desert ecosystem near portal, arizona (1977–2013). *Ecology*, 97(4),
618 1082–1082.

619 Hampton, S. E., Anderson, S. S., Bagby, S. C., Gries, C., Han, X., Hart, E. M., ...
620 others. (2015). The tao of open science for ecology. *Ecosphere*, 6(7), 1–13.

621 Harris, D. J., Taylor, S. D., & White, E. P. (2018). Forecasting biodiversity in breeding
622 birds using best practices. *PeerJ*.

623 Hooten, M. B., & Hobbs, N. (2015). A guide to bayesian model selection for ecologists.

624 *Ecological Monographs*, 85(1), 3–28.

625 Jolliffe, I. T., & Stephenson, D. B. (Eds.). (2003). *Forecast verification: a practitioner's*
626 *guide in atmospheric science*. John Wiley; Sons, Ltd. Retrieved from
627 <http://linkinghub.elsevier.com/retrieve/pii/S0169207005001214>

628 Kalnay, E. (2003). *Atmospheric modeling, data assimilation and predictability*.
629 Cambridge university press.

630 Liboschik, T., Fokianos, K., & Fried, R. (2015). *Tscount: An r package for analysis of*
631 *count time series following generalized linear models*. Universitätsbibliothek
632 Dortmund.

633 Luo, Y., Ogle, K., Tucker, C., Fei, S., Gao, C., LaDeau, S., . . . Schimel, D. S. (2011).
634 Ecological forecasting and data assimilation in a data-rich era. *Ecological Applications*,
635 21(5), 1429–1442.

636 McGill, B. J. (2012). Ecologists need to do a better job of prediction – part ii – partly
637 cloudy and a 20% chance of extinction (or the 6 p's of good prediction). Retrieved from
638 [https://dynamicecology.wordpress.com/2013/01/09/](https://dynamicecology.wordpress.com/2013/01/09/ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/)
639 [ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/](https://dynamicecology.wordpress.com/2013/01/09/ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/)

640 Merkel, D. (2014). Docker: Lightweight linux containers for consistent development
641 and deployment. *Linux J.*, 2014(239). Retrieved from
642 <http://dl.acm.org/citation.cfm?id=2600239.2600241>

643 Morris, B. D., & White, E. P. (2013). The ecodata retriever: Improving access to
644 existing ecological data. *PloS One*, 8(6), e65848.

645 NOAA. (2016). NOAA fisheries (2016) status of stocks 2016: Annual report to
646 congress on the status of u.s. fisheries (noaa fisheries, washington, dc). Retrieved from
647 <https://doi.org/10.5281/zenodo.833438>

648 Pace, M. L., Batt, R. D., Buelo, C. D., Carpenter, S. R., Cole, J. J., Kurtzweil, J. T., &

649 Wilkinson, G. M. (2017). Reversal of a cyanobacterial bloom in response to early
650 warnings. *Proceedings of the National Academy of Sciences*, 114(2), 352–357.

651 Petchey, O. L., Pontarp, M., Massie, T. M., Kéfi, S., Ozgul, A., Weilenmann, M., ...
652 others. (2015). The ecological forecast horizon, and examples of its uses and
653 determinants. *Ecology Letters*, 18(7), 597–611.

654 Read, E. K., O'Rourke, M., Hong, G., Hanson, P., Winslow, L. A., Crowley, S., ...
655 Weathers, K. (2016). Building the team for team science. *Ecosphere*, 7(3).

656 Senyondo, H., Morris, B. D., Goel, A., Zhang, A., Narasimha, A., Negi, S., ... White,
657 E. P. (2017). Retriever: Data retrieval tool. *The Journal of Open Source Software*, 2(19),
658 451. doi:10.21105/joss.00451

659 Stodden, V., & Miguez, S. (2014). Best practices for computational science: Software
660 infrastructure and environments for reproducible and extensible research. *Journal of*
661 *Open Research Software*, 2(1).

662 Strasser, C., Cook, R., Michener, W., Budden, A., & Koskela, R. (2011). Promoting data
663 stewardship through best practices. In *Proceedings of the environmental information*
664 *management conference 2011 (eim 2011)*. Oak Ridge National Laboratory (ORNL).

665 Stumpf, R. P., Tomlinson, M. C., Calkins, J. A., Kirkpatrick, B., Fisher, K., Nierenberg,
666 K., ... Wynne, T. T. (2009). Skill assessment for an operational algal bloom forecast
667 system. *Journal of Marine Systems*, 76(1-2), 151–161.

668 Tallis, H. M., & Kareiva, P. (2006). Shaping global environmental decisions using
669 socio-ecological models. *Trends in Ecology & Evolution*, 21(10), 562–568.

670 Teal, T. K., Cranston, K. A., Lapp, H., White, E., Wilson, G., Ram, K., & Pawlik, A.
671 (2015). Data carpentry: Workshops to increase data literacy for researchers.
672 *International Journal of Digital Curation*, 10(1), 135–143.

673 Tredennick, A. T., Hooten, M. B., Aldridge, C. L., Homer, C. G., Kleinhesselink, A. R.,

674 & Adler, P. B. (2016). Forecasting climate change impacts on plant populations over
675 large spatial extents. *Ecosphere*, 7(10).

676 Vargas, R., Alcaraz-Segura, D., Birdsey, R., Brunsell, N. A., Cruz-Gaistardo, C. O.,
677 Jong, B. de, ... others. (2017). Enhancing interoperability to facilitate implementation
678 of redd+: Case study of mexico. *Carbon Management*, 8(1), 57–65.

679 Weigel, A. P., Liniger, M., & Appenzeller, C. (2008). Can multi-model combination
680 really enhance the prediction skill of probabilistic ensemble forecasts? *Quarterly*
681 *Journal of the Royal Meteorological Society*, 134(630), 241–260.

682 White, E. P. (2015). Some thoughts on best publishing practices for scientific software.
683 *Ideas in Ecology and Evolution*, 8(1).

684 White, E. P., Baldridge, E., Brym, Z. T., Locey, K. J., McGlinn, D. J., & Supp, S. R.
685 (2013). Nine simple ways to make it easier to (re) use your data. *Ideas in Ecology and*
686 *Evolution*, 6(2).

687 White, E. P., Bledsoe, E. K., Christensen, E. M., Harris, D. J., Simonis, J. L., Taylor, S.
688 D., ... Ernest, S. K. M. (2018a, February). Weecology/forecasts: 2018-02-17.
689 doi:10.5281/zenodo.839580

690 White, E. P., Yenni, G. M., Taylor, C., Shawn D and, Bledsoe, E. K., Simonis, J. L., &
691 Ernest, S. K. M. (2018b, February). Weecology/portalPredictions 2018-02-16.
692 doi:10.5281/zenodo.833438

693 Wickham, H. (2011). Testthat: Get started with testing. *The R Journal*, 3, 5–10.
694 Retrieved from
695 http://journal.r-project.org/archive/2011-1/RJournal_2011-1_Wickham.pdf

696 Wilson, G., Aruliah, D. A., Brown, C. T., Hong, N. P. C., Davis, M., Guy, R. T., ...
697 others. (2014). Best practices for scientific computing. *PLoS Biology*, 12(1), e1001745.

698 Winowiecki, L., Smukler, S., Shirley, K., Remans, R., Peltier, G., Lothes, E., ...

699 Alkema, L. (2011). Tools for enhancing interdisciplinary communication.
700 *Sustainability: Science, Practice and Policy*, 7(1), 74–80.

701 Xie, Y. (2015). *Dynamic documents with r and knitr* (Vol. 29). CRC Press.

702 Yenni, G. M., Christensen, E. M., Bledsoe, E. K., Supp, S. R., Diaz, R. M., White, E. P.,
703 & Ernest, S. M. (2018). Developing a modern data workflow for living data. *bioRxiv*,
704 344804.