

Automated iterative forecasting for the Portal Project

Introduction

Forecasting the future state of ecological systems is important for management, conservation, and evaluation of our basic understanding of how ecology operates (Clark et al., 2001; Tallis & Kareiva, 2006; Díaz et al., 2015; Dietze, 2017). In 2001, Clark et al. [-clark2001] called for a more central role of forecasting in ecology, and since then an increasing number of ecological forecasts are being published. However, most of these forecasts are made once, published, and never assessed or updated. Without assessment, we have limited information on how much confidence to place in our predictions; without regular updates, forecasts lack the most up-to-date information as conditions change (Dietze et al., 2016). This lack of both regular assessment and active updating has limited the progress of ecological forecasting and hindered our ability to make useful and reliable predictions. More regular updating and assessment will advance ecological forecasting as a field by speeding up the identification of the best models for forecasting specific issues in specific ecosystems and contribute overall to our understanding of how to best design forecasting approaches for ecology. For ecological forecasting to mature as a field, we need to change how we produce and interact with forecasts, creating a more dynamic interplay between model development, prediction generation, and incorporation of new data and information (Dietze et al., 2016).

With the goal of making ecological forecasting more dynamic and responsive, Dietze et al. [-dietze2018] recently called for an increase in iterative near-term forecasting.

Iterative near-term forecasting means making forecasts for the near future and making

24 these forecasts repeatedly through a cycle of forecast evaluation, integration of updated
25 data, and generation of new forecasts. This approach to forecasting has a number of
26 advantages. Because forecasts are made ‘near-term’—daily to annual forecasting
27 instead of multi-decadal—predictions can be assessed more quickly and frequently,
28 leading to more rapid model improvements (Dietze et al., 2016; Tredennick et al., 2016).
29 Because the forecasts are made repeatedly through time, new data can be integrated
30 with each new iteration of the forecast cycle. This iterative approach to forecasting
31 allows any changes in the state of the system that have occurred since the previous
32 forecast to be incorporated and accounted for, thereby improving the accuracy of future
33 forecasts (Dietze et al., 2016). Iterative near-term forecasting has the potential to
34 promote rapid improvement in the state of ecological forecasting by quickly identifying
35 how models are failing, facilitating rapid testing of improved models, and incorporating
36 updated data so models run with the most up-to-date information on the system
37 available. While use of iterative near-term forecasting is often discussed in the context
38 of policy and management (???; Clark et al., 2001; Petchey et al., 2015), this approach
39 to model testing can also be used to improve our basic understanding of ecological
40 systems (Dietze, 2017). For example, alternative mechanistic models can be competed
41 against each other to see which model provides the best forecasts for near-term
42 dynamics, thus providing insights into the relative importance of different processes
43 driving dynamics of ecological systems (Dietze et al., 2016). Whether deployed for
44 basic or applied uses, iterative near-term forecasting incorporates a more dynamic
45 interplay between models, predictions, and data that is clearly needed to improve
46 ecological forecasting and our understanding of ecological systems more broadly.

47 Because iterative near-term forecasting requires a dynamic interplay of models,
48 predictions, and data, Dietze et al [-dietze2018] highlight approaches to data
49 management, model construction and evaluation, and cyberinfrastructure that are
50 necessary to effectively implement this type of forecasting (Box 1). Data to be used for

51 iterative near-term forecasting needs to be widely accessible, which requires data to be
52 released quickly under open licenses (Dietze et al., 2016; Vargas et al., 2017) and
53 structured so that it can be used easily by a variety of researchers and in multiple
54 modeling approaches (Borer, Seabloom, Jones, & Schildhauer, 2009; Strasser, Cook,
55 Michener, Budden, & Koskela, 2011). Models need to be able to deal with uncertainty,
56 in the predictors and the predictions, to properly convey uncertainty in the resulting
57 forecasts (Diniz-Filho et al., 2009). Multiple models should be developed, both to assess
58 which models are performing best (Dietze et al., 2016) and to facilitate combining
59 models to form ensemble predictions which tend to perform better than single models
60 (Araujo & New, 2007; Diniz-Filho et al., 2009). Ensuring that data and models are
61 regularly updated and new forecasts are made requires cyberinfrastructure to automate
62 data processing, model fitting, prediction, model evaluation, forecast visualization, and
63 archiving. In combination, these approaches should allow forecasts to be easily rerun
64 and evaluated as new data becomes available (Box 1; Dietze et al., 2016).

65 While iterative near-term forecasting is an important next step in the evolution of
66 ecological forecasting, the requirements outlined by Dietze et al (Box 1) are not trivial
67 to implement; few of their recommendations are in widespread use in ecology today.
68 We examined what it would entail to operationalize Dietze et al's recommendations by
69 constructing our own iterative near-term forecasting pipeline for an on-going long-term
70 (~40 year) ecological study that collects high-frequency data on desert rodent
71 abundances (J. Brown, 1998; S. M. Ernest, Brown, Thibault, White, & Goheen, 2008).
72 We constructed our forecasting pipeline with the goal of being able to forecast rodent
73 abundances and evaluate our predictions on a monthly basis. In this paper, we discuss
74 our approach for creating this iterative near-term forecasting pipeline, the challenges we
75 encountered, the tools we used, and the lessons we learned so that others can create
76 their own iterative forecasting systems.

System Background

Iterative forecasting requires data that is collected repeatedly, and it benefits most from data that is collected frequently, as this provides more opportunities for updating model results and assessing (and potentially improving) model performance (Box 1; Dietze et al., 2016). The Portal Project is a long-term ecological study situated in the Chihuahuan Desert (2 km north and 6.5 km east of Portal, Arizona, US). Researchers have been continuously collecting data at the site since 1977, including data on the abundance of rodent and plant species (monthly and twice yearly, respectively) and climactic factors such as air temperature and precipitation (daily) (J. Brown, 1998; S. Ernest, Valone, & Brown, 2009; S. M. Ernest et al., 2016). The site consists of 24 50m x 50m experimental plots. Each plot contains 49 permanently marked trapping stations laid out in a 7 x 7 grid, and all plots are trapped with Sherman live traps for one night each month. For all rodents caught during a trapping session, information on species identity, size, and reproductive condition is collected, and new individuals are given identification tags. This data on rodent populations is high-frequency, uses consistent trapping methodology, and has an extended time-series (469 monthly samples and counting), making this study an ideal case for near-term iterative forecasting.

Implementing an automated iterative forecasting system

Implementation of iterative forecasting requires the regular rebuilding of models with new raw data as it becomes available and the presentation of those forecasts in usable forms; in our case, this occurs monthly. Rebuilding models in an efficient and maintainable way relies on developing an automated pipeline to handle the six stages of converting raw data into new forecasts: data collection, data sharing, data manipulation, modeling and forecasting, archiving, and presenting of the forecasts (Figure 1a). To implement the pipeline outlined in Figure 1a, we used a “continuous analysis”

102 framework (*sensu* Beaulieu-Jones & Greene, 2017) that automatically processes the
103 most up-to-date data, refits the models, makes new forecasts, archives the forecasts, and
104 updates a website with analysis of current and previous forecasts. In this section we
105 describe our approach to streamlining and automating the multiple components of the
106 forecasting pipeline and the tools and infrastructure we employed to execute each stage
107 of the pipeline.

108 **Continuous Analysis Framework**

109 A core aspect of iterative near-term forecasting is the regular rerunning of the
110 forecasting pipeline. We chose to have the computer run our pipeline and employed
111 “continuous analysis” (*sensu* Beaulieu-Jones & Greene, 2017) to drive the automation of
112 both the full pipeline and a number of its individual components. Continuous analysis
113 uses a set of tools originally designed for software development called “continuous
114 integration” (CI). CI combines computing environments for running code with
115 monitoring systems to identify changes in data or code. Essentially, CI is a computer
116 helper whose job is to watch the pipeline and, when it sees a change in the code or data,
117 it runs all the computer scripts needed to ensure that the forecasting pipeline runs from
118 beginning to end. This is useful for iterative near-term forecasting because it does not
119 rely on humans to remember to create forecasts when new models or data are added.
120 These tools are common in the area of software development where they are used to
121 automate software testing and integrate work by multiple developers working on the
122 same code base. However, these tools can be used for any computational task that needs
123 to be regularly repeated or run after changes to code or data (Beaulieu-Jones & Greene,
124 2017). Our forecasting pipeline currently runs on a publicly available continuous
125 integration service (Travis CI; <https://travis-ci.org/>) that is free for open source projects
126 (up to a limited amount of computing time). Because of the widespread use of CI in
127 software development, alternative services that can run code on local or cloud-based

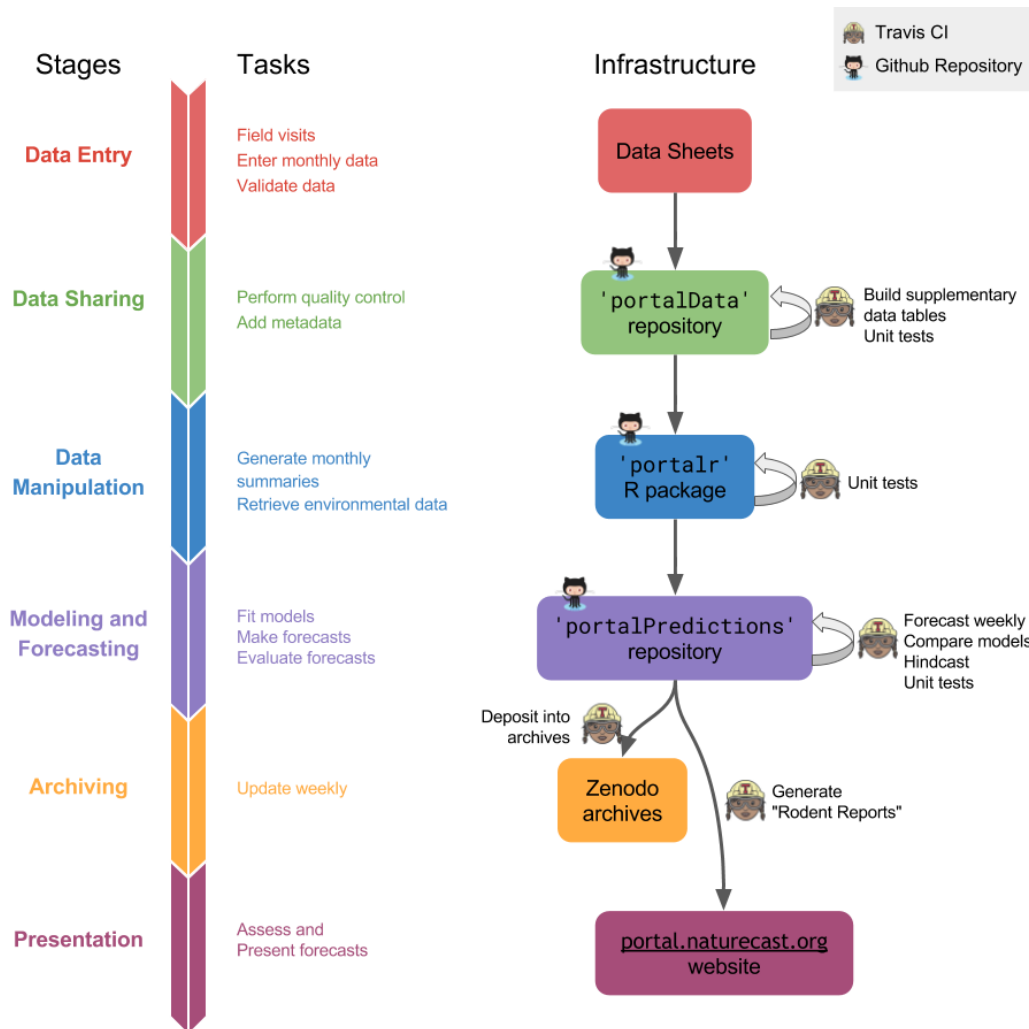


Figure 1: Figure 1. a) Stages of the forecasting pipeline. To go from raw data to forecast presentation involves a number of stages, each of which requires unique tasks, tools and infrastructure. While each stage has unique tasks associated with it, the stages of our pipeline are also interdependent. Our pipeline is primarily a linear structure where outputs from one stage form the inputs for the subsequent stage. All stages require code written in R. b) Continuous integration system. Each box denotes the core infrastructure used for each stage of the forecasting pipeline denoted in Figure 1a. Continuous integration, denoted here with the Travis icon (a girl wearing safety glasses and hardhat), plugs into our pipeline at every stage. Travis triggers the code involved in major events that involve integration across stages of the pipeline, such as taking the output from the forecasting stage (purple box) to create an updated presentation (rose box). Travis also runs tasks within a stage, such as conducting tests to make sure code changes have not introduced errors (icons on arrows originating and ending on the same box).

128 computational infrastructure also exist (e.g., Drone (<http://try.drone.io/>))
129 (Beaulieu-Jones & Greene, 2017). As detailed below, we use CI to quality check data,
130 test code using “unit tests” (Wilson et al., 2014), build models, make forecasts, and
131 publicly present and archive the results (Figure 1b).

132 In addition to automatically running software pipelines, the other key component of
133 “continuous analysis” is making sure that the pipelines will continue to run even as
134 software dependencies change (Beaulieu-Jones & Greene, 2017). Many of us have
135 experienced the frustrations that can occur when software updates (e.g., changes in R
136 package versions) create errors in previously functional code. We experienced this issue
137 when a package one of our models relies on, `tscount` (Liboschik, Fokianos, & Fried,
138 2015), was temporarily removed from CRAN (the R package repository) and, therefore,
139 could not be installed in the usual way. This broke our forecasting pipeline because we
140 could no longer run models that used that package. To make our pipeline robust to
141 changes in external software dependencies, we follow Beaulieu and Greene’s (2017)
142 recommendation to use software containers. Software containers are standalone
143 packages that contain copies of everything you need to run some piece of software,
144 including the operating system. Once created, a software container is basically a time
145 capsule, it contains all the software dependencies in the exact state used to develop and
146 run the software. If those dependencies change (or disappear) in the wider world, they
147 still exist, unchanged, in your container. We use an existing platform, Docker (???), to
148 store an exact image of the complete software environment for running the forecasts.
149 Docker also allows a specified set of packages to be used consistently across different
150 computer and server environments. Using containers allows us to update to new
151 package versions after testing and making any necessary changes to the data processing
152 and analysis code. We use a container created by the Rocker project which is a Docker
153 image with many important R packages (i.e. tidyverse) pre-installed (Boettiger &
154 Eddelbuettel, 2017). We add our code and dependencies to this existing Rocker image

155 to create a software container for our forecasting pipeline. In combination, the
156 automated running of the pipeline (continuous integration) and the guarantee it will not
157 stop working unexpectedly due to software dependencies (via a software container)
158 allows continuous analysis to serve as the glue that connects all stages of the forecasting
159 pipeline.

160 **Data Collection, Entry, and Processing**

161 Iterative forecasting benefits from frequently updated data so that state changes can be
162 quickly incorporated into new forecasts (Dietze et al., 2016). Frequent data collection
163 and rapid processing are both important for providing timely forecasts. Since our data is
164 collected monthly, ensuring that the models have access to the newest data requires a
165 data latency period of less than 1 month from collection to availability for modeling. To
166 accomplish this, we automated components of the data processing and quality
167 assurance/quality control (QA/QC) process to reduce the time needed to add new data
168 to the database (Figure 1).

169 New data is double-entered into Microsoft Excel using the “data validation” feature.
170 The two versions are then compared in an R script to control for errors in data entry.
171 Quality control (QC) checks using the `testthat` R package (Wickham, 2011) are run
172 on the data to test for validity and consistency both within the new data and between the
173 new and archived data. The local use of the QC scripts to flag problematic data greatly
174 reduces the time spent error-checking and ensures that the quality of data is consistent.
175 The data is then uploaded to the GitHub-based Portal Data repository
176 (<https://github.com/weecology/PortalData>). GitHub (<https://github.com/>) is a software
177 development tool for managing computer code development, but we have also found it
178 useful for data management. On GitHub, changes to data can be tracked through the Git
179 version control system which logs all changes made to any files in the repository. All
180 updates to data are processed through “pull requests”, which are notifications that

181 someone has a modified version of the data to contribute. QA/QC checks are
182 automatically rerun on the submitted data using continuous integration to ensure that
183 these checks have been run and that no avoidable errors reach the official version of the
184 dataset.

185 We also automated the updating of supplementary data tables, including information on
186 weather and trapping history, that were previously updated manually. As soon as new
187 field data is merged into the repository, continuous integration updates all
188 supplementary files. Weather data is automatically fetched from our cellular-connected
189 weather station, cleaned, and appended to the weather data table. Supplementary data
190 tables related to trapping history are updated based on the data added to the main data
191 tables. Using CI for this ensures that all supplementary data tables are always
192 up-to-date with the core data.

193 **Data Sharing**

194 The Portal Project has a long history of making its data publicly available, which means
195 that anyone can use it for forecasting or other projects. Historically the publication of
196 the data was conducted through data papers (S. Ernest et al., 2009, S. M. Ernest et al.
197 (2016)), the most common approach in ecology; however, this approach caused years of
198 data latency. Recently, the project has switched to posting data directly to a public
199 GitHub repository with a CC0 (i.e. no rights reserved) license (Figure 1). This
200 immediate posting reduces that data latency to less than one month and, therefore,
201 makes meaningful iterative near-term forecasting possible for not only our group but
202 other interested parties, as well.

203 **Data Manipulation**

204 Once data is available, it needs to be processed into a form appropriate for modeling
205 (Figure 1). For many ecological datasets, this requires not only simple data
206 manipulation but also a good understanding of the data to facilitate appropriate
207 aggregation. Data manipulation steps are often conducted using custom one-off code to
208 convert the raw data into the desired form (Morris & White, 2013), but this approach
209 has several limitations. First, each researcher must develop and maintain their own data
210 manipulation code, which is inefficient and can result in different researchers producing
211 different versions of the data for the same task. Subtle differences in data processing
212 decisions have led to confusion when reproducing results for the Portal data in the past.
213 Second, this kind of code is rarely robust to changes in data structure and location.
214 Based on our experience developing and maintaining the Data Retriever (Morris &
215 White, 2013; Senyondo et al., 2017), these kinds of changes are common. Finally, this
216 kind of code is generally poorly tested, which can lead to errors based on mistakes in
217 data manipulation. To avoid these issues for the Portal Project data, the Portal team has
218 been developing an R package (portalR; <http://github.com/weecology/portalr>) for
219 acquiring the data and handling common data cleaning and aggregation tasks. As a
220 result, our modeling and forecasting code only needs to install this package and run the
221 data manipulation and summary functions to get the appropriate data (Figure 1b). The
222 package undergoes thorough automated unit tests to ensure that data manipulations are
223 achieving the desired results. Having data manipulation code maintained in a separate
224 package that focuses on consistently providing properly summarized forms of the most
225 recent data has made maintaining the forecasting code itself much more straightforward.

226 **Modeling and Forecasting**

227 Iterative near-term forecasting involves regularly refitting a variety of different models
228 (Figure 1). Ideally, new models should be easy to incorporate to allow for iterative
229 improvements to the general modeling structure and approach. We use CI to refit the
230 models and make new forecasts each time the modeling code changes and when new
231 data becomes available (Figure 1b). We use a plugin infrastructure to allow new models
232 to be easily added to the system. This approach treats each model as an interchangeable
233 black boxes - all models have access to the same input data and generate the same
234 structure for model outputs (Figure 2). Each iteration of the forecasting cycle, all
235 models that are present are run and the standardized outputs are combined into a single
236 file to store the results of the different models' forecasts. A weighted ensemble model is
237 then added with weights based on how well individual models fit the training data. This
238 plugin infrastructure makes it easy to add and compare very different types of models,
239 from the basic time-series approaches currently implemented to complex state-space
240 and machine learning models. As long as a model script can load the provided data and
241 produce the appropriate output it will be run and its results incorporated into the rest of
242 the forecasting system.

243 In addition to flexibility in model internals, we also wanted to support flexibility in what
244 the models predict. Allowing models to make forecasts for system properties ranging
245 from individual species' population abundances to total community biomass facilitates
246 exploration of differences in forecastability across different aspects of ecological
247 systems. We designed a forecast output format to support this. Each forecast output file
248 contains the date being forecast, the collection date of the data used for fitting the
249 models, the date the forecast was made, the state variable being forecast (e.g., rodent
250 biomass, the abundance of a species), and the forecast value and associated uncertainty
251 of that forecast (Figure 2). This allows us to store a variety of different forecasts in a
252 common format and may serve as a useful starting point for developing a standard for

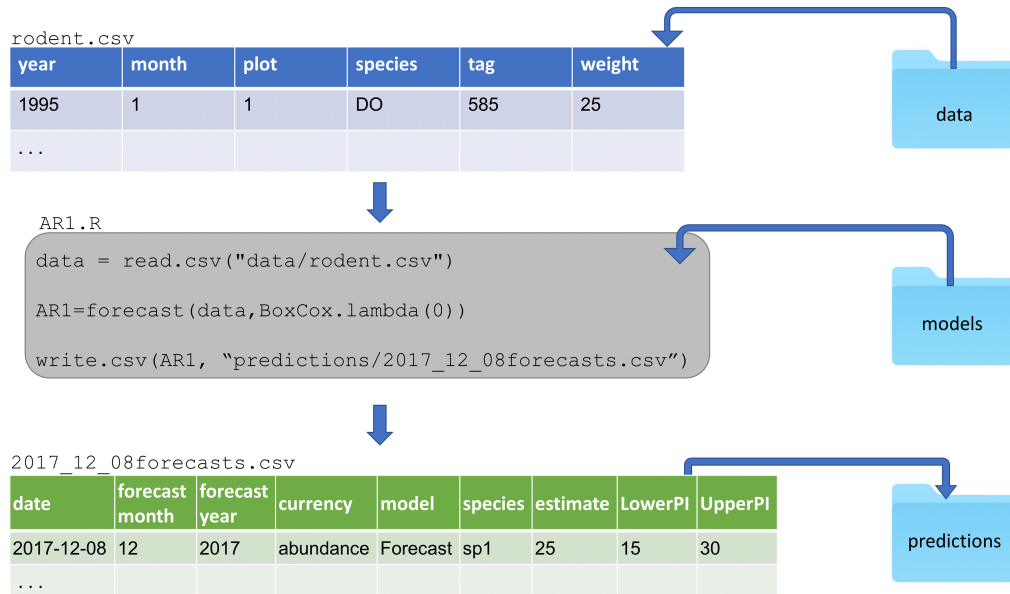


Figure 2: Figure 2. Demonstration of plugin infrastructure. All model scripts (represented here by the example AR1.R) are housed in a single folder. Each model script uses data provided by the core forecasting code (represented here by rodent.csv) and returns its forecast outputs in a predefined structure that is consistent across models (represented here by the example 2017_12_08forecasts.csv). Outputs from all models run on a particular date are combined into the same file (i.e. 2017_12_08forecasts.csv) to allow cross-model evaluations. Model output files are housed in a folder containing all forecast outputs from all previous dates to facilitate our archiving process and forecast assessment activities

253 storing ecological forecasts more generally.

254 Forecasts are currently evaluated using root mean square error (RMSE) to evaluate
255 point forecasts and coverage to evaluate uncertainty. We plan to add additional metrics
256 in the future. In addition to evaluating the actual forecasts, we also use hindcasting
257 (forecasting on already collected data) to gain additional insight into the methods that
258 work best for forecasting this system. For example, a model is fit using rodent
259 observations up to June 2005, then used to make a forecast 12 months out to May 2006.
260 The observations of that 12 month period can immediately be used to evaluate the
261 model. Hindcasting can be conducted using data that was already been collected, thus
262 allowing model comparison of large numbers of hindcasts and giving insight into which
263 models make the best forecasts without needing to wait for new data to be collected. It
264 can also be used to quickly evaluate new models instead of waiting for an adequate
265 amount of data to accumulate.

266 **Archiving**

267 Publicly archiving forecasts before new data is collected allows the field to assess,
268 compare, and build on forecasts made by different groups (McGill, 2012; Tredennick et
269 al., 2016; Dietze, 2017; Harris, Taylor, & White, 2018) (Figure 1). Archiving serves as
270 a form of pre-registration for model predictions, helping facilitate unbiased
271 interpretation of model performance. We wanted our archives to not only be accessible
272 to others but to also be a permanent record whose continued existence was not
273 dependent on any particular researcher or funding stream. Storing forecast outputs in
274 our forecasting GitHub repository was not sufficient for archival purposes because
275 repositories can be deleted. We explored three major repositories for archiving our
276 forecasts: FigShare (<https://figshare.com/>), Zenodo (<https://zenodo.org/>), and Open
277 Science Framework (<https://osf.io/>). While all three repositories allowed for easy
278 manual submissions (i.e., a human uploading files after each forecast), automating this

279 process was substantially more difficult. Various combinations of repositories, APIs
280 (i.e., interfaces that allow automated ways interacting with these websites) and
281 associated R packages had issues with: 1) integrating authorization with continuous
282 integration; 2) automatically making archived files public; 3) adding new files to an
283 existing location; or 4) automatically permanently archiving the files. Our eventual
284 solution was to leverage the GitHub-Zenodo integration
285 (<https://guides.github.com/activities/citable-code/>) and automatically push forecasts to
286 the GitHub repository from the CI server. The GitHub-Zenodo integration is designed
287 to automatically create versioned archives of GitHub repositories. There is an existing
288 one-time process for linking our forecasting repository on GitHub with Zenodo. Once
289 this link is created, each time a new forecast is created, our pipeline adds the new
290 forecasts to the GitHub repository and uses the GitHub API to create a new “release” of
291 our repository. This triggers the Zenodo-GitHub integration, which automatically
292 archives the resulting forecasts and the code that generated them under a top-level DOI
293 that refers to all archived forecasts (<https://doi.org/10.5281/zenodo.833438>). Through
294 this process, we automatically archive every forecast made with a documented history
295 of the archive. While this approach is functional because everything in the repository is
296 archived when a new forecast is made, these archives have complicated folder
297 structures, making it more harder for users to find and access the forecasting results.
298 Finding simpler approaches to automated archiving would increase the transparency of
299 the forecast results and their ease of use by external researchers.

300 **Presentation**

301 In addition to archiving the results of each forecast, we present them on a website that
302 displays monthly rodent forecasts, model evaluation metrics, monthly reports, and
303 information about the study site (Figure 3; <http://portal.naturecast.org>). The website
304 includes a graphical presentation of the most recent month’s forecasts (including

Portal Forecast

Total Abundance Forecast

This is the forecast for next month's sampling of rodents at Portal.

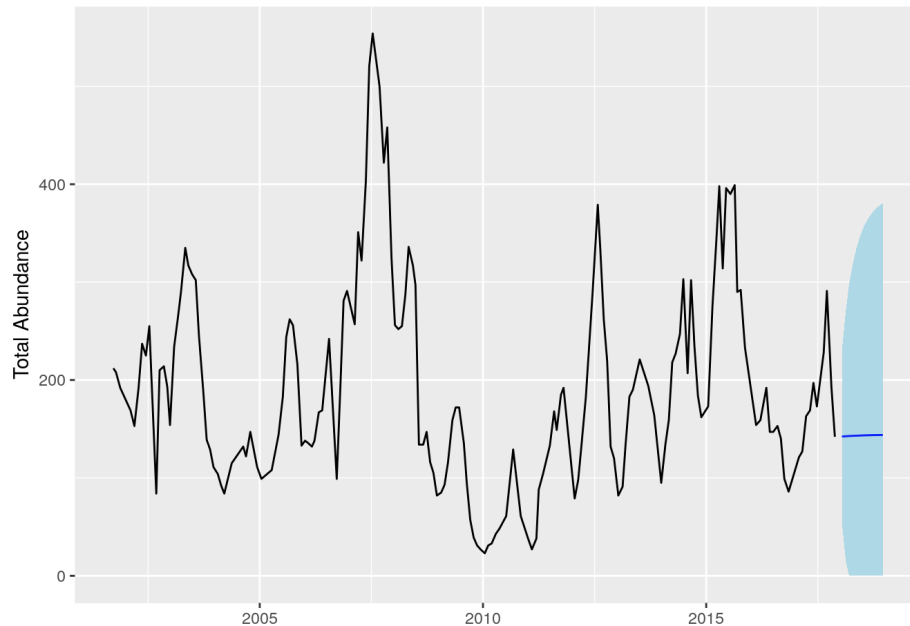


Figure 3: Figure 3. Screen capture of the homepage of the Portal Forecasting website; <http://portal.naturecast.org>. This site contains information on the most current forecasts, evaluation of forecast performance, and general information about the species being forecast.

305 uncertainty) and compares the latest data to the previous forecasts. Information on the
 306 the species and the field site targeted to a general audience are also included. The site is
 307 built using Rmarkdown (Allaire et al., 2017), which naturally integrates into the
 308 pipeline, and is automatically updated after each forecast. The `knitr` R package (Xie,
 309 2015) compiles the code into HTML, which is then published using Github Pages
 310 (<https://pages.github.com/>). The files for the website are stored in a subdirectory of the
 311 forecasting repository. As a result, the website is also archived automatically as part of
 312 the forecast archiving.

313 Discussion

314 Following the recommendations of Dietze et al [-dietze2018], we developed an
315 automated iterative forecasting system (Figure 1) to support repeated forecasting of an
316 ecological system. Our forecasting system automatically acquires and processes the
317 newest data, refits the models, makes new forecasts, publicly archives those forecasts,
318 and presents both the current forecast and information on how previous forecasts
319 performed. Every week our forecasting system generates a new set of forecasts with no
320 human intervention, except for the entry of new field data. This ensures that forecasts
321 based on the most recent data are always available and allows us to rapidly assess the
322 performance of multiple forecasting models for a number of different states of the
323 system, including the abundances of individual species and community-level variables
324 such as total abundance. To create this iterative near-term forecasting system, we used
325 R to process data and conduct analyses, and we leveraged already existing services
326 (i.e. GitHub, Travis, Docker) for more complicated cyberinfrastructure tasks. Thus, our
327 approach to developing iterative near-term forecasting infrastructure provides an
328 example for how short-term ecological forecasting systems can be initially developed.

329 We designed this forecasting system with the goal of making it relatively easy to build,
330 maintain, and extend. We used existing technology for both running the pipeline and
331 building individual components, which allowed us to build the system relatively cheaply
332 in terms of both time and money. This included the use of tools like Docker for
333 reproducibility, the Travis CI continuous integration system for automatically running
334 the pipeline, Rmarkdown and `knitr` for generating the website, and the already
335 existing integration between Github and Zenodo to archive the forecasts. By using this
336 “continuous analysis” approach (Beaulieu-Jones & Greene, 2017), where analyses are
337 automatically rerun when changes are made to data, models, or associated code, we
338 have reduced the time required by scientists to run and maintain the forecasting pipeline.
339 To make the system extensible so that new models could be easily incorporated, we

340 used a plugin-based infrastructure so that adding a new model to the system is as easy
341 as adding a single file to the ‘models’ folder in our repository (Figure 2). This should
342 substantially lower the barriers to other scientists contributing models to this forecasting
343 effort. We also automatically archive the resulting forecasts publicly so that the
344 performance of these forecasts can be assessed by both us and other researchers as new
345 data is collected. This serves as a form of pre-registration by providing a quantitative
346 record of the forecast before the data being predicted were collected.

347 While building this system was facilitated by the use of existing technological solutions,
348 there were still a number of challenges in making existing tools work for automated
349 iterative forecasting. Continuous integration is designed primarily for running
350 automated tests on software, not for running a coordinated forecasting pipeline. As a
351 result, extra effort was sometimes necessary to figure out how to get these systems to
352 work properly in non-standard situations, like running code that was not part of a
353 software package. In addition, hosted continuous integration solutions, like Travis,
354 provide only limited computational resources. As the number and complexity of the
355 models we fit has grown, we have had to continually invest effort in reducing our total
356 compute time so we can stay within these limits. Finally, we found no satisfactory
357 existing solution for archiving our results. All approaches we tried had limitations when
358 it came to automatically generating publicly versioned archives of forecasts on a
359 repeated basis. Overall, we found existing technology to be sufficient to the task, but it
360 required greater expertise and a greater investment of time than is ideal. Tool
361 development to reduce the effort required for scientists to set up their own short-term
362 forecasting systems would clearly be useful, but our efforts show that it is possible for
363 scientists to use existing tools to develop initial iterative systems as a method for both
364 advancing scientific understanding and developing proof of concept forecasting systems.
365 Developing systems that produce automated iterative forecasts will make it easier to
366 convert research focused forecasts into fully operationalized forecasting systems that

367 are relied on for decision making (Dietze et al., 2016, other paper by operationalization
368 co-authors).

369 Because of the breadth of expertise needed to set up our forecasting pipeline, our effort
370 required a team with diverse skills and perspectives, ranging from software
371 development to field site expertise. It is rare to find such breadth within a single
372 research group, and our system was developed as a collaboration between the lab
373 collecting the data and a computational ecology lab. When teams have a breadth of
374 expertise, communication can be challenging. We found a shared base of knowledge
375 related to both the field research and fundamental computational skills was important
376 for the success of the group. Everyone on the team had received training in fundamental
377 data management and computing skills through a combination of university courses,
378 Software and Data Carpentry workshops (Teal et al., 2015), and informal lab training
379 efforts. In addition, everyone was broadly familiar with the study site and methods of
380 data collection, and most team members had participated in field work at the site on
381 multiple occasions. This provided a shared set of knowledge and vocabulary that
382 actively facilitated interdisciplinary interactions. Given the current state of tools for
383 forecasting, forecasting teams will need people with significant experience in working
384 with continuous integration and APIs, which means interdisciplinary teams will
385 generally be required for creating these pipelines until tool development improves.

386 We developed this infrastructure for automatically making iterative forecasts with the
387 goals of making accurate forecasts for this well-studied system, learning what methods
388 work well for ecological forecasting more generally, and improving our understanding
389 of the processes driving ecological dynamics. The most obvious application of
390 automated iterative ecological forecasting is for speeding up development of forecasting
391 models by providing the most recent data available to models and by quickly iterating to
392 improve the models used for forecasting. By learning what works best for forecasting in
393 this and other ecological systems, we will better understand what the best approaches

are for ecological forecasting more generally. By designing the pipeline so that it can forecast many different aspects of the ecological community, we also hope to learn about what aspects of ecology are more forecastable. Finally, automated forecasting infrastructures like this one also provide a core foundation for faster scientific inquiry more broadly because new models can quickly be applied to data and compared to existing models. The forecasting infrastructure does the time-consuming work of data processing, data integration, and model assessment, allowing new research to focus on the models being developed and the inferences about the system that can be drawn from them (Dietze et al., 2016). We plan to use this pipeline to drive future research into understanding the processes that govern the dynamics of individual populations and the community as a whole. By regularly running different models for population and community dynamics, a near-term iterative pipeline such as ours should also make it possible to rapidly detect changes in how the system is operating, which should allow the rapid identification of ecological transitions or even possibly allow them to be prevented (??? example). By building an automated iterative near-term forecasting infrastructure we can improve our ability to forecast natural systems, our understanding of the biology driving ecological dynamics, and detect or even predict changes in system state that are important for conservation and management.

Acknowledgements

This research was supported by the National Science Foundation through grant 1622425 to S.K.M. Ernest and by the Gordon and Betty Moore Foundation's Data-Driven Discovery Initiative through grant GBMF4563 to E.P. White. We thank all of the graduate students, postdocs, and volunteers who have collected the Portal Project over the last 40 years and the developers of all of the software and tools that made this project possible.

419 **Box 1. Key practices for automated iterative near-term** 420 **ecological forecasting**

421 A list of some of the key practices developed by Dietze et al [-dietze2018] for
422 facilitating iterative near-term ecological forecasting and discussion of why these
423 practices are important.

424 **Data**

425 **1. Frequent data collection**

426 Frequent data collection allows models to be regularly updated and forecasts to be
427 frequently evaluated (Dietze et al., 2016). Depending on the system being studied, this
428 frequency could range from sub-daily to annual, but typically the more frequently the
429 data is collected the better.

430 **2. Rapid data release under open licenses**

431 Data should be released as quickly as possible (low latency) under open licenses so that
432 forecasts can be made frequently and data can be accessed by a community of
433 forecasters (Dietze et al., 2016; Vargas et al., 2017).

434 **3. Best practices in data structure**

435 To reduce the time and effort needed to incorporate this data into models, best practices
436 in data structure need to be employed for managing and storing collected data to ensure
437 it is easy to integrate into other systems (interoperability) (Borer et al., 2009; Strasser et
438 al., 2011; White et al., 2013).

439 **Models**

440 **4. Focus on uncertainty**

441 Understanding the uncertainty of forecasts is crucial to interpreting the forecasts and
442 understanding their utility. Models used for forecasting should be probabilistic to
443 properly quantify uncertainty and to convey how this uncertainty increases through time.
444 Evaluation of forecast models should include assessment of how accurately they
445 quantify uncertainty as well as point estimates. This can be done using “proper and
446 local” scores (Hooten & Hobbs, 2015).

447 **5. Compare forecasts to simple baselines**

448 Understanding how much information is present in a forecast requires comparing its
449 accuracy to simple baselines to see if the models yield improvements over the naive
450 expectation that the system is static (Harris et al., 2018).

451 **6. Compare and combine multiple modeling approaches**

452 To quickly learn about the best approaches to forecasting different aspects of ecology,
453 multiple modeling approaches should be compared for forecasting tasks (Harris et al.,
454 2018). Different modeling approaches should also be combined into ensemble models,
455 which are known to outperform single models for many forecasting and prediction tasks
456 (Weigel, Liniger, & Appenzeller, 2008).

457 **Cyberinfrastructure**

458 In addition to improvements in data and models, iterative near-term forecasting requires
459 improved infrastructure and approaches to support continuous model development and
460 iterative forecasting (Dietze et al., 2016).

461 **7. Best practices in software development**

Best practices should be followed in the development of scientific software and modeling to make it easier to maintain, integrate into pipelines, and build on by other researchers. Key best practices include using open licenses, good documentation, version control, and cross-platform support (Wilson et al., 2014; Hampton et al., 2015).

8. Support easy inclusion of new models

To facilitate the comparison and ensembling of different modeling approaches, code for fitting models and making forecasts should be easily extensible, allowing models developed by different groups to be easily integrated into a single framework (Dietze et al., 2016).

9. Automated end-to-end reproducibility

Iteratively making forecasts requires that acquiring the newest data, refitting the models, and making new forecasts is simple. Ideally, this should be done automatically without requiring human intervention. Therefore, the process of making forecasts should emphasize end-to-end reproducibility, including data, models, and evaluation (Stodden & Miguez, 2014), to allow the forecasts to be easily rerun as new data becomes available (Dietze et al., 2016). Ideally, the entire forecasting pipeline will be rerun automatically as new data becomes available.

10. Publicly archive forecasts

Forecasts should be openly archived to demonstrate that the forecasts were made without knowledge of the outcomes (i.e., as a form of pre-registration *sensu*) and to allow the community to assess and compare the performance of different forecasting approaches both now and in the future (McGill, 2012; Dietze et al., 2016; Tredennick et al., 2016; Harris et al., 2018). Ideally, the forecasts and evaluation of their performance should be automatically posted publicly in a manner that is understandable by both interested scientists and other stakeholders.

References

- Allaire, J., Cheng, J., Xie, Y., McPherson, J., Chang, W., Allen, J., . . . Arslan, R. (2017). *Rmarkdown: Dynamic documents for r*. Retrieved from <https://CRAN.R-project.org/package=rmarkdown>
- Araujo, M. B., & New, M. (2007). Ensemble forecasts of species distributions. *Trends in Ecology and Evolution*, 22(1), 42–47.
- Beaulieu-Jones, B. K., & Greene, C. S. (2017). Reproducibility of computational workflows is automated using continuous analysis. *Nature Biotechnology*, 35(4), 342–346.
- Boettiger, C., & Eddelbuettel, D. (2017). An introduction to rocker: Docker containers for r. *arXiv Preprint arXiv:1710.03675*.
- Borer, E. T., Seabloom, E. W., Jones, M. B., & Schildhauer, M. (2009). Some simple guidelines for effective data management. *The Bulletin of the Ecological Society of America*, 90(2), 205–214.
- Brown, J. (1998). The desert granivory experiments at portal. *Experimental Ecology. Oxford University Press, Oxford, UK*, 71–95.
- Clark, J. S., Carpenter, S. R., Barber, M., Collins, S., Dobson, A., Foley, J. A., . . . others. (2001). Ecological forecasts: An emerging imperative. *Science*, 293(5530), 657–660.
- Dietze, M. C. (2017). *Ecological forecasting*. Princeton University Press.
- Dietze, M. C., Fox, A., Betancourt, J. L., Hooten, M. B., Jarnevich, C. S., Keitt, T. H., . . . others. (2016). Iterative ecological forecasting: Needs, opportunities, and challenges. *Proceedings of the National Academy of Sciences*.
- Diniz-Filho, J. A. F., Bini, L. M., Rangel, T. F., Loyola, R. D., Hof, C., Nogues-Bravo, D., & Araujo, M. B. (2009). Partitioning and mapping uncertainties in ensembles of

511 forecasts of species turnover under climate change. *Ecography*, 32(6), 897–906.

512 Díaz, S., Demissew, S., Carabias, J., Joly, C., Lonsdale, M., Ash, N., . . . others. (2015).
 513 The ipbes conceptual framework—connecting nature and people. *Current Opinion in*
 514 *Environmental Sustainability*, 14, 1–16.

515 Ernest, S. M., Brown, J. H., Thibault, K. M., White, E. P., & Goheen, J. R. (2008). Zero
 516 sum, the niche, and metacommunities: Long-term dynamics of community assembly.
 517 *The American Naturalist*, 172(6), E257–E269.

518 Ernest, S. M., Yenni, G. M., Allington, G., Christensen, E. M., Geluso, K., Goheen, J.
 519 R., . . . others. (2016). Long-term monitoring and experimental manipulation of a
 520 chihuahuan desert ecosystem near portal, arizona (1977–2013). *Ecology*, 97(4),
 521 1082–1082.

522 Ernest, S., Valone, T. J., & Brown, J. H. (2009). Long-term monitoring and
 523 experimental manipulation of a chihuahuan desert ecosystem near portal, arizona, usa.
 524 *Ecology*, 90(6), 1708–1708.

525 Hampton, S. E., Anderson, S. S., Bagby, S. C., Gries, C., Han, X., Hart, E. M., . . .
 526 others. (2015). The tao of open science for ecology. *Ecosphere*, 6(7), 1–13.

527 Harris, D. J., Taylor, S. D., & White, E. P. (2018). Forecasting biodiversity in breeding
 528 birds using best practices. *PeerJ*.

529 Hooten, M. B., & Hobbs, N. (2015). A guide to bayesian model selection for ecologists.
 530 *Ecological Monographs*, 85(1), 3–28.

531 Liboschik, T., Fokianos, K., & Fried, R. (2015). *Tscount: An r package for analysis of*
 532 *count time series following generalized linear models*. Universitätsbibliothek
 533 Dortmund.

534 McGill, B. J. (2012). Ecologists need to do a better job of prediction – part ii – partly
 535 cloudy and a 20% chance of extinction (or the 6 p’s of good prediction). Retrieved from

536 <https://dynamicecology.wordpress.com/2013/01/09/>
 537 [ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/](#)

538 Morris, B. D., & White, E. P. (2013). The ecodata retriever: Improving access to
 539 existing ecological data. *PloS One*, 8(6), e65848.

540 Petchey, O. L., Pontarp, M., Massie, T. M., Kéfi, S., Ozgul, A., Weilenmann, M., ...
 541 others. (2015). The ecological forecast horizon, and examples of its uses and
 542 determinants. *Ecology Letters*, 18(7), 597–611.

543 Senyondo, H., Morris, B. D., Goel, A., Zhang, A., Narasimha, A., Negi, S., ... White,
 544 E. P. (2017). Retriever: Data retrieval tool. *The Journal of Open Source Software*, 2(19),
 545 451. doi:10.21105/joss.00451

546 Stodden, V., & Miguez, S. (2014). Best practices for computational science: Software
 547 infrastructure and environments for reproducible and extensible research. *Journal of*
 548 *Open Research Software*, 2(1).

549 Strasser, C., Cook, R., Michener, W., Budden, A., & Koskela, R. (2011). Promoting data
 550 stewardship through best practices. In *Proceedings of the environmental information*
 551 *management conference 2011 (eim 2011)*. Oak Ridge National Laboratory (ORNL).

552 Tallis, H. M., & Kareiva, P. (2006). Shaping global environmental decisions using
 553 socio-ecological models. *Trends in Ecology & Evolution*, 21(10), 562–568.

554 Teal, T. K., Cranston, K. A., Lapp, H., White, E., Wilson, G., Ram, K., & Pawlik, A.
 555 (2015). Data carpentry: Workshops to increase data literacy for researchers.
 556 *International Journal of Digital Curation*, 10(1), 135–143.

557 Tredennick, A. T., Hooten, M. B., Aldridge, C. L., Homer, C. G., Kleinhesselink, A. R.,
 558 & Adler, P. B. (2016). Forecasting climate change impacts on plant populations over
 559 large spatial extents. *Ecosphere*, 7(10).

560 Vargas, R., Alcaraz-Segura, D., Birdsey, R., Brunsell, N. A., Cruz-Gaistardo, C. O.,

561 Jong, B. de, ... others. (2017). Enhancing interoperability to facilitate implementation
 562 of redd+: Case study of mexico. *Carbon Management*, 8(1), 57–65.

563 Weigel, A. P., Liniger, M., & Appenzeller, C. (2008). Can multi-model combination
 564 really enhance the prediction skill of probabilistic ensemble forecasts? *Quarterly*
 565 *Journal of the Royal Meteorological Society*, 134(630), 241–260.

566 White, E. P., Baldrige, E., Brym, Z. T., Locey, K. J., McGlinn, D. J., & Supp, S. R.
 567 (2013). Nine simple ways to make it easier to (re) use your data. *Ideas in Ecology and*
 568 *Evolution*, 6(2).

569 Wickham, H. (2011). Testthat: Get started with testing. *The R Journal*, 3, 5–10.
 570 Retrieved from
 571 http://journal.r-project.org/archive/2011-1/RJournal_2011-1_Wickham.pdf

572 Wilson, G., Aruliah, D. A., Brown, C. T., Hong, N. P. C., Davis, M., Guy, R. T., ...
 573 others. (2014). Best practices for scientific computing. *PLoS Biology*, 12(1), e1001745.

574 Xie, Y. (2015). *Dynamic documents with r and knitr* (Vol. 29). CRC Press.