

# Automated iterative near-term forecasting for the Portal Project

## Introduction

Forecasting the future state of ecological systems is important for management, conservation, and evaluation of how well models capture the processes governing ecological systems (Clark et al., 2001; Tallis & Kareiva, 2006; Díaz et al., 2015; Dietze, 2017). In 2001, Clark et al. (2001) called for a more central role of forecasting in ecology, and since then an increasing number of ecological forecasts are being published. However, most of these forecasts are made once, published, and never assessed or updated. This lack of both regular assessment and active updating has limited the progress of ecological forecasting and hindered our ability to make useful and reliable predictions. The lack of active assessment results in limited information on how much confidence to place in forecasts and makes it difficult to determine which forecasting methods to build on. Without regular updates, forecasts lack the most current data and forecast accuracy will decay over time (???; Dietze et al., 2016). More regular updating and assessment will advance ecological forecasting as a field by accelerating the identification of the best models for individual forecasts and improving our understanding of how to best design forecasting approaches for ecology in general. For ecological forecasting to mature as a field, we need to change how we produce and interact with forecasts, creating a more dynamic interplay between model development, prediction generation, and incorporation of new data and information (Dietze et al., 2016).

With the goal of making ecological forecasting more dynamic and responsive, Dietze et

al (2016) recently called for an increase in iterative near-term forecasting. Iterative near-term forecasting is defined by making predictions for the near future and repeatedly updating those predictions through a cycle of evaluation, integration of new data, and generation of new forecasts. Because forecasts are made ‘near-term’—daily to annual time scales instead of multi-decadal—predictions can be assessed more quickly and frequently, leading to more rapid model improvements (Dietze et al., 2016; Tredennick et al., 2016). In addition, since forecasts are made repeatedly through time, new data can be continuously integrated with each iteration (Dietze et al., 2016). As a result, iterative near-term forecasting has the potential to promote rapid improvement in the state of ecological forecasting by quickly identifying how models are failing, facilitating rapid testing of improved models, and incorporating the most up-to-date data available. In addition to yielding improved information for guiding policy and management (Clark et al., 2001; Luo et al., 2011; Petchey et al., 2015), this iterative approach will help improve our basic understanding of ecological systems (Dietze et al., 2016). For example, alternative mechanistic models can be compared to determine which model provides the best forecasts, thus providing insights into the importance of different ecological processes (Dietze et al., 2016). Iterative near-term forecasting provides a more dynamic interplay between models, predictions, and data that will improve ecological forecasting and our understanding of ecological systems more broadly. Because iterative near-term forecasting requires a dynamic integration of models, predictions, and data, Dietze et al (2016) highlight approaches to data management, model construction and evaluation, and cyberinfrastructure that are necessary to effectively implement this type of forecasting (Box 1). Data needs to be released quickly under open licenses (Dietze et al., 2016; Vargas et al., 2017) and structured so that it can be used easily by a variety of researchers and in multiple modeling approaches (Borer, Seabloom, Jones, & Schildhauer, 2009; Strasser, Cook, Michener, Budden, & Koskela, 2011). Models need to be able to deal with uncertainty, in the

51 predictors and the predictions, to properly convey uncertainty in the resulting forecasts  
52 (Diniz-Filho et al., 2009). Multiple models should be developed, both to assess which  
53 models are performing best (Dietze et al., 2016) and to facilitate combining models to  
54 form ensemble predictions which tend to perform better than single models (Araujo &  
55 New, 2007; Diniz-Filho et al., 2009). Ensuring that data and models are regularly  
56 updated and new forecasts are made requires cyberinfrastructure to automate data  
57 processing, model fitting, prediction, model evaluation, forecast visualization, and  
58 archiving. In combination, these approaches should allow forecasts to be easily rerun  
59 and evaluated as new data becomes available (Box 1; Dietze et al., 2016).

60 While iterative near-term forecasting is an important next step in the evolution of  
61 ecological forecasting, the requirements outlined by Dietze et al (Box 1) are not trivial  
62 to implement and few of their recommendations are in widespread use in ecology today.  
63 We explored what it would entail to operationalize Dietze et al's recommendations by  
64 constructing our own iterative near-term forecasting pipeline for an on-going long-term  
65 ecological study that collects high-frequency data on desert rodent abundances (J.  
66 Brown, 1998; S. M. Ernest, Brown, Thibault, White, & Goheen, 2008). We constructed  
67 an automated forecasting pipeline with the goal of being able to forecast rodent  
68 abundances and evaluate our predictions on a monthly basis. In this paper, we discuss  
69 our approach for creating this iterative near-term forecasting pipeline, the challenges we  
70 encountered, the tools we used, and the lessons we learned so that others can create  
71 their own iterative forecasting systems.

## 72 **System Background**

73 Iterative forecasting is most effective with frequently collected data, since it provides  
74 more opportunities for updating model results and assessing (and potentially improving)  
75 model performance (Box 1; Dietze et al., 2016). The Portal Project is a long-term

ecological study situated in the Chihuahuan Desert (2 km north and 6.5 km east of Portal, Arizona, US). Researchers have been continuously collecting data at the site since 1977, including data on the abundance of rodent and plant species (monthly and twice yearly, respectively) and climatic factors such as air temperature and precipitation (daily) (J. Brown, 1998; S. Ernest, Valone, & Brown, 2009; S. M. Ernest et al., 2016). The site consists of 24 50m x 50m experimental plots. Each plot contains 49 permanently marked trapping stations laid out in a 7 x 7 grid, and all plots are trapped with Sherman live traps for one night each month. For all rodents caught during a trapping session, information on species identity, size, and reproductive condition is collected, and new individuals are given identification tags. This information on rodent populations is high-frequency, uses consistent trapping methodology, and has an extended time-series (469 monthly samples and counting), making this study an ideal case for near-term iterative forecasting.

## **Implementing an automated iterative forecasting system**

Implementation of iterative forecasting requires the regular rebuilding of models with new raw data as it becomes available and the presentation of those forecasts in usable forms; in our case, this occurs monthly. Rebuilding models in an efficient and maintainable way relies on developing an automated pipeline to handle the six stages of converting raw data into new forecasts: data collection, data sharing, data manipulation, modeling and forecasting, archiving, and presentation of the forecasts (Figure 1a). To implement the pipeline outlined in Figure 1a, we used a “continuous analysis” framework (*sensu* Beaulieu-Jones & Greene, 2017) that automatically processes the most up-to-date data, refits the models, makes new forecasts, archives the forecasts, and updates a website with analysis of current and previous forecasts. In this section we describe our approach to streamlining and automating the multiple components of the

101 forecasting pipeline and the tools and infrastructure we employed to execute each  
102 component.

## 103 **Continuous Analysis Framework**

104 A core aspect of iterative near-term forecasting is the regular rerunning of the  
105 forecasting pipeline. We employed “continuous analysis” (*sensu* Beaulieu-Jones &  
106 Greene, 2017) to drive the automation of both the full pipeline and a number of its  
107 individual components. Continuous analysis uses a set of tools originally designed for  
108 software development called “continuous integration” (CI). CI combines computing  
109 environments for running code with monitoring systems to identify changes in data or  
110 code. Essentially, CI is a computer helper whose job is to watch the pipeline and, when  
111 it sees a change in the code or data, it runs all the computer scripts needed to ensure that  
112 the forecasting pipeline runs from beginning to end. This is useful for iterative  
113 near-term forecasting because it does not rely on humans to create new forecasts  
114 whenever new models or data are added. These tools are common in the area of  
115 software development where they are used to automate software testing and integrate  
116 work by multiple developers working on the same code base. However, these tools can  
117 be used for any computational task that needs to be regularly repeated or run after  
118 changes to code or data (Beaulieu-Jones & Greene, 2017). Our forecasting pipeline  
119 currently runs on a publicly available continuous integration service (Travis CI;  
120 <https://travis-ci.org/>) that is free for open source projects (up to a limited amount of  
121 computing time). Because of the widespread use of CI in software development,  
122 alternative services that can run code on local or cloud-based computational  
123 infrastructure also exist (Beaulieu-Jones & Greene, 2017). We use CI to quality check  
124 data, test code using “unit tests” (Wilson et al., 2014), build models, make forecasts,  
125 and publicly present and archive the results (Figure 1b).

126 In addition to automatically running software pipelines, the other key component of

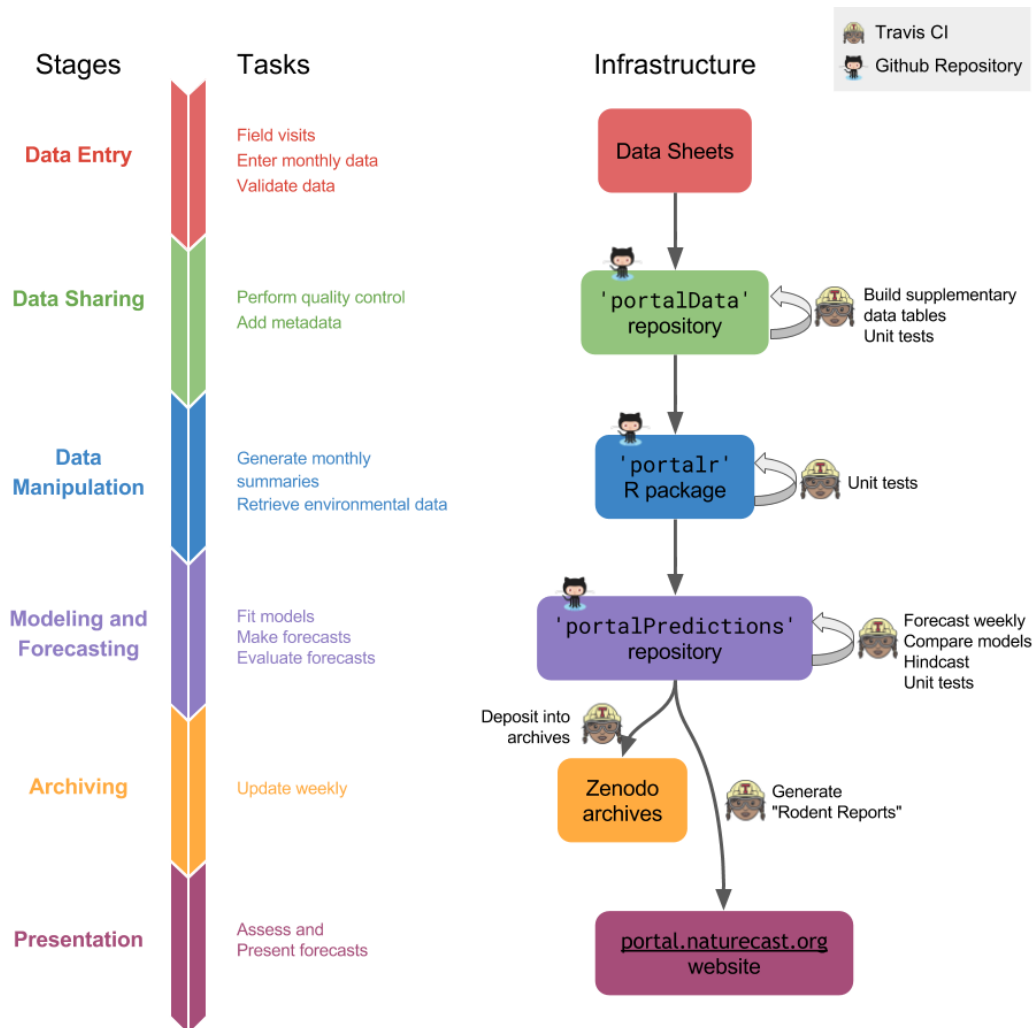


Figure 1: Figure 1. a) Stages of the forecasting pipeline. To go from raw data to forecast presentation involves a number of stages, each of which requires unique tasks, tools and infrastructure. The stages of are interdependent with outputs from one stage forming the inputs for the subsequent stage. Tasks in all stages are run using code written in R. b) Continuous integration system. Each box denotes the core infrastructure used for each stage of the forecasting pipeline. Continuous integration (denoted by the Travis icon; a woman wearing safety glasses and hardhat) triggers the code involved in events that link the stages of the pipeline, such as using the output from the forecasting stage (purple box) to create an updated website (rose box). Travis also runs tasks within a stage, such as testing code and adding weather data (icons on arrows originating and ending on the same box).

127 “continuous analysis” is making sure that the pipelines will continue to run even as  
128 software dependencies change (Beaulieu-Jones & Greene, 2017). Many of us have  
129 experienced the frustrations that can occur when software updates (e.g., changes in R  
130 package versions) create errors in previously functional code. We experienced this issue  
131 when the `tscount` package (Liboschik, Fokianos, & Fried, 2015), used by one of our  
132 forecasting models, was temporarily removed from CRAN (the R package repository)  
133 and could not be installed in the usual way. This broke our forecasting pipeline because  
134 we could no longer run models that used that package. To make our pipeline robust to  
135 changes in external software dependencies, we follow Beaulieu and Greene’s (2017)  
136 recommendation to use software containers. Software containers are standalone  
137 packages that contain copies of everything you need to run some piece of software,  
138 including the operating system. Once created, a software container is basically a time  
139 capsule, it contains all the software dependencies in the exact state used to develop and  
140 run the software. If those dependencies change (or disappear) in the wider world, they  
141 still exist, unchanged, in your container. We use an existing platform, Docker (Merkel,  
142 2014), to store an exact image of the complete software environment for running the  
143 forecasts. Docker also allows a specified set of packages to be used consistently across  
144 different computer and server environments. Using containers allows us to control  
145 transitions to new package versions, implementing them only after we have tested them  
146 and made any necessary changes to the data processing and analysis code. We use a  
147 container created by the Rocker project which is a Docker image with many important  
148 R packages (i.e. tidyverse) pre-installed (Boettiger & Eddelbuettel, 2017). We add our  
149 code and dependencies to this existing Rocker image to create a software container for  
150 our forecasting pipeline. In combination, the automated running of the pipeline  
151 (continuous integration) and the guarantee it will not stop working unexpectedly due to  
152 software dependencies (via a software container) allows continuous analysis to serve as  
153 the glue that connects all stages of the forecasting pipeline.

## 154 **Data Collection, Entry, and Processing**

155 Iterative forecasting benefits from frequently updated data so that state changes can be  
156 quickly incorporated into new forecasts (Dietze et al., 2016). Frequent data collection  
157 and rapid processing are both important for providing timely forecasts. Since we collect  
158 data monthly, ensuring that the models have access to the newest data requires a data  
159 latency period of less than 1 month from collection to availability for modeling. To  
160 accomplish this, we automated components of the data processing and quality  
161 assurance/quality control (QA/QC) process to reduce the time needed to add new data  
162 to the database (Figure 1).

163 New data are double-entered into Microsoft Excel using the “data validation” feature.  
164 The two versions are then compared using an R script to control for errors in data entry.  
165 Quality control (QC) checks using the `testthat` R package (Wickham, 2011) are run  
166 on the data to test for validity and consistency both within the new data and between the  
167 new and archived data. The local use of the QC scripts to flag problematic data greatly  
168 reduces the time spent error-checking and ensures that the quality of data is consistent.  
169 The data are then uploaded to the GitHub-based Portal Data repository  
170 (<https://github.com/weecology/PortalData>). GitHub (<https://github.com/>) is a software  
171 development tool for managing computer code development, but we have also found it  
172 useful for data management. On GitHub, changes to data can be tracked through the Git  
173 version control system which logs all changes made to any files in the repository -  
174 giving us a record of exactly of when specific lines of data were changed or added. All  
175 updates to data are processed through “pull requests”, which are notifications that  
176 someone has a modified version of the data to contribute. QA/QC checks are  
177 automatically run on the submitted data using continuous integration to ensure that  
178 these checks are run and that no avoidable errors reach the official version of the dataset.  
179 We also automated the updating of supplementary data tables, including information on  
180 weather and trapping history, that were previously updated manually. As soon as new



181 field data is merged into the repository, continuous integration updates all  
182 supplementary files. Weather data is automatically fetched from our cellular-connected  
183 weather station, cleaned, and appended to the weather data table. Supplementary data  
184 tables related to trapping history are updated based on the data added to the main data  
185 tables. Using CI for this ensures that all supplementary data tables are always  
186 up-to-date with the core data.

## 187 **Data Sharing**

188 The Portal Project has a long history of making its data publicly available, so that  
189 anyone can use it for forecasting or other projects. Historically the publication of the  
190 data was conducted through data papers (S. Ernest et al., 2009, S. M. Ernest et al.  
191 (2016)), the most common approach in ecology; however, this approach caused years of  
192 data latency. With the recent switch to posting data directly to a public GitHub  
193 repository (Figure 1) with a CC0 waiver (i.e. no restrictions on data use;  
194 <https://creativecommons.org/publicdomain/zero/1.0/>) data latency for everyone has been  
195 reduced to less than one month making meaningful iterative near-term forecasting  
196 possible for not only our group but other interested parties, as well.

## 197 **Data Manipulation**

198 Once data is available, it must be processed into a form appropriate for modeling  
199 (Figure 1). For many ecological datasets, this requires not only simple data  
200 manipulation but also a good understanding of the data to facilitate appropriate  
201 aggregation. Data manipulation steps are often conducted using custom one-off code to  
202 convert the raw data into the desired form (Morris & White, 2013), but this approach  
203 has several limitations. First, each researcher must develop and maintain their own data  
204 manipulation code, which is inefficient and can result in different researchers producing

different versions of the data for the same task. Subtle differences in data processing decisions have led to confusion when reproducing results for the Portal data in the past. Second, this kind of code is rarely robust to changes in data structure and location. Based on our experience developing and maintaining the Data Retriever (Morris & White, 2013; Senyondo et al., 2017), these kinds of changes are common. Finally, this kind of code is generally poorly tested, which can lead to errors based on mistakes in data manipulation. To avoid these issues for the Portal Project data, the Portal team has been developing an R package (portalr; <http://github.com/weecology/portalr>) for acquiring the data and handling common data cleaning and aggregation tasks. As a result, our modeling and forecasting code only needs to install this package and run the data manipulation and summary functions to get the appropriate data (Figure 1b). The package undergoes thorough automated unit testing to ensure that data manipulations are achieving the desired results. Having data manipulation code maintained in a separate package that focuses on consistently providing properly summarized forms of the most recent data has made maintaining the forecasting code itself much more straightforward.

## Modeling and Forecasting

Iterative near-term forecasting involves regularly refitting a variety of different models (Figure 1). Ideally, new models should be easy to incorporate to allow for iterative improvements to the general modeling structure and approach. We use CI to refit the models and make new forecasts each time the modeling code changes and when new data become available (Figure 1b). We use a plugin infrastructure to allow new models to be easily added to the system. This approach treats each model as an interchangeable black box - all models have access to the same input data and generate the same structure for model outputs (Figure 2). During each run of the forecasting code, all existing models are run and the standardized outputs are combined into a single file to

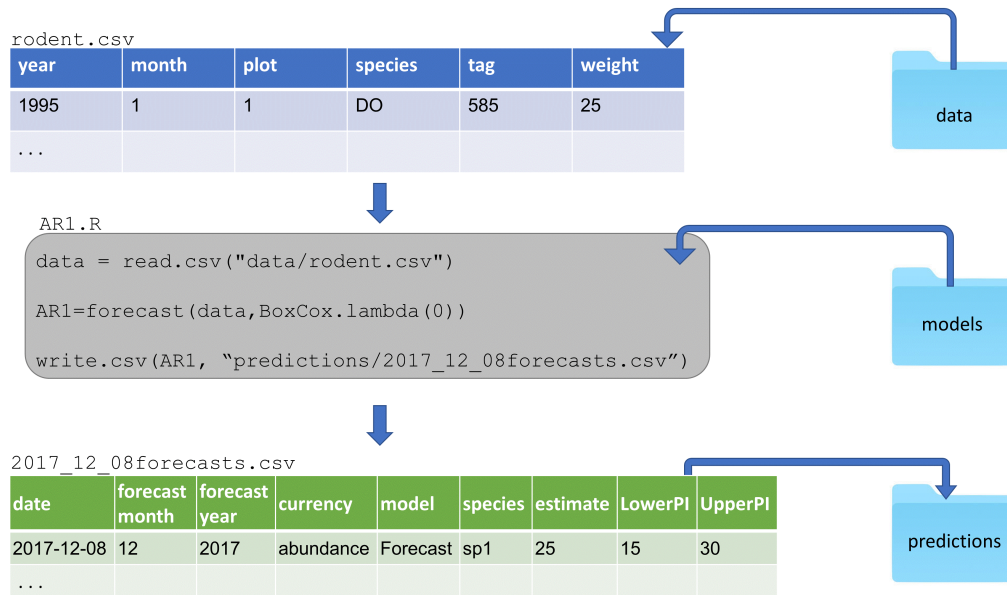


Figure 2: Figure 2. Demonstration of plugin infrastructure. All model scripts (represented here by the example AR1.R) are housed in a single folder. Each model script uses data provided by the core forecasting code (represented here by rodent.csv) and returns its forecast outputs in a predefined structure that is consistent across models (represented here by the example 2017\_12\_08forecasts.csv). Outputs from all models run on a particular date are combined into the same file (i.e. 2017\_12\_08forecasts.csv) to allow cross-model evaluations. Model output files are housed in a folder containing all forecast outputs from all previous dates to facilitate archiving and forecast assessment.

231 store the results of the different models' forecasts. A weighted ensemble model is then  
 232 added with weights based on how well individual models fit the training data. This  
 233 plugin infrastructure makes it easy to add and compare very different types of models,  
 234 from the basic time-series approaches currently implemented to the more complex  
 235 state-space and machine learning models we hope to implement in the future. As long  
 236 as a model script can load the provided data and produce the appropriate output it will  
 237 be run and its results incorporated into the rest of the forecasting system.

238 In addition to flexibility in what model structures can be supported, we also wanted to  
 239 support flexibility in what the models predict. Allowing models to make forecasts for  
 240 system properties ranging from individual species' population abundances to total  
 241 community biomass facilitates exploration of differences in forecastability across  
 242 different aspects of ecological systems. We designed a forecast output format to support

243 this. Each forecast output file contains the date being forecast, the collection date of the  
244 data used for fitting the models, the model name, the date the forecast was made, the  
245 state variable being forecast (e.g., rodent biomass, the abundance of a species), and the  
246 forecast value and associated uncertainty of that forecast (Figure 2). This allows us to  
247 store a variety of different forecasts in a common format and may serve as a useful  
248 starting point for developing a standard for storing ecological forecasts more generally.

249 Forecasts are currently evaluated using root mean square error (RMSE) to evaluate  
250 point forecasts and coverage to evaluate uncertainty. We plan to add additional metrics  
251 in the future. In addition to evaluating the actual forecasts, we also use hindcasting  
252 (forecasting on already collected data; Jolliffe & Stephenson, 2003) to gain additional  
253 insight into the methods that work best for forecasting this system. For example, a  
254 model is fit using rodent observations up to June 2005, then used to make a forecast 12  
255 months out to May 2006. The observations of that 12 month period can immediately be  
256 used to evaluate the model. Since hindcasting is conducted using data that has already  
257 been collected, it allows model comparisons to be conducted on large numbers of  
258 hindcasts and provides insight into which models make the best forecasts without  
259 needing to wait for new data to be collected (Harris, Taylor, & White, 2018). It can also  
260 be used to quickly evaluate new models instead of waiting for an adequate amount of  
261 data to accumulate.

## 262 **Archiving**

263 Publicly archiving forecasts before new data is collected allows the field to assess,  
264 compare, and build on forecasts made by different groups (McGill, 2012; Dietze et al.,  
265 2016; Tredennick et al., 2016; Harris et al., 2018) (Figure 1). Archiving serves as a  
266 form of pre-registration for model predictions, because the forecasts cannot be modified  
267 once the data to assess them has been collected. This helps facilitate an unbiased  
268 interpretation of model performance. To serve this role archives should be publicly

269 accessible and be a permanent record that cannot be changed or deleted. The second  
270 criteria means that GitHub is not sufficient for archival purposes because repositories  
271 can be changed or deleted (Bergman, 2012; White, 2015). We explored three major  
272 repositories for archiving forecasts: FigShare (<https://figshare.com/>), Zenodo  
273 (<https://zenodo.org/>), and Open Science Framework (<https://osf.io/>). While all three  
274 repositories allowed for easy manual submissions (i.e., a human uploading files after  
275 each forecast), automating this process was substantially more difficult. Various  
276 combinations of repositories, APIs (i.e., interfaces for automatically interacting with the  
277 archiving websites) and associated R packages had issues with: 1) integrating  
278 authorization with continuous integration; 2) automatically making archived files public;  
279 3) adding new files to an existing location; or 4) automatically permanently archiving  
280 the files. Our eventual solution was to leverage the GitHub-Zenodo integration  
281 (<https://guides.github.com/activities/citable-code/>) and automatically push forecasts to a  
282 GitHub repository from the CI server and release them via the GitHub API. The  
283 GitHub-Zenodo integration is designed to automatically create versioned archives of  
284 GitHub repositories. We created a repository for storing forecasts  
285 (<https://github.com/weecology/forecasts>) and linked this repository with Zenodo (a  
286 one-time manual process). Each time a new forecast is created, our pipeline adds the  
287 new forecasts to the GitHub repository and uses the GitHub API to create a new  
288 “release” for that repository. This triggers the Zenodo-GitHub integration, which  
289 automatically archives the resulting forecasts under a top-level DOI that refers to all  
290 archived forecasts (<https://doi.org/10.5281/zenodo.839580>). Through this process, we  
291 automatically archive every forecast made with a documented time-stamp. In addition,  
292 we also archive the full state of the modeling and forecasting repository  
293 (<https://doi.org/10.5281/zenodo.833438>). This ensures that every forecast is fully  
294 reproducible since the exact code used to generate every forecast is preserved. Early  
295 forecasts from this system are archived in this modeling and forecasting code archive,  
296 not the newer forecasts repository.

## 297 **Presentation**

298 Each month, we present our forecasts on a website that displays monthly rodent  
299 forecasts, model evaluation metrics, monthly reports, and information about the study  
300 site (Figure 3; <http://portal.naturecast.org>). The website includes a graphical  
301 presentation of the most recent month's forecasts (including uncertainty) and compares  
302 the latest data to the previous forecasts. Information on the species and the field site are  
303 also included. The site is built using Rmarkdown (Allaire et al., 2017), which naturally  
304 integrates into the pipeline, and is automatically updated after each forecast. The  
305 `knitr` R package (Xie, 2015) compiles the code into HTML, which is then published  
306 using Github Pages (<https://pages.github.com/>). The files for the website are stored in a  
307 subdirectory of the forecasting repository. As a result, the website is also archived  
308 automatically as part of archiving the forecast results.

## 309 **Discussion**

310 Following the recommendations of Dietze et al (2016), we developed an automated  
311 iterative forecasting system (Figure 1) to support repeated forecasting of an ecological  
312 system. Our forecasting system automatically acquires and processes the newest data,  
313 refits the models, makes new forecasts, publicly archives those forecasts, and presents  
314 both the current forecast and information on how previous forecasts performed. Every  
315 week our forecasting system generates a new set of forecasts with no human  
316 intervention, except for the entry of new field data. Our forecasting system ensures that  
317 forecasts based on the most recent data are always available. It is also designed to allow  
318 rapid assessment of the performance of multiple forecasting models for a number of  
319 different states of the system, including the abundances of individual species and  
320 community-level variables such as total abundance. To create this iterative near-term  
321 forecasting system, we used R to process data and conduct analyses, and leveraged

## Portal Forecast

### Total Abundance Forecast

This is the forecast for next month's sampling of rodents at Portal.

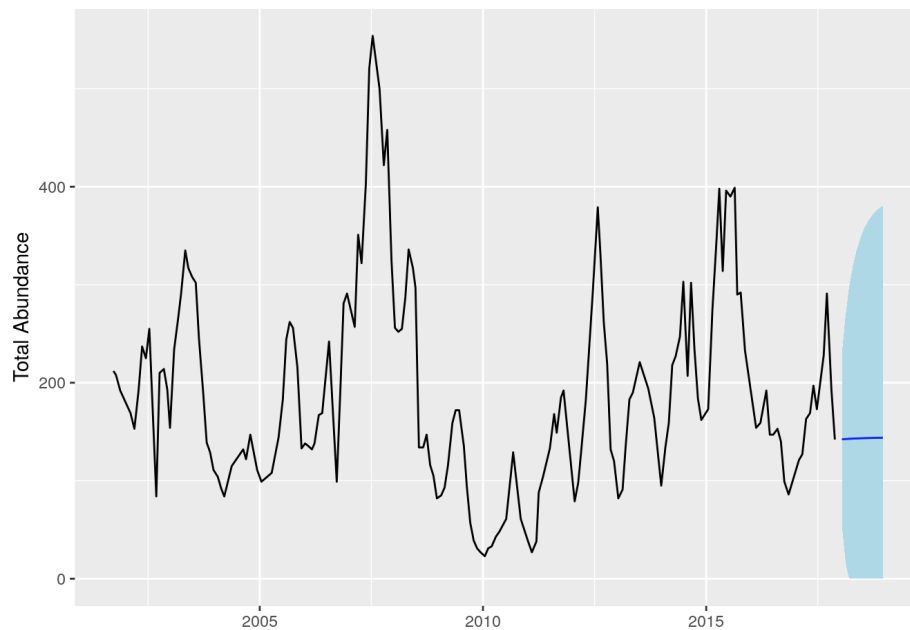


Figure 3: Figure 3. Screen capture of the homepage of the Portal Forecasting website (<http://portal.naturecast.org>). This site contains information on the most current forecasts, evaluation of forecast performance, and general information about the species being forecast.

322 existing tools and services (i.e. GitHub, Travis, Docker) for more complicated  
323 cyberinfrastructure tasks. Thus, our approach to developing iterative near-term  
324 forecasting infrastructure provides an example for how short-term ecological  
325 forecasting systems can be developed.

326 We designed this forecasting system with the goal of making it relatively easy to build,  
327 maintain, and extend. We used existing technology for both running the pipeline and  
328 building individual components, which allowed us to build the system relatively cheaply  
329 in terms of both time and money. This included the use of tools like Docker for  
330 reproducibility, Travis CI continuous integration for automatically running the pipeline,  
331 Rmarkdown and `knitr` for generating the website, and the already existing integration  
332 between Github and Zenodo to archive the forecasts. By using this “continuous analysis”  
333 approach (Beaulieu-Jones & Greene, 2017), where analyses are automatically rerun  
334 when changes are made to data, models, or associated code, we have reduced the time  
335 required by scientists to run and maintain the forecasting pipeline. To make the system  
336 extensible so that new models could be easily incorporated, we used a plugin-based  
337 infrastructure so that adding a new model to the system is as easy as adding a single file  
338 to the ‘models’ folder in our repository (Figure 2). This should substantially lower the  
339 barriers to other scientists contributing models to this forecasting effort. We also  
340 automatically archive the resulting forecasts publicly so that the performance of these  
341 forecasts can be assessed by both us and other researchers as new data is collected. This  
342 serves as a form of pre-registration by providing a quantitative record of the forecast  
343 before the data being predicted were collected.

344 While building this system was facilitated by the use of existing technological solutions,  
345 there were still a number of challenges in making existing tools work for automated  
346 iterative forecasting. Continuous integration is designed primarily for running  
347 automated tests on software, not for running a coordinated forecasting pipeline. As a  
348 result, extra effort was sometimes necessary to figure out how to get these systems to



349 work properly in non-standard situations, like running code that was not part of a  
350 software package. In addition, hosted continuous integration solutions, like Travis,  
351 provide only limited computational resources. As the number and complexity of the  
352 models we fit has grown, we have had to continually invest effort in reducing our total  
353 compute time so we can stay within these limits. Finally, we found no satisfactory  
354 existing solution for archiving our results. All approaches we tried had limitations when  
355 it came to automatically generating publicly versioned archives of forecasts on a  
356 repeated basis, and our eventual solution was difficult to configure to such a degree that  
357 it will be an impediment for most researchers. Overall, we found existing technology to be  
358 sufficient to the task of creating an iterative forecasting pipeline, but it required greater  
359 expertise and a greater investment of time than is ideal. Additional tool development to  
360 reduce the effort required for scientists to set up their own short-term forecasting  
361 systems would clearly be useful. However, our efforts show that it is possible to use  
362 existing tools to develop initial iterative systems as a method for both advancing  
363 scientific understanding and developing proof of concept forecasting systems.

364 Because of the breadth of expertise needed to set up our forecasting pipeline, our effort  
365 required a team with diverse skills and perspectives, ranging from software  
366 development to field site expertise. It is rare to find such breadth within a single  
367 research group, and our system was developed as a collaboration between the lab  
368 collecting the data and a computational ecology lab. When teams have a breadth of  
369 expertise, communication can be challenging (Winowiecki et al., 2011). We found a  
370 shared base of knowledge related to both the field research and fundamental  
371 computational skills was important for the success of the group. The two labs are part of  
372 a joint interdisciplinary ecology group that has a mission of breaking down barriers  
373 between field and computational/theoretical ecologists (<http://weecology.org>). Everyone  
374 on the team had received training in fundamental data management and computing  
375 skills through a combination of university courses, Software and Data Carpentry

workshops (Teal et al., 2015), and informal lab training efforts. In addition, everyone was broadly familiar with the study site and methods of data collection, and most team members had participated in field work at the site on multiple occasions. This provided a shared set of knowledge and vocabulary that actively facilitated interdisciplinary interactions. Given the current state of tools for forecasting, forecasting teams will need people with significant experience in working with continuous integration and APIs. This means interdisciplinary teams will generally be required for creating these pipelines until tool development improves. To improve the success of these diverse groups, we believe efforts at providing ‘team science’ training to scientists interested in forecasting will be beneficial for the success of iterative forecasting attempts for the foreseeable future (???)

We developed infrastructure for automatically making iterative forecasts with the goals of making accurate forecasts for this well-studied system, learning what methods work well for ecological forecasting more generally, and improving our understanding of the processes driving ecological dynamics. The most obvious application of automated iterative ecological forecasting is for speeding up development of forecasting models by using the most recent data available and by quickly iterating to improve the models used for forecasting. By learning what works best for forecasting in this and other ecological systems, we will better understand what the best approaches are for ecological forecasting more generally. By designing the pipeline so that it can forecast many different aspects of the ecological community, we also hope to learn about what aspects of ecology are more forecastable. Finally, automated forecasting infrastructures like this one also provide a core foundation for faster scientific inquiry because new models can quickly be applied to data and compared to existing models. The forecasting infrastructure does the time-consuming work of data processing, data integration, and model assessment, allowing new research to focus on the models being developed and the inferences about the system that can be drawn from them (Dietze et al., 2016). We

403 plan to use this pipeline to drive future research into understanding the processes that  
404 govern the dynamics of individual populations and the community as a whole. By  
405 regularly running different models for population and community dynamics, a near-term  
406 iterative pipeline such as ours should also make it possible to rapidly detect changes in  
407 how the system is operating, which should allow the rapid identification of ecological  
408 transitions or even possibly allow them to be prevented (Pace et al., 2017). By building  
409 an automated iterative near-term forecasting infrastructure we can improve our ability to  
410 forecast natural systems, understand the biology driving ecological dynamics, and  
411 detect or even predict changes in system state that are important for conservation and  
412 management.

## 413 **Acknowledgements**

414 We thank Henry Senyondo for help with continuous integration and Hao Ye and Juniper  
415 Simonis for discussions and feedback on the manuscript. We thank all of the graduate  
416 students, postdocs, and volunteers who have collected the Portal Project over the last 40  
417 years and the developers of the software and tools that made this project possible. This  
418 research was supported by the National Science Foundation through grant 1622425 to  
419 S.K.M. Ernest and by the Gordon and Betty Moore Foundation's Data-Driven  
420 Discovery Initiative through grant GBMF4563 to E.P. White.

## 421 **Box 1. Key practices for automated iterative near-term** 422 **ecological forecasting**

423 A list of some of the key practices developed by Dietze et al (2016) for facilitating  
424 iterative near-term ecological forecasting and discussion of why these practices are  
425 important.

## 426 **Data**

### 427 **1. Frequent data collection**

428 Frequent data collection allows models to be regularly updated and forecasts to be  
429 frequently evaluated (Dietze et al., 2016). Depending on the system being studied, this  
430 frequency could range from sub-daily to annual, but typically the more frequently the  
431 data is collected the better.

### 432 **2. Rapid data release under open licenses**

433 Data should be released as quickly as possible (low latency) under open licenses so that  
434 forecasts can be made frequently and data can be accessed by a community of  
435 forecasters (Dietze et al., 2016; Vargas et al., 2017).

### 436 **3. Best practices in data structure**

437 To reduce the time and effort needed to incorporate data into models, best practices in  
438 data structure should be employed for managing and storing collected data to ensure it  
439 is easy to integrate into other systems (interoperability) (Borer et al., 2009; Strasser et  
440 al., 2011; White et al., 2013).

## 441 **Models**

### 442 **4. Focus on uncertainty**

443 Understanding the uncertainty of forecasts is crucial to interpreting and understanding  
444 their utility. Models used for forecasting should be probabilistic to properly quantify  
445 uncertainty and to convey how this uncertainty increases through time. Evaluation of  
446 forecast models should include assessment of how accurately they quantify uncertainty  
447 as well as point estimates (Hooten & Hobbs, 2015).

### 448 **5. Compare forecasts to simple baselines**

449 Understanding how much information is present in a forecast requires comparing its  
450 accuracy to simple baselines to see if the models yield improvements over the naive  
451 expectation that the system is static (Harris et al., 2018).

## 452 **6. Compare and combine multiple modeling approaches**

453 To quickly learn about the best approaches to forecasting different aspects of ecology,  
454 multiple modeling approaches should be compared (Harris et al., 2018). Different  
455 modeling approaches should also be combined into ensemble models, which often  
456 outperform single models for prediction (Weigel, Liniger, & Appenzeller, 2008).

## 457 **Cyberinfrastructure**

458 In addition to improvements in data and models, iterative near-term forecasting requires  
459 improved infrastructure and approaches to support continuous model development and  
460 iterative forecasting (Dietze et al., 2016).

## 461 **7. Best practices in software development**

462 Best practices should be followed in the development of scientific software and  
463 modeling to make it easier to maintain, integrate into pipelines, and build on by other  
464 researchers. Key best practices include open licenses, good documentation, version  
465 control, and cross-platform support (Wilson et al., 2014; Hampton et al., 2015).

## 466 **8. Support easy inclusion of new models**

467 To facilitate the comparison and ensembling of different modeling approaches, code for  
468 fitting models and making forecasts should be easily extensible, to allow models  
469 developed by different groups to be integrated into a single framework (Dietze et al.,  
470 2016).

## 471 **9. Automated end-to-end reproducibility**

Each forecast iteration involves acquiring new data, refitting the models, and making new forecasts. This should be done automatically without requiring human intervention. Therefore, the process of making forecasts should emphasize end-to-end reproducibility, including data, models, and evaluation (Stodden & Miguez, 2014), to allow the forecasts to be easily rerun as new data becomes available (Dietze et al., 2016).

## 10. Publicly archive forecasts

Forecasts should be openly archived to demonstrate that the forecasts were made without knowledge of the outcomes and to allow the community to assess and compare the performance of different forecasting approaches both now and in the future (McGill, 2012; Dietze et al., 2016; Tredennick et al., 2016; Harris et al., 2018). Ideally, the forecasts and evaluation of their performance should be automatically posted publicly in a manner that is understandable by both scientists and the broader stakeholder community.

## References

- Allaire, J., Cheng, J., Xie, Y., McPherson, J., Chang, W., Allen, J., . . . Arslan, R. (2017). *Rmarkdown: Dynamic documents for r*. Retrieved from <https://CRAN.R-project.org/package=rmarkdown>
- Araujo, M. B., & New, M. (2007). Ensemble forecasts of species distributions. *Trends in Ecology and Evolution*, 22(1), 42–47.
- Beaulieu-Jones, B. K., & Greene, C. S. (2017). Reproducibility of computational workflows is automated using continuous analysis. *Nature Biotechnology*, 35(4), 342–346.
- Bergman, C. (2012). On the preservation of published bioinformatics code on github. Retrieved from <https://caseybergman.wordpress.com/2012/11/08/>

on-the-preservation-of-published-bioinformatics-code-on-github/

Boettiger, C., & Eddelbuettel, D. (2017). An introduction to rocker: Docker containers for R. *arXiv Preprint arXiv:1710.03675*.

Borer, E. T., Seabloom, E. W., Jones, M. B., & Schildhauer, M. (2009). Some simple guidelines for effective data management. *The Bulletin of the Ecological Society of America*, 90(2), 205–214.

Brown, J. (1998). The desert granivory experiments at Portal. *Experimental Ecology*. Oxford University Press, Oxford, UK, 71–95.

Clark, J. S., Carpenter, S. R., Barber, M., Collins, S., Dobson, A., Foley, J. A., . . . others. (2001). Ecological forecasts: An emerging imperative. *Science*, 293(5530), 657–660.

Dietze, M. C. (2017). *Ecological forecasting*. Princeton University Press.

Dietze, M. C., Fox, A., Betancourt, J. L., Hooten, M. B., Jarnevich, C. S., Keitt, T. H., . . . others. (2016). Iterative ecological forecasting: Needs, opportunities, and challenges. *Proceedings of the National Academy of Sciences*.

Diniz-Filho, J. A. F., Bini, L. M., Rangel, T. F., Loyola, R. D., Hof, C., Noguees-Bravo, D., & Araujo, M. B. (2009). Partitioning and mapping uncertainties in ensembles of forecasts of species turnover under climate change. *Ecography*, 32(6), 897–906.

Díaz, S., Demissew, S., Carabias, J., Joly, C., Lonsdale, M., Ash, N., . . . others. (2015). The ipbes conceptual framework—connecting nature and people. *Current Opinion in Environmental Sustainability*, 14, 1–16.

Ernest, S. M., Brown, J. H., Thibault, K. M., White, E. P., & Goheen, J. R. (2008). Zero sum, the niche, and metacommunities: Long-term dynamics of community assembly. *The American Naturalist*, 172(6), E257–E269.

Ernest, S. M., Yenni, G. M., Allington, G., Christensen, E. M., Geluso, K., Goheen, J. R., . . . others. (2016). Long-term monitoring and experimental manipulation of a

521 chihuahuan desert ecosystem near portal, arizona (1977–2013). *Ecology*, 97(4),  
522 1082–1082.

523 Ernest, S., Valone, T. J., & Brown, J. H. (2009). Long-term monitoring and  
524 experimental manipulation of a chihuahuan desert ecosystem near portal, arizona, usa.  
525 *Ecology*, 90(6), 1708–1708.

526 Hampton, S. E., Anderson, S. S., Bagby, S. C., Gries, C., Han, X., Hart, E. M., ...  
527 others. (2015). The tao of open science for ecology. *Ecosphere*, 6(7), 1–13.

528 Harris, D. J., Taylor, S. D., & White, E. P. (2018). Forecasting biodiversity in breeding  
529 birds using best practices. *PeerJ*.

530 Hooten, M. B., & Hobbs, N. (2015). A guide to bayesian model selection for ecologists.  
531 *Ecological Monographs*, 85(1), 3–28.

532 Jolliffe, I. T., & Stephenson, D. B. (Eds.). (2003). *Forecast verification: a practitioner's*  
533 *guide in atmospheric science*. John Wiley; Sons, Ltd. Retrieved from  
534 <http://linkinghub.elsevier.com/retrieve/pii/S0169207005001214>

535 Liboschik, T., Fokianos, K., & Fried, R. (2015). *Tscount: An r package for analysis of*  
536 *count time series following generalized linear models*. Universitätsbibliothek  
537 Dortmund.

538 Luo, Y., Ogle, K., Tucker, C., Fei, S., Gao, C., LaDeau, S., ... Schimel, D. S. (2011).  
539 Ecological forecasting and data assimilation in a data-rich era. *Ecological Applications*,  
540 21(5), 1429–1442.

541 McGill, B. J. (2012). Ecologists need to do a better job of prediction – part ii – partly  
542 cloudy and a 20% chance of extinction (or the 6 p's of good prediction). Retrieved from  
543 [https://dynamicecology.wordpress.com/2013/01/09/](https://dynamicecology.wordpress.com/2013/01/09/ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/)  
544 [ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/](https://dynamicecology.wordpress.com/2013/01/09/ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/)

545 Merkel, D. (2014). Docker: Lightweight linux containers for consistent development



546 and deployment. *Linux J.*, 2014(239). Retrieved from  
 547 <http://dl.acm.org/citation.cfm?id=2600239.2600241>

548 Morris, B. D., & White, E. P. (2013). The ecodata retriever: Improving access to  
 549 existing ecological data. *PloS One*, 8(6), e65848.

550 Pace, M. L., Batt, R. D., Buelo, C. D., Carpenter, S. R., Cole, J. J., Kurtzweil, J. T., &  
 551 Wilkinson, G. M. (2017). Reversal of a cyanobacterial bloom in response to early  
 552 warnings. *Proceedings of the National Academy of Sciences*, 114(2), 352–357.

553 Petchey, O. L., Pontarp, M., Massie, T. M., Kéfi, S., Ozgul, A., Weilenmann, M., ...  
 554 others. (2015). The ecological forecast horizon, and examples of its uses and  
 555 determinants. *Ecology Letters*, 18(7), 597–611.

556 Senyondo, H., Morris, B. D., Goel, A., Zhang, A., Narasimha, A., Negi, S., ... White,  
 557 E. P. (2017). Retriever: Data retrieval tool. *The Journal of Open Source Software*, 2(19),  
 558 451. doi:10.21105/joss.00451

559 Stodden, V., & Miguez, S. (2014). Best practices for computational science: Software  
 560 infrastructure and environments for reproducible and extensible research. *Journal of*  
 561 *Open Research Software*, 2(1).

562 Strasser, C., Cook, R., Michener, W., Budden, A., & Koskela, R. (2011). Promoting data  
 563 stewardship through best practices. In *Proceedings of the environmental information*  
 564 *management conference 2011 (eim 2011)*. Oak Ridge National Laboratory (ORNL).

565 Tallis, H. M., & Kareiva, P. (2006). Shaping global environmental decisions using  
 566 socio-ecological models. *Trends in Ecology & Evolution*, 21(10), 562–568.

567 Teal, T. K., Cranston, K. A., Lapp, H., White, E., Wilson, G., Ram, K., & Pawlik, A.  
 568 (2015). Data carpentry: Workshops to increase data literacy for researchers.  
 569 *International Journal of Digital Curation*, 10(1), 135–143.

570 Tredennick, A. T., Hooten, M. B., Aldridge, C. L., Homer, C. G., Kleinhesselink, A. R.,

571 & Adler, P. B. (2016). Forecasting climate change impacts on plant populations over  
572 large spatial extents. *Ecosphere*, 7(10).

573 Vargas, R., Alcaraz-Segura, D., Birdsey, R., Brunsell, N. A., Cruz-Gaistardo, C. O.,  
574 Jong, B. de, ... others. (2017). Enhancing interoperability to facilitate implementation  
575 of redd+: Case study of mexico. *Carbon Management*, 8(1), 57–65.

576 Weigel, A. P., Liniger, M., & Appenzeller, C. (2008). Can multi-model combination  
577 really enhance the prediction skill of probabilistic ensemble forecasts? *Quarterly*  
578 *Journal of the Royal Meteorological Society*, 134(630), 241–260.

579 White, E. P. (2015). Some thoughts on best publishing practices for scientific software.  
580 *Ideas in Ecology and Evolution*, 8(1).

581 White, E. P., Baldridge, E., Brym, Z. T., Locey, K. J., McGlinn, D. J., & Supp, S. R.  
582 (2013). Nine simple ways to make it easier to (re) use your data. *Ideas in Ecology and*  
583 *Evolution*, 6(2).

584 Wickham, H. (2011). Testthat: Get started with testing. *The R Journal*, 3, 5–10.  
585 Retrieved from  
586 [http://journal.r-project.org/archive/2011-1/RJournal\\_2011-1\\_Wickham.pdf](http://journal.r-project.org/archive/2011-1/RJournal_2011-1_Wickham.pdf)

587 Wilson, G., Aruliah, D. A., Brown, C. T., Hong, N. P. C., Davis, M., Guy, R. T., ...  
588 others. (2014). Best practices for scientific computing. *PLoS Biology*, 12(1), e1001745.

589 Winowiecki, L., Smukler, S., Shirley, K., Remans, R., Peltier, G., Lothes, E., ...  
590 Alkema, L. (2011). Tools for enhancing interdisciplinary communication.  
591 *Sustainability: Science, Practice and Policy*, 7(1), 74–80.

592 Xie, Y. (2015). *Dynamic documents with r and knitr* (Vol. 29). CRC Press.