

Automated iterative near-term forecasting for the Portal Project

*Ethan P. White^{1,2,3}, Glenda M. Yenni¹, Shawn D. Taylor⁴, Erica M.
Christensen¹, Ellen K. Bledsoe⁴, Juniper L. Simonis¹, S. K. Morgan
Ernest^{1,3}*

¹ Department of Wildlife Ecology and Conservation, University of Florida, Gainesville,
FL, United States

² Informatics Institute, University of Florida, Gainesville, FL, United States

³ Biodiversity Institute, University of Florida, Gainesville, FL, United States

⁴ School of Natural Resources and Environment, University of Florida Gainesville, FL,
United States

Abstract

1. Most forecasts for the future state of ecological systems are conducted once and never updated or assessed. As a result, many available ecological forecasts are not based on the most up-to-date data, and the scientific progress of ecological forecasting models is slowed by a lack of feedback on how well the forecasts perform.
2. Iterative near-term ecological forecasting involves repeated daily to annual scale forecasts of an ecological system as new data becomes available and regular assessment of the resulting forecasts. We demonstrate how automated iterative

near-term forecasting systems for ecology can be constructed by building one to conduct monthly forecasts of rodent abundances at the Portal Project, a long-term study with over 40 years of monthly data. This system automates most aspects of the six stages of converting raw data into new forecasts: data collection, data sharing, data manipulation, modeling and forecasting, archiving, and presentation of the forecasts.

3. The forecasting system uses R code for working with data, fitting models, making forecasts, and archiving and presenting these forecasts. The resulting pipeline is automated using continuous integration (a software development tool) to run the entire pipeline once a week. The cyberinfrastructure is designed for long-term maintainability and to allow the easy addition of new models. Constructing this forecasting system required a team with expertise ranging from field site experience to software development.

4. Automated near-term iterative forecasting systems will allow the science of ecological forecasting to advance more rapidly and provide the most up-to-date forecasts possible for conservation and management. These forecasting systems will also accelerate basic science by allowing new models of natural systems to be quickly implemented and compared to existing models. Using existing technology, and teams with diverse skill sets, it is possible for ecologists to build these systems and use them to advance our understanding of natural systems.

Key-words: forecasting, prediction, mammals, iterative forecasting, Portal Project

Introduction

Forecasting the future state of ecological systems is important for management, conservation, and evaluation of how well models capture the processes governing ecological systems (Clark et al., 2001; Tallis & Kareiva, 2006; Díaz et al., 2015; Dietze,

2017). In 2001, Clark et al. (2001) called for a more central role of forecasting in ecology. Since then, an increasing number of ecological forecasts are being published. Most of these forecasts, however, are made once, published, and never assessed or updated. This lack of both regular assessment and active updating has limited the progress of ecological forecasting and hindered our ability to make useful and reliable predictions. The lack of active assessment results in limited information on how much confidence to place in forecasts and makes it difficult to determine on which forecasting methods to build. Without regular updates, forecasts lack the most current data, and the longer a forecast remains out of date, the less accurate it becomes (Petchey et al., 2015; Dietze et al., 2016). More regular updating and assessment will advance ecological forecasting as a field by accelerating the identification of the best models for individual forecasts and improving our understanding of how to best design forecasting approaches for ecology in general. For ecological forecasting to mature as a field, we need to change how we produce and interact with forecasts, creating a more dynamic interplay between model development, prediction generation, and incorporation of new data and information (Dietze et al., 2016).

With the goal of making ecological forecasting more dynamic and responsive, Dietze et al (2016) recently called for an increase in iterative near-term forecasting. Iterative near-term forecasting is defined as making predictions for the near future and repeatedly updating those predictions through a cycle of evaluation, integration of new data, and generation of new forecasts. Because forecasts are made ‘near-term’—daily to annual time scales instead of multi-decadal—predictions can be assessed more quickly and frequently, leading to more rapid model improvements (Dietze et al., 2016; Tredennick et al., 2016). Since forecasts are made repeatedly through time, new data can be continuously integrated with each iteration (Dietze et al., 2016). By quickly identifying how models are failing, facilitating rapid testing of improved models, and incorporating the most up-to-date data available, iterative near-term forecasting has the potential to

73 promote rapid improvement in the state of ecological forecasting. In addition to
74 yielding improved information for guiding policy and management (Clark et al., 2001;
75 Luo et al., 2011; Petchey et al., 2015), this iterative approach will help improve our
76 basic understanding of ecological systems (Dietze et al., 2016). For example, alternative
77 mechanistic models can be compared to determine which model provides the best
78 forecasts, thus providing insights into the importance of different ecological processes
79 (Dietze et al., 2016). Iterative near-term forecasting provides the more dynamic
80 interplay between models, predictions, and data that has been identified as necessary for
81 improving ecological forecasting and our understanding of ecological systems more
82 broadly.

83 Because iterative near-term forecasting requires a dynamic integration of models,
84 predictions, and data, Dietze et al (2016) highlight approaches to data management,
85 model construction and evaluation, and cyberinfrastructure that are necessary to
86 effectively implement this type of forecasting (Box 1). Data needs to be released
87 quickly under open licenses (Dietze et al., 2016; Vargas et al., 2017) and structured so
88 that it can be used easily by a variety of researchers and in multiple modeling
89 approaches (Borer, Seabloom, Jones, & Schildhauer, 2009; Strasser, Cook, Michener,
90 Budden, & Koskela, 2011). Models need to be able to deal with uncertainty, in both the
91 predictors and the predictions, to properly convey uncertainty in the resulting forecasts
92 (Diniz-Filho et al., 2009). Multiple models should be developed, both to assess which
93 models are performing best (Dietze et al., 2016) and to facilitate combining models to
94 form ensemble predictions which tend to perform better than single models (Araujo &
95 New, 2007; Diniz-Filho et al., 2009). Ensuring that data and models are regularly
96 updated and new forecasts are made requires cyberinfrastructure to automate data
97 processing, model fitting, prediction, model evaluation, forecast visualization, and
98 archiving. In combination, these approaches should allow forecasts to be easily rerun
99 and evaluated as new data becomes available (Box 1; Dietze et al., 2016).

100 While iterative near-term forecasting is an important next step in the evolution of
101 ecological forecasting, the requirements outlined by Dietze et al (Box 1) are not trivial
102 to implement, and few of their recommendations are in widespread use in ecology today.
103 We explored what it would entail to operationalize Dietze et al's recommendations by
104 constructing our own iterative near-term forecasting pipeline for an on-going, long-term
105 ecological study that collects high-frequency data on desert rodent abundances (J.
106 Brown, 1998; S. M. Ernest, Brown, Thibault, White, & Goheen, 2008). We constructed
107 an automated forecasting pipeline with the goal of being able to forecast rodent
108 abundances and evaluate our predictions on a monthly basis. In this paper, we discuss
109 our approach for creating this iterative near-term forecasting pipeline, the challenges we
110 encountered, the tools we used, and the lessons we learned so that others can create
111 their own iterative forecasting systems.

112 **System Background**

113 Iterative forecasting is most effective with frequently collected data, since it provides
114 more opportunities for updating model results and assessing (and potentially improving)
115 model performance (Box 1; Dietze et al., 2016). The Portal Project is a long-term
116 ecological study situated in the Chihuahuan Desert (2 km north and 6.5 km east of
117 Portal, Arizona, US). Researchers have been continuously collecting data at the site
118 since 1977, including data on the abundance of rodent and plant species (monthly and
119 twice yearly, respectively) and climatic factors such as air temperature and precipitation
120 (daily) (J. Brown, 1998; S. Ernest, Valone, & Brown, 2009; S. M. Ernest et al., 2016).
121 The site consists of 24 50m x 50m experimental plots. Each plot contains 49
122 permanently marked trapping stations laid out in a 7 x 7 grid, and all plots are trapped
123 with Sherman live traps for one night each month. For all rodents caught during a
124 trapping session, information on species identity, size, and reproductive condition is

125 collected, and new individuals are given identification tags. This information on rodent
126 populations is high-frequency, uses consistent trapping methodology, and has an
127 extended time-series (470 monthly samples and counting), making this study an ideal
128 case for near-term iterative forecasting.

129 **Implementing an automated iterative forecasting system**

130 Implementation of iterative forecasting requires the regular rebuilding of models with
131 new raw data as it becomes available and the presentation of those forecasts in usable
132 forms; in our case, this occurs monthly. Rebuilding models in an efficient and
133 maintainable way relies on developing an automated pipeline to handle the six stages of
134 converting raw data into new forecasts: data collection, data sharing, data manipulation,
135 modeling and forecasting, archiving, and presentation of the forecasts (Figure 1a). To
136 implement the pipeline outlined in Figure 1a, we used a “continuous analysis”
137 framework (*sensu* Beaulieu-Jones & Greene, 2017) that automatically processes the
138 most up-to-date data, refits the models, makes new forecasts, archives the forecasts, and
139 updates a website with analysis of current and previous forecasts. In this section we
140 describe our approach to streamlining and automating the multiple components of the
141 forecasting pipeline and the tools and infrastructure we employed to execute each
142 component.

143 **Continuous Analysis Framework**

144 A core aspect of iterative near-term forecasting is the regular rerunning of the
145 forecasting pipeline. We employed “continuous analysis” (*sensu* Beaulieu-Jones &
146 Greene, 2017) to drive the automation of both the full pipeline and a number of its
147 individual components. Continuous analysis uses a set of tools originally designed for
148 software development called “continuous integration” (CI). CI combines computing

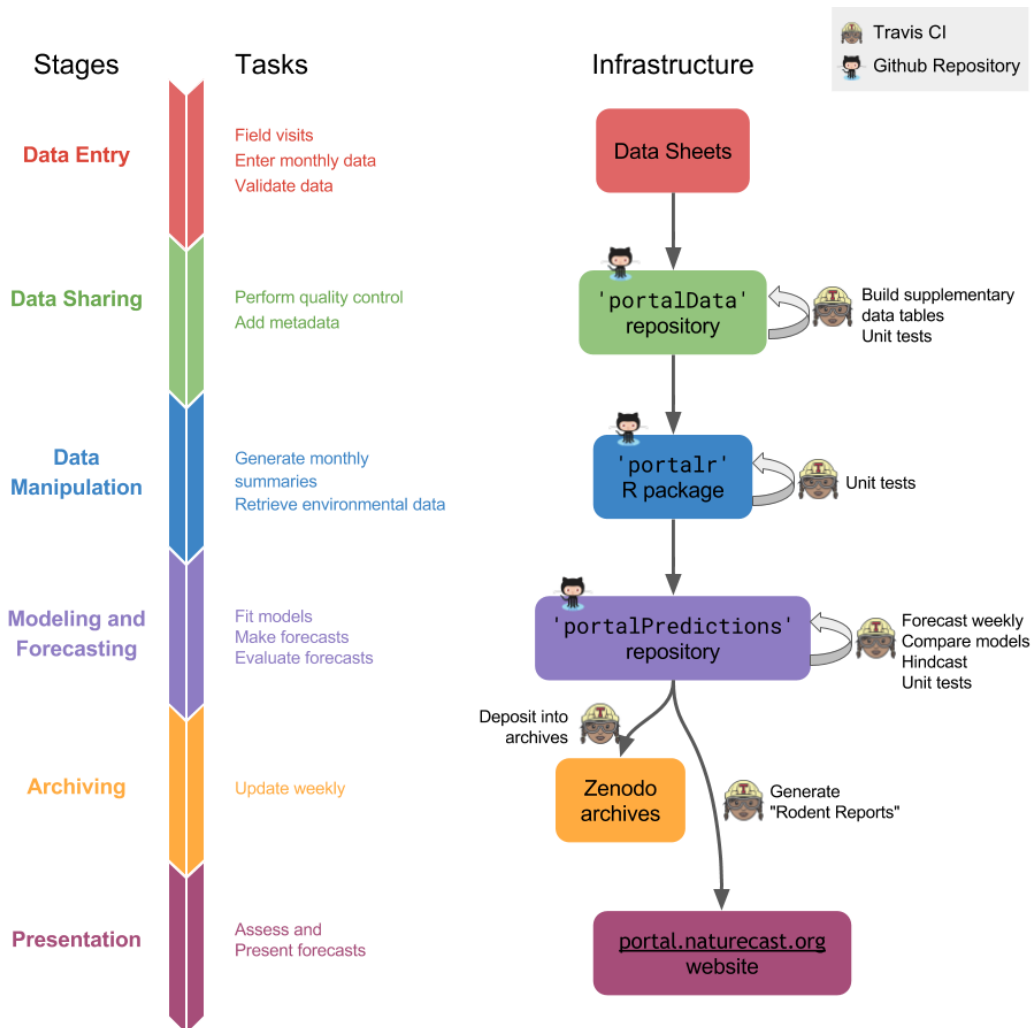


Figure 1: Figure 1. a) Stages of the forecasting pipeline. To go from raw data to forecast presentation involves a number of stages, each of which requires unique tasks, tools and infrastructure. The stages are interdependent, with outputs from one stage forming the inputs for the subsequent stage. Tasks in all stages are run using code written in R. b) Continuous integration system. Each box denotes the core infrastructure used for each stage of the forecasting pipeline. Continuous integration (denoted by the Travis icon, a woman wearing safety glasses and hardhat) triggers the code involved in events that link the stages of the pipeline, such as using the output from the forecasting stage (purple box) to create an updated website (rose box). Travis also runs tasks within a stage, such as testing code and adding weather data (icons on arrows originating and ending on the same box).

149 environments for running code with monitoring systems to identify changes in data or
150 code. Essentially, CI is a computer helper who watches the pipeline and, when it sees a
151 change in the code or data, runs all the computer scripts needed to ensure that the
152 forecasting pipeline runs from beginning to end. This is useful for iterative near-term
153 forecasting because it does not rely on humans to create new forecasts whenever new
154 models or data are added. These tools are common in the area of software development,
155 where they are used to automate software testing and integrate work by multiple
156 developers working on the same code base. However, these tools can be used for any
157 computational task that needs to be regularly repeated or run after changes to code or
158 data (Beaulieu-Jones & Greene, 2017). Our forecasting pipeline currently runs on a
159 publicly available continuous integration service (Travis CI; <https://travis-ci.org/>) that is
160 free for open source projects (up to a limited amount of computing time). Because of the
161 widespread use of CI in software development, alternative services that can run code on
162 local or cloud-based computational infrastructure also exist (Beaulieu-Jones & Greene,
163 2017). We use CI to quality check data, test code using “unit tests” (Wilson et al., 2014),
164 build models, make forecasts, and publicly present and archive the results (Figure 1b).

165 In addition to automatically running software pipelines, the other key component of
166 “continuous analysis” is making sure that the pipelines will continue to run even as
167 software dependencies change (Beaulieu-Jones & Greene, 2017). Many of us have
168 experienced the frustrations that can occur when software updates (e.g., changes in R
169 package versions) create errors in previously functional code. We experienced this issue
170 when the `tscount` package (Liboschik, Fokianos, & Fried, 2015), used by one of our
171 forecasting models, was temporarily removed from CRAN (the R package repository)
172 and could not be installed in the usual way. This broke our forecasting pipeline, as we
173 could no longer run models that used that package. To make our pipeline robust to
174 changes in external software dependencies, we follow Beaulieu and Greene’s (2017)
175 recommendation to use software containers. Software containers are standalone

176 packages that contain copies of everything needed to run a given piece of software,
177 including the operating system. Once created, a software container is basically a time
178 capsule, containing all the software dependencies in the exact state used to develop and
179 run the software. If those dependencies change (or disappear) in the wider world, they
180 still exist, unchanged, in the container. We use an existing platform, Docker (Merkel,
181 2014), to store an exact image of the complete software environment for running the
182 forecasts. Docker also allows a specified set of packages to be used consistently across
183 different computer and server environments. Using containers allows us to control
184 transitions to new package versions, implementing them only after we have tested them
185 and made any necessary changes to the data processing and analysis code. We use a
186 container created by the Rocker project, which is a Docker image with many important
187 R packages (i.e. tidyverse) pre-installed (Boettiger & Eddelbuettel, 2017). We add our
188 code and dependencies to this existing Rocker image to create a software container for
189 our forecasting pipeline. In combination, the automated running of the pipeline
190 (continuous integration) and the guarantee it will not stop working unexpectedly due to
191 software dependencies (via a software container) allows continuous analysis to serve as
192 the glue that connects all stages of the forecasting pipeline.

193 **Data Collection, Entry, and Processing**

194 Iterative forecasting benefits from frequently updated data so that state changes can be
195 quickly incorporated into new forecasts (Dietze et al., 2016). Both frequent data
196 collection and rapid processing are important for providing timely forecasts. Since we
197 collect data monthly, ensuring that the models have access to the newest data requires a
198 data latency period of less than 1 month from collection to availability for modeling. To
199 accomplish this, we automated components of the data processing and quality
200 assurance/quality control (QA/QC) process to reduce the time needed to add new data
201 to the database (Figure 1).

202 New data are double-entered into Microsoft Excel using the “data validation” feature.
203 The two versions are then compared using an R script to control for errors in data entry.
204 Quality control (QC) checks using the `testthat` R package (Wickham, 2011) are run
205 on the data to test for validity and consistency both within the new data and between the
206 new and archived data. The local use of the QC scripts to flag problematic data greatly
207 reduces the time spent error-checking and ensures that the quality of data is consistent.
208 The cleaned data are then uploaded to the GitHub-based PortalData repository
209 (<https://github.com/weecology/PortalData>). GitHub (<https://github.com/>) is a software
210 development tool for managing computer code development, but we have also found it
211 useful for data management. On GitHub, changes to data can be tracked through the Git
212 version control system which logs all changes made to any files in the repository, giving
213 us a record of exactly of when specific lines of data were changed or added. All updates
214 to data are processed through “pull requests,” which are notifications that someone has a
215 modified version of the data to contribute. QA/QC checks are automatically run on the
216 submitted data using continuous integration to ensure that no avoidable errors reach the
217 official version of the dataset.

218 We also automated the updating of supplementary data tables, including information on
219 weather and trapping history, that were previously updated manually. As soon as new
220 field data is merged into the repository, continuous integration updates all
221 supplementary files. Weather data is automatically fetched from our cellular-connected
222 weather station, cleaned, and appended to the weather data table. Supplementary data
223 tables related to trapping history are updated based on the data added to the main data
224 tables. Using CI for this ensures that all supplementary data tables are always
225 up-to-date with the core data.

226 **Data Sharing**

227 The Portal Project has a long history of making its data publicly available so that anyone
228 can use it for forecasting or other projects. Historically, the publication of the data was
229 conducted through data papers (S. Ernest et al., 2009, S. M. Ernest et al. (2016)), the
230 most common approach in ecology; this approach, however, caused years of data
231 latency. With the recent switch to posting data directly to a public GitHub repository
232 (Figure 1) with a CC0 waiver (i.e. no restrictions on data use;
233 <https://creativecommons.org/publicdomain/zero/1.0/>), data latency for everyone has
234 been reduced to less than one month, making meaningful iterative near-term forecasting
235 possible for not only our group but other interested parties, as well.

236 **Data Manipulation**

237 Once data is available, it must be processed into a form appropriate for modeling
238 (Figure 1). For many ecological datasets, this requires not only simple data
239 manipulation but also a good understanding of the data to facilitate appropriate
240 aggregation. Data manipulation steps are often conducted using custom one-off code to
241 convert the raw data into the desired form (Morris & White, 2013), but this approach
242 has several limitations. First, each researcher must develop and maintain their own data
243 manipulation code, which is inefficient and can result in different researchers producing
244 different versions of the data for the same task. Subtle differences in data processing
245 decisions have led to confusion when reproducing results for the Portal data in the past.
246 Second, this kind of code is rarely robust to changes in data structure and location.
247 Based on our experience developing and maintaining the Data Retriever (Morris &
248 White, 2013; Senyondo et al., 2017), these kinds of changes are common. Finally, this
249 kind of code is generally poorly tested, which can lead to errors based on mistakes in
250 data manipulation. To avoid these issues for the Portal Project data, the Portal team has

251 been developing an R package (portalr; <http://github.com/weecology/portalr>) for
252 acquiring the data and handling common data cleaning and aggregation tasks. As a
253 result, our modeling and forecasting code only needs to install this package and run the
254 data manipulation and summary functions to get the appropriate data (Figure 1b). The
255 package undergoes thorough automated unit testing to ensure that data manipulations
256 are achieving the desired results. Having data manipulation code maintained in a
257 separate package that focuses on consistently providing properly summarized forms of
258 the most recent data has made maintaining the forecasting code itself much more
259 straightforward.

260 **Modeling and Forecasting**

261 Iterative near-term forecasting involves regularly refitting a variety of different models
262 (Figure 1). Ideally, new models should be easy to incorporate to allow for iterative
263 improvements to the general modeling structure and approach. We use CI to refit the
264 models and make new forecasts each time the modeling code changes and when new
265 data become available (Figure 1b). We use a plugin infrastructure to allow new models
266 to be easily added to the system. This approach treats each model as an interchangeable
267 black box; all models have access to the same input data and generate the same structure
268 for model outputs (Figure 2). During each run of the forecasting code, all existing
269 models are run and the standardized outputs are combined into a single file to store the
270 results of the different models' forecasts. A weighted ensemble model is then added
271 with weights based on how well individual models fit the training data. This plugin
272 infrastructure makes it easy to add and compare very different types of models, from the
273 basic time-series approaches currently implemented to the more complex state-space
274 and machine learning models we hope to implement in the future. As long as a model
275 script can load the provided data and produce the appropriate output, it will be run and
276 its results incorporated into the rest of the forecasting system.

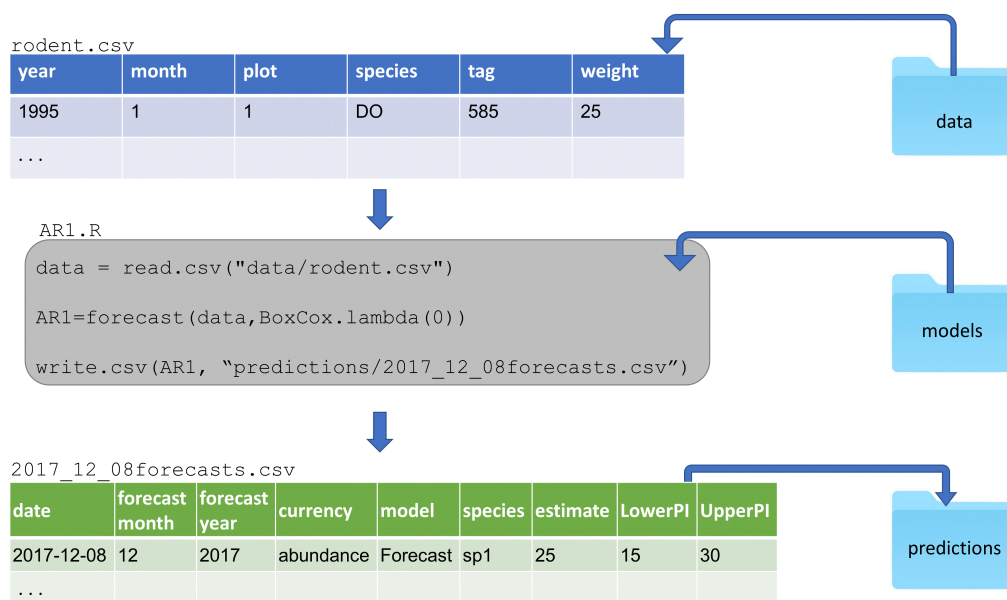


Figure 2: Figure 2. Demonstration of plugin infrastructure. All model scripts (represented here by the example AR1.R) are housed in a single folder. Each model script uses data provided by the core forecasting code (represented here by rodent.csv) and returns its forecast outputs in a predefined structure that is consistent across models (represented here by the example 2017_12_08forecasts.csv). Outputs from all models run on a particular date are combined into the same file (i.e. 2017_12_08forecasts.csv) to allow cross-model evaluations. Model output files are housed in a folder containing all forecast outputs from all previous dates to facilitate archiving and forecast assessment.

277 In addition to flexibility in what model structures can be supported, we also wanted to
278 support flexibility in what the models predict. Allowing models to make forecasts for
279 system properties ranging from individual species' population abundances to total
280 community biomass facilitates exploration of differences in forecastability across
281 different aspects of ecological systems. We designed a forecast output format to support
282 this. Each forecast output file contains the date being forecast, the collection date of the
283 data used for fitting the models, the model name, the date the forecast was made, the
284 state variable being forecast (e.g., rodent biomass, the abundance of a species), and the
285 forecast value and associated uncertainty of that forecast (Figure 2). This allows us to
286 store a variety of different forecasts in a common format and may serve as a useful
287 starting point for developing a standard for storing ecological forecasts more generally.

288 Forecasts are currently evaluated using root mean square error (RMSE) to evaluate
289 point forecasts and coverage to evaluate uncertainty. We plan to add additional metrics
290 in the future. In addition to evaluating the actual forecasts, we also use hindcasting
291 (forecasting on already collected data; Jolliffe & Stephenson, 2003) to gain additional
292 insight into the methods that work best for forecasting this system. For example, a
293 model is fit using rodent observations up to June 2005, then used to make a forecast 12
294 months out to May 2006. The observations of that 12-month period can immediately be
295 used to evaluate the model. Since hindcasting is conducted using data that has already
296 been collected, it allows model comparisons to be conducted on large numbers of
297 hindcasts and provides insight into which models make the best forecasts without
298 needing to wait for new data to be collected (Harris, Taylor, & White, 2018). It can also
299 be used to quickly evaluate new models instead of waiting for an adequate amount of
300 data to accumulate.

301 Archiving

302 Publicly archiving forecasts before new data is collected allows the field to assess,
303 compare, and build on forecasts made by different groups (McGill, 2012; Dietze et al.,
304 2016; Tredennick et al., 2016; Harris et al., 2018) (Figure 1). Archiving serves as a
305 form of pre-registration for model predictions because the forecasts cannot be modified
306 once the data to assess them has been collected. This helps facilitate an unbiased
307 interpretation of model performance. To serve this role, archives should be publicly
308 accessible and be a permanent record that cannot be changed or deleted. This second
309 criterion means that GitHub is not sufficient for archival purposes because repositories
310 can be changed or deleted (Bergman, 2012; White, 2015). We explored three major
311 repositories for archiving forecasts: FigShare (<https://figshare.com/>), Zenodo
312 (<https://zenodo.org/>), and Open Science Framework (<https://osf.io/>). While all three
313 repositories allowed for easy manual submissions (i.e., a human uploading files after
314 each forecast), automating this process was substantially more difficult. Various
315 combinations of repositories, APIs (i.e., interfaces for automatically interacting with the
316 archiving websites), and associated R packages had issues with: 1) integrating
317 authorization with continuous integration; 2) automatically making archived files public;
318 3) adding new files to an existing location; or 4) automatically permanently archiving
319 the files. Our eventual solution was to leverage the GitHub-Zenodo integration
320 (<https://guides.github.com/activities/citable-code/>) and automatically push forecasts to a
321 GitHub repository from the CI server and release them via the GitHub API. The
322 GitHub-Zenodo integration is designed to automatically create versioned archives of
323 GitHub repositories. We created a repository for storing forecasts
324 (<https://github.com/weecology/forecasts>) and linked this repository with Zenodo (a
325 one-time manual process). Each time a new forecast is created, our pipeline adds the
326 new forecasts to the GitHub repository and uses the GitHub API to create a new
327 “release” for that repository. This triggers the GitHub-Zenodo integration, which

328 automatically archives the resulting forecasts under a top-level DOI that refers to all
329 archived forecasts (<https://doi.org/10.5281/zenodo.839580>). Through this process, we
330 automatically archive every forecast made with a documented time-stamp. In addition,
331 we also archive the full state of the modeling and forecasting repository
332 (<https://doi.org/10.5281/zenodo.833438>). This ensures that every forecast is fully
333 reproducible since the exact code used to generate every forecast is preserved. Early
334 forecasts from this system are archived in the modeling and forecasting code archive,
335 not in the newer repository ‘forecasts’.

336 **Presentation**

337 Each month, we present our forecasts on a website that displays monthly rodent
338 forecasts, model evaluation metrics, monthly reports, and information about the study
339 site (Figure 3; <http://portal.naturecast.org>). The website includes a graphical
340 presentation of the most recent month’s forecasts (including uncertainty) and compares
341 the latest data to the previous forecasts. Information on the species and the field site are
342 also included. The site is built using Rmarkdown (Allaire et al., 2017), which naturally
343 integrates into the pipeline and is automatically updated after each forecast. The `knitr`
344 R package (Xie, 2015) compiles the code into HTML, which is then published using
345 Github Pages (<https://pages.github.com/>). The files for the website are stored in a
346 subdirectory of the forecasting repository. As a result, the website is also archived
347 automatically as part of archiving the forecast results.

348 **Discussion**

349 Following the recommendations of Dietze et al (2016), we developed an automated
350 iterative forecasting system (Figure 1) to support repeated forecasting of an ecological
351 system. Our forecasting system automatically acquires and processes the newest data,

Portal Forecast

Total Abundance Forecast

This is the forecast for next month's sampling of rodents at Portal.




Figure 3: Figure 3. Screen capture of the homepage of the Portal Forecasting website (<http://portal.naturecast.org>). This site contains information on the most current forecasts, evaluation of forecast performance, and general information about the species being forecast.

352 refits the models, makes new forecasts, publicly archives those forecasts, and presents
353 both the current forecast and information on how previous forecasts performed. Every
354 week, the forecasting system generates a new set of forecasts with no human
355 intervention, except for the entry of new field data. Our forecasting system ensures that
356 forecasts based on the most recent data are always available and is designed to allow
357 rapid assessment of the performance of multiple forecasting models for a number of
358 different states of the system, including the abundances of individual species and
359 community-level variables such as total abundance. To create this iterative near-term
360 forecasting system, we used R to process data and conduct analyses and leveraged
361 existing tools and services (i.e. GitHub, Travis, Docker) for more complicated
362 cyberinfrastructure tasks. Thus, our approach to developing iterative near-term
363 forecasting infrastructure provides an example for how short-term ecological
364 forecasting systems can be developed.

365 We designed this forecasting system with the goal of making it relatively easy to build,
366 maintain, and extend. We used existing technology for both running the pipeline and
367 building individual components, which allowed us to build the system relatively cheaply
368 in terms of both time and money. This included the use of tools like Docker for
369 reproducibility, Travis CI continuous integration for automatically running the pipeline,
370 Rmarkdown and `knitr` for generating the website, and the already existing integration
371 between Github and Zenodo to archive the forecasts. By using this “continuous analysis”
372 approach (Beaulieu-Jones & Greene, 2017), where analyses are automatically rerun
373 when changes are made to data, models, or associated code, we have reduced the time
374 required by scientists to run and maintain the forecasting pipeline. To make the system
375 extensible so that new models could be easily incorporated, we used a plugin-based
376 infrastructure so that adding a new model to the system is as easy as adding a single file
377 to the ‘models’ folder in our repository (Figure 2). This should substantially lower the
378 barriers to other scientists contributing models to this forecasting effort. We also

379 automatically archive the resulting forecasts publicly so that the performance of these
380 forecasts can be assessed by both us and other researchers as new data is collected. This
381 serves as a form of pre-registration by providing a quantitative record of the forecast
382 before the data being predicted were collected.

383 While building this system was facilitated by the use of existing technological solutions,
384 there were still a number of challenges in making existing tools work for automated
385 iterative forecasting. Continuous integration is designed primarily for running
386 automated tests on software, not for running a coordinated forecasting pipeline. As a
387 result, extra effort was sometimes necessary to figure out how to get these systems to
388 work properly in non-standard situations, like running code that was not part of a
389 software package. In addition, hosted continuous integration solutions, like Travis,
390 provide only limited computational resources. As the number and complexity of the
391 models we fit has grown, we have had to continually invest effort in reducing our total
392 compute time so we can stay within these limits. Finally, we found no satisfactory
393 existing solution for archiving our results. All approaches we tried had limitations when
394 it came to automatically generating publicly-versioned archives of forecasts on a
395 repeated basis, and our eventual solution was difficult to configure to such a degree that
396 it will remain an impediment for most researchers. Overall, we found existing
397 technology to be sufficient to the task of creating an iterative forecasting pipeline, but it
398 required greater expertise and a greater investment of time than is ideal. Additional tool
399 development to reduce the effort required for scientists to set up their own short-term
400 forecasting systems would clearly be useful. Our efforts, however, show that it is
401 possible to use existing tools to develop initial iterative systems as a method for both
402 advancing scientific understanding and developing proof of concept forecasting systems.

403 Because of the breadth of expertise needed to set up our forecasting pipeline, our effort
404 required a team with diverse skills and perspectives, ranging from software
405 development to field site expertise. It is rare to find such breadth within a single

406 research group, and our system was developed as a collaboration between the lab
407 collecting the data and a computational ecology lab. When teams have a breadth of
408 expertise, communication can be challenging (Winowiecki et al., 2011). We found a
409 shared base of knowledge related to both the field research and fundamental
410 computational skills was important for the success of the group. The two labs are part of
411 a joint interdisciplinary ecology group that has a mission of breaking down barriers
412 between field and computational/theoretical ecologists (<http://weecology.org>). Everyone
413 on the team had received training in fundamental data management and computing
414 skills through a combination of university courses, Software and Data Carpentry
415 workshops (Teal et al., 2015), and informal lab training efforts. In addition, everyone
416 was broadly familiar with the study site and methods of data collection, and most team
417 members had participated in field work at the site on multiple occasions. This provided
418 a shared set of knowledge and vocabulary that actively facilitated interdisciplinary
419 interactions. Given the current state of tools for forecasting, forecasting teams will need
420 people with significant experience in working with continuous integration and APIs.
421 This means interdisciplinary teams will generally be required for creating these
422 pipelines until tool development improves. To improve the success of these diverse
423 groups, we believe efforts at providing ‘team science’ training to scientists interested in
424 forecasting will be beneficial for the success of iterative forecasting attempts for the
425 foreseeable future (Read et al., 2016).

426 We developed infrastructure for automatically making iterative forecasts with the goals
427 of making accurate forecasts for this well-studied system, learning what methods work
428 well for ecological forecasting more generally, and improving our understanding of the
429 processes driving ecological dynamics. The most obvious application of automated
430 iterative ecological forecasting is for speeding up development of forecasting models by
431 using the most recent data available and by quickly iterating to improve the models used
432 for forecasting. By learning what works best for forecasting in this and other ecological

433 systems, we will better understand what the best approaches are for ecological
434 forecasting more generally. By designing the pipeline so that it can forecast many
435 different aspects of the ecological community, we also hope to learn about what aspects
436 of ecology are more forecastable. Finally, automated forecasting infrastructures like this
437 one also provide a core foundation for faster scientific inquiry because new models can
438 quickly be applied to data and compared to existing models. The forecasting
439 infrastructure does the time-consuming work of data processing, data integration, and
440 model assessment, allowing new research to focus on the models being developed and
441 the inferences about the system that can be drawn from them (Dietze et al., 2016). We
442 plan to use this pipeline to drive future research into understanding the processes that
443 govern the dynamics of individual populations and the community as a whole. By
444 regularly running different models for population and community dynamics, a near-term
445 iterative pipeline such as ours should also make it possible to rapidly detect changes in
446 how the system is operating, which should allow the rapid identification of ecological
447 transitions or even possibly allow them to be prevented (Pace et al., 2017). By building
448 an automated iterative near-term forecasting infrastructure, we can improve our ability
449 to forecast natural systems, understand the biology driving ecological dynamics, and
450 detect or even predict changes in system state that are important for conservation and
451 management.

452 **Acknowledgements**

453 We thank Henry Senyondo for help with continuous integration and Hao Ye and Juniper
454 Simonis for discussions and feedback on the manuscript. We thank all of the graduate
455 students, postdocs, and volunteers who have collected the Portal Project over the last 40
456 years and the developers of the software and tools that made this project possible. This
457 research was supported by the National Science Foundation through grant 1622425 to

458 S.K.M. Ernest and by the Gordon and Betty Moore Foundation's Data-Driven
459 Discovery Initiative through grant GBMF4563 to E.P. White.

460 **Data Accessibility**

461 The data used in this study is from the Portal Project and is openly available (CC0) on
462 GitHub (<https://github.com/weecology/PortalData>). Code for reproducing all analyses is
463 available on GitHub (<https://github.com/weecology/portalPredictions>) and archived on
464 Zenodo (White et al., 2018b). Forecasts made by this system are all archived to Zenodo
465 (White et al., 2018a).

466 **Box 1. Key practices for automated iterative near-term** 467 **ecological forecasting**

468 A list of some of the key practices developed by Dietze et al (2016) for facilitating
469 iterative near-term ecological forecasting and discussion of why these practices are
470 important.

471 **Data**

472 **1. Frequent data collection**

473 Frequent data collection allows models to be regularly updated and forecasts to be
474 frequently evaluated (Dietze et al., 2016). Depending on the system being studied, this
475 frequency could range from sub-daily to annual, but typically the more frequently the
476 data is collected the better.

477 **2. Rapid data release under open licenses**

478 Data should be released as quickly as possible (low latency) under open licenses so that
479 forecasts can be made frequently and data can be accessed by a community of
480 forecasters (Dietze et al., 2016; Vargas et al., 2017).

481 **3. Best practices in data structure**

482 To reduce the time and effort needed to incorporate data into models, best practices in
483 data structure should be employed for managing and storing collected data to ensure it
484 is easy to integrate into other systems (interoperability) (Borer et al., 2009; Strasser et
485 al., 2011; White et al., 2013).

486 **Models**

487 **4. Focus on uncertainty**

488 Understanding the uncertainty of forecasts is crucial to interpreting and understanding
489 their utility. Models used for forecasting should be probabilistic to properly quantify
490 uncertainty and to convey how this uncertainty increases through time. Evaluation of
491 forecast models should include assessment of how accurately they quantify uncertainty
492 as well as point estimates (Hooten & Hobbs, 2015).

493 **5. Compare forecasts to simple baselines**

494 Understanding how much information is present in a forecast requires comparing its
495 accuracy to simple baselines to see if the models yield improvements over the naive
496 expectation that the system is static (Harris et al., 2018).

497 **6. Compare and combine multiple modeling approaches**

498 To quickly learn about the best approaches to forecasting different aspects of ecology,
499 multiple modeling approaches should be compared (Harris et al., 2018). Different
500 modeling approaches should also be combined into ensemble models, which often
501 outperform single models for prediction (Weigel, Liniger, & Appenzeller, 2008).

502 **Cyberinfrastructure**

503 In addition to improvements in data and models, iterative near-term forecasting requires
504 improved infrastructure and approaches to support continuous model development and
505 iterative forecasting (Dietze et al., 2016).

506 **7. Best practices in software development**

507 Best practices should be followed in the development of scientific software and
508 modeling to make it easier to maintain, integrate into pipelines, and build on by other
509 researchers. Key best practices include open licenses, good documentation, version
510 control, and cross-platform support (Wilson et al., 2014; Hampton et al., 2015).

511 **8. Support easy inclusion of new models**

512 To facilitate the comparison and ensembling of different modeling approaches, code for
513 fitting models and making forecasts should be easily extensible, to allow models
514 developed by different groups to be integrated into a single framework (Dietze et al.,
515 2016).

516 **9. Automated end-to-end reproducibility**

517 Each forecast iteration involves acquiring new data, refitting the models, and making
518 new forecasts. This should be done automatically without requiring human intervention.
519 Therefore, the process of making forecasts should emphasize end-to-end reproducibility,
520 including data, models, and evaluation (Stodden & Miguez, 2014), to allow the
521 forecasts to be easily rerun as new data becomes available (Dietze et al., 2016).

522 **10. Publicly archive forecasts**

523 Forecasts should be openly archived to demonstrate that the forecasts were made
524 without knowledge of the outcomes and to allow the community to assess and compare
525 the performance of different forecasting approaches both now and in the future (McGill,
526 2012; Dietze et al., 2016; Tredennick et al., 2016; Harris et al., 2018). Ideally, the

527 forecasts and evaluation of their performance should be automatically posted publicly in
528 a manner that is understandable by both scientists and the broader stakeholder
529 community.

530 **References**

- 531 Allaire, J., Cheng, J., Xie, Y., McPherson, J., Chang, W., Allen, J., . . . Arslan, R. (2017).
532 *Rmarkdown: Dynamic documents for r*. Retrieved from
533 <https://CRAN.R-project.org/package=rmarkdown>
- 534 Araujo, M. B., & New, M. (2007). Ensemble forecasts of species distributions. *Trends*
535 *in Ecology and Evolution*, 22(1), 42–47.
- 536 Beaulieu-Jones, B. K., & Greene, C. S. (2017). Reproducibility of computational
537 workflows is automated using continuous analysis. *Nature Biotechnology*, 35(4),
538 342–346.
- 539 Bergman, C. (2012). On the preservation of published bioinformatics code on github.
540 Retrieved from [https://caseybergman.wordpress.com/2012/11/08/](https://caseybergman.wordpress.com/2012/11/08/on-the-preservation-of-published-bioinformatics-code-on-github/)
541 [on-the-preservation-of-published-bioinformatics-code-on-github/](https://caseybergman.wordpress.com/2012/11/08/on-the-preservation-of-published-bioinformatics-code-on-github/)
- 542 Boettiger, C., & Eddelbuettel, D. (2017). An introduction to rocker: Docker containers
543 for r. *ArXiv Preprint ArXiv:1710.03675*.
- 544 Borer, E. T., Seabloom, E. W., Jones, M. B., & Schildhauer, M. (2009). Some simple
545 guidelines for effective data management. *The Bulletin of the Ecological Society of*
546 *America*, 90(2), 205–214.
- 547 Brown, J. (1998). The desert granivory experiments at portal. *Experimental Ecology*.
548 *Oxford University Press, Oxford, UK*, 71–95.
- 549 Clark, J. S., Carpenter, S. R., Barber, M., Collins, S., Dobson, A., Foley, J. A., . . . others.

550 (2001). Ecological forecasts: An emerging imperative. *Science*, 293(5530), 657–660.

551 Dietze, M. C. (2017). *Ecological forecasting*. Princeton University Press.

552 Dietze, M. C., Fox, A., Betancourt, J. L., Hooten, M. B., Jarnevich, C. S., Keitt, T. H.,
553 ... others. (2016). Iterative ecological forecasting: Needs, opportunities, and
554 challenges. *Proceedings of the National Academy of Sciences*.

555 Diniz-Filho, J. A. F., Bini, L. M., Rangel, T. F., Loyola, R. D., Hof, C., Nogues-Bravo,
556 D., & Araujo, M. B. (2009). Partitioning and mapping uncertainties in ensembles of
557 forecasts of species turnover under climate change. *Ecography*, 32(6), 897–906.

558 Díaz, S., Demissew, S., Carabias, J., Joly, C., Lonsdale, M., Ash, N., ... others. (2015).
559 The ipbes conceptual framework—connecting nature and people. *Current Opinion in*
560 *Environmental Sustainability*, 14, 1–16.

561 Ernest, S. M., Brown, J. H., Thibault, K. M., White, E. P., & Goheen, J. R. (2008). Zero
562 sum, the niche, and metacommunities: Long-term dynamics of community assembly.
563 *The American Naturalist*, 172(6), E257–E269.

564 Ernest, S. M., Yenni, G. M., Allington, G., Christensen, E. M., Geluso, K., Goheen, J.
565 R., ... others. (2016). Long-term monitoring and experimental manipulation of a
566 chihuahuan desert ecosystem near portal, arizona (1977–2013). *Ecology*, 97(4),
567 1082–1082.

568 Ernest, S., Valone, T. J., & Brown, J. H. (2009). Long-term monitoring and
569 experimental manipulation of a chihuahuan desert ecosystem near portal, arizona, usa.
570 *Ecology*, 90(6), 1708–1708.

571 Hampton, S. E., Anderson, S. S., Bagby, S. C., Gries, C., Han, X., Hart, E. M., ...
572 others. (2015). The tao of open science for ecology. *Ecosphere*, 6(7), 1–13.

573 Harris, D. J., Taylor, S. D., & White, E. P. (2018). Forecasting biodiversity in breeding

574 birds using best practices. *PeerJ*.

575 Hooten, M. B., & Hobbs, N. (2015). A guide to bayesian model selection for ecologists.
576 *Ecological Monographs*, 85(1), 3–28.

577 Jolliffe, I. T., & Stephenson, D. B. (Eds.). (2003). *Forecast verification: a practitioner's*
578 *guide in atmospheric science*. John Wiley; Sons, Ltd. Retrieved from
579 <http://linkinghub.elsevier.com/retrieve/pii/S0169207005001214>

580 Liboschik, T., Fokianos, K., & Fried, R. (2015). *Tscount: An r package for analysis of*
581 *count time series following generalized linear models*. Universitätsbibliothek
582 Dortmund.

583 Luo, Y., Ogle, K., Tucker, C., Fei, S., Gao, C., LaDeau, S., . . . Schimel, D. S. (2011).
584 Ecological forecasting and data assimilation in a data-rich era. *Ecological Applications*,
585 21(5), 1429–1442.

586 McGill, B. J. (2012). Ecologists need to do a better job of prediction – part ii – partly
587 cloudy and a 20% chance of extinction (or the 6 p's of good prediction). Retrieved from
588 [https://dynamicecology.wordpress.com/2013/01/09/](https://dynamicecology.wordpress.com/2013/01/09/ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/)
589 [ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/](https://dynamicecology.wordpress.com/2013/01/09/ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/)

590 Merkel, D. (2014). Docker: Lightweight linux containers for consistent development
591 and deployment. *Linux J.*, 2014(239). Retrieved from
592 <http://dl.acm.org/citation.cfm?id=2600239.2600241>

593 Morris, B. D., & White, E. P. (2013). The ecodata retriever: Improving access to
594 existing ecological data. *PloS One*, 8(6), e65848.

595 Pace, M. L., Batt, R. D., Buelo, C. D., Carpenter, S. R., Cole, J. J., Kurtzweil, J. T., &
596 Wilkinson, G. M. (2017). Reversal of a cyanobacterial bloom in response to early
597 warnings. *Proceedings of the National Academy of Sciences*, 114(2), 352–357.

598 Petchey, O. L., Pontarp, M., Massie, T. M., Kéfi, S., Ozgul, A., Weilenmann, M., . . .

599 others. (2015). The ecological forecast horizon, and examples of its uses and
 600 determinants. *Ecology Letters*, 18(7), 597–611.

601 Read, E. K., O'Rourke, M., Hong, G., Hanson, P., Winslow, L. A., Crowley, S., ...
 602 Weathers, K. (2016). Building the team for team science. *Ecosphere*, 7(3).

603 Senyondo, H., Morris, B. D., Goel, A., Zhang, A., Narasimha, A., Negi, S., ... White,
 604 E. P. (2017). Retriever: Data retrieval tool. *The Journal of Open Source Software*, 2(19),
 605 451. doi:10.21105/joss.00451

606 Stodden, V., & Miguez, S. (2014). Best practices for computational science: Software
 607 infrastructure and environments for reproducible and extensible research. *Journal of*
 608 *Open Research Software*, 2(1).

609 Strasser, C., Cook, R., Michener, W., Budden, A., & Koskela, R. (2011). Promoting data
 610 stewardship through best practices. In *Proceedings of the environmental information*
 611 *management conference 2011 (eim 2011)*. Oak Ridge National Laboratory (ORNL).

612 Tallis, H. M., & Kareiva, P. (2006). Shaping global environmental decisions using
 613 socio-ecological models. *Trends in Ecology & Evolution*, 21(10), 562–568.

614 Teal, T. K., Cranston, K. A., Lapp, H., White, E., Wilson, G., Ram, K., & Pawlik, A.
 615 (2015). Data carpentry: Workshops to increase data literacy for researchers.
 616 *International Journal of Digital Curation*, 10(1), 135–143.

617 Tredennick, A. T., Hooten, M. B., Aldridge, C. L., Homer, C. G., Kleinhesselink, A. R.,
 618 & Adler, P. B. (2016). Forecasting climate change impacts on plant populations over
 619 large spatial extents. *Ecosphere*, 7(10).

620 Vargas, R., Alcaraz-Segura, D., Birdsey, R., Brunsell, N. A., Cruz-Gaistardo, C. O.,
 621 Jong, B. de, ... others. (2017). Enhancing interoperability to facilitate implementation
 622 of redd+: Case study of mexico. *Carbon Management*, 8(1), 57–65.

623 Weigel, A. P., Liniger, M., & Appenzeller, C. (2008). Can multi-model combination

624 really enhance the prediction skill of probabilistic ensemble forecasts? *Quarterly*
625 *Journal of the Royal Meteorological Society*, 134(630), 241–260.

626 White, E. P. (2015). Some thoughts on best publishing practices for scientific software.
627 *Ideas in Ecology and Evolution*, 8(1).

628 White, E. P., Baldridge, E., Brym, Z. T., Locey, K. J., McGlinn, D. J., & Supp, S. R.
629 (2013). Nine simple ways to make it easier to (re) use your data. *Ideas in Ecology and*
630 *Evolution*, 6(2).

631 White, E. P., Bledsoe, E. K., Christensen, E. M., Harris, D. J., Simonis, J. L., Taylor, S.
632 D., . . . Ernest, S. K. M. (2018a, February). Weecology/forecasts: 2018-02-17.
633 doi:10.5281/zenodo.839580

634 White, E. P., Yenni, G. M., Taylor, C., Shawn D and, Bledsoe, E. K., Simonis, J. L., &
635 Ernest, S. K. M. (2018b, February). Weecology/portalPredictions 2018-02-16.
636 doi:10.5281/zenodo.833438

637 Wickham, H. (2011). Testthat: Get started with testing. *The R Journal*, 3, 5–10.
638 Retrieved from
639 http://journal.r-project.org/archive/2011-1/RJournal_2011-1_Wickham.pdf

640 Wilson, G., Aruliah, D. A., Brown, C. T., Hong, N. P. C., Davis, M., Guy, R. T., . . .
641 others. (2014). Best practices for scientific computing. *PLoS Biology*, 12(1), e1001745.

642 Winowiecki, L., Smukler, S., Shirley, K., Remans, R., Peltier, G., Lothes, E., . . .
643 Alkema, L. (2011). Tools for enhancing interdisciplinary communication.
644 *Sustainability: Science, Practice and Policy*, 7(1), 74–80.

645 Xie, Y. (2015). *Dynamic documents with r and knitr* (Vol. 29). CRC Press.