


Automated iterative forecasting for the Portal Project

Introduction

Forecasting the future state of ecological systems is important for management, conservation, and evaluation of our fundamental understanding of ecology (Clark et al., 2001; Tallis & Kareiva, 2006; Díaz et al., 2015; Dietze, 2017). Since Clark et al. [Clark2001] called for a more central role of forecasting in ecology, an increasing number of ecological forecasts are being published. However, most of these forecasts are made once, published, and never assessed or updated. Without assessment, we have limited information on how much confidence to place in our predictions; without regular updates, forecasts lack the most up-to-date information as conditions change (Dietze et al., 2016). This lack of both regular assessment and active updating has limited the progress of ecological forecasting and hindered our ability to make useful and reliable predictions. For ecological forecasting to mature as a field, we need to change how we produce and interact with forecasts, creating a more dynamic interplay between model development, prediction generation, and incorporation of new data and information (Dietze et al., 2016).

With the goal of making ecological forecasting more dynamic and responsive, Dietze et al. [dietze2018] recently called for an increase in iterative near-term forecasting. Iterative near-term forecasting means making forecasts for the near future and making these forecasts repeatedly through a cycle of forecast evaluation, integration of updated data, and generation of new forecasts. This approach to forecasting has a number of advantages. Because forecasts are made ‘near-term’—daily to annual forecasting

24 instead of multi-decadal—predictions can be assessed more quickly and frequently,
25 leading to more rapid model improvements (Dietze et al., 2016; Tredennick et al., 2016).
26 Because the forecasts are made repeatedly through time, new data can be integrated with
27 each new forecast cycle. This iterative approach to forecasting allows any changes in
28 the state of the system that have occurred since the previous forecast to be incorporated
29 and accounted for (Dietze et al., 2016). Iterative near-term forecasting has the potential
30 to promote rapid improvement in the state of ecological forecasting by quickly
31 identifying how models are failing, facilitating rapid testing of improved models, and
32 incorporating updated data so models run with the most up-to-date information on the
33 system available. While use of iterative near-term forecasting is often contextualized as
34 a management tool s approach to model testing can also be used to improve our
35 basic understanding of ecological systems. For example, alternative mechanistic models
36 can be competed to see ~~which~~ model provides the best forecasts for near-term dynamics,
37 thus providing insights into the relative importance of different processes driving
38 dynamics of ecological systems (Dietze et al., 2016). Whether deployed for basic or
39 applied uses, iterative near-term forecasting incorporates a more dynamic interplay
40 between models, predictions, and data that is clearly needed to improve ecological
41 forecasting and our understanding of ecological systems more broadly.

42 Because iterative near-term forecasting requires a dynamic interplay of models,
43 predictions, and data, Dietze et al [-dietze2018] highlight approaches to data
44 management, model construction and evaluation, and cyberinfrastructure that are
45 necessary to effectively implement this type of forecasting (Box 1). Data to be used for
46 iterative near-term forecasting needs to be widely accessible, which requires data to be
47 released quickly under open licenses (Dietze et al., 2016; Vargas et al., 2017) and
48 structured so that it can be used easily by a variety of researchers and in multiple
49 modeling approaches (Borer, Seabloom, Jones, & Schildhauer, 2009; Strasser, Cook,
50 Michener, Budden, & Koskela, 2011). Models need to be able to deal with uncertainty,

51 ~~in both the predictors and the predictions~~, to properly convey uncertainty in the
52 resulting forecasts. Multiple models should be ~~compared~~ to assess which models are
53 performing best and to ~~allow for combining~~ models to form ensemble predictions.
54 Ensuring that data and models are regularly updated and new forecasts are made
55 requires cyberinfrastructure to automate data processing, model fitting, prediction,
56 model evaluation, forecast visualization, and archiving. In combination, these
57 approaches should allow forecasts to be easily rerun and evaluated as new data becomes
58 available (Box 1; Dietze et al., 2016).

59 While iterative near-term forecasting is an important next step in the evolution of
60 ecological forecasting, the requirements outlined by Dietze et al (Box 1) are not trivial
61 to implement; few of their recommendations are in widespread use in ecology today.
62 We examined what it would entail to operationalize Dietze et al's recommendations by
63 constructing our own iterative near-term forecasting pipeline for an on-going long-term
64 (~40 year) ecological study that collects high-frequency data on desert rodent
65 abundances. We constructed our forecasting pipeline with the goal of being able to
66 forecast rodent abundances and evaluate our predictions on a monthly basis. In this
67 paper, we discuss our approach for creating this iterative near-term forecasting pipeline,
68 the challenges we encountered, the tools we used, and the lessons we learned that may
69 help others to create their own iterative forecasting systems.

70 **System Background**

71 Iterative forecasting requires data that is collected repeatedly, and it benefits most from
72 data that is collected frequently, as this provides more opportunities for updating model
73 results and assessing (and potentially improving) model performance (Box 1; Dietze et
74 al., 2016). The Portal Project is a long-term ecological study situated in the Chihuahuan
75 Desert (2 km north and 6.5 km east of Portal, Arizona, US). Researchers have been

continuously collecting data at the site since 1977, including data on the abundance of rodent and plant species (monthly and twice yearly, respectively) and climactic factors such as air temperature and precipitation (daily). The site consists of 24 50m x 50m experimental plots. Each plot contains 49 permanently marked trapping stations laid out in a 7 x 7 grid, and all plots are trapped with Sherman live traps for one night each month. For all rodents caught during a trapping session, information on species identity, size, and reproductive condition is collected, and new individuals are given identification tags. We use the data from the control plots at ~~this site~~, where rodent populations are not experimentally manipulated. ~~This~~ data on rodent populations ~~is~~ high-frequency, ~~uses~~ consistent trapping methodology, and ~~has~~ an extended time-series (469 monthly samples and counting), making this study an ideal case for near-term iterative forecasting.

Implementing an automated iterative forecasting system

Implementation of iterative forecasting requires the regular rebuilding of models with new raw data as ~~it~~ becomes available and the presentation of those forecasts in usable forms; in our case, this occurs monthly. ~~Doing this~~ in an efficient and maintainable way relies on developing an automated pipeline to handle the six stages of converting raw data into new forecasts: data collection, data sharing, data manipulation, modeling and forecasting, archiving, and ~~presentation of~~ the forecasts (Figure 1). To implement the pipeline outlined in Figure 1, we used a “continuous analysis” framework (*sensu* Beaulieu-Jones & Greene, 2017) that automatically processes the most up-to-date data, refits the models, makes new forecasts, archives the forecasts, and updates a website with analysis of current and previous forecasts. In this section we describe our approach to streamlining and automating the multiple components of the forecasting pipeline and the tools and infrastructure we employed to execute each stage of the pipeline.

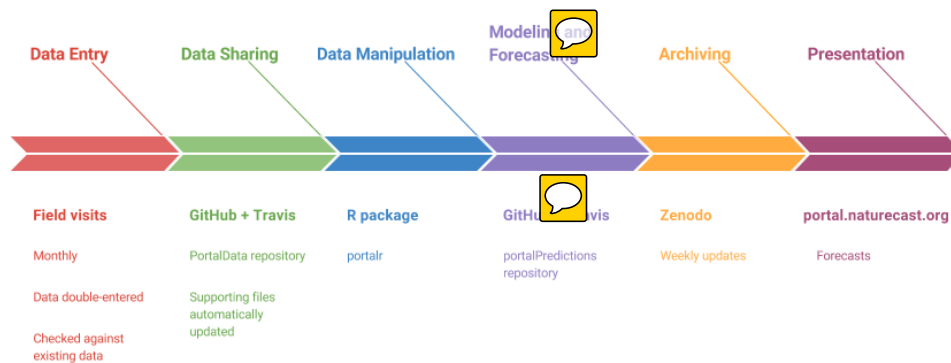


Figure 1: Figure 1. Stages of the forecasting pipeline. To go from raw data to forecast presentation involves a number of stages, each of which required unique tasks, tools and infrastructure. While each stage has unique tasks associated with it, the stages of our pipeline are also interdependent. Our pipeline is primarily a linear structure where outputs from one stage form the inputs for the subsequent stage

Continuous Analysis Framework

A core component of iterative near-term forecasting is the regular rerunning of the forecasting pipeline. This can be conducted manually—with a human making sure all code is run and all tables and files are updated—or automatically by having the computer conduct those tasks. We chose to have the computer run our pipeline and employed “continuous analysis” (*sensu* Beaulieu-Jones & Greene, 2017) to drive the automation of both the full pipeline and a number of its individual components. Continuous analysis uses a set of tools originally designed for software development called “continuous integration” (CI). CI combines computing environments for running code with monitoring systems to identify changes in data or code. Essentially, CI is a computer helper whose job is to watch the pipeline and, when it sees a change in the code or data, it runs all the tasks needed to ensure that the forecasting pipeline runs from beginning to end. This is useful for iterative near-term forecasting because it does

114 not rely on humans to remember to create forecasts when new models or data are added.
115 These tools are common in the area of software development where they are used to
116 automate software testing and integrate work by multiple software developers working
117 on the same software. However, these tools can be used for any computational task that
118 needs to be regularly repeated or run after changes to code or data (Beaulieu-Jones &
119 Greene, 2017). Because of the widespread use of CI in software development, several
120 CI services already exist. Our forecasting pipeline currently runs on a publicly available
121 continuous integration service (Travis CI; <https://travis-ci.org/>) that is free for open
122 source projects (up to a limited amount of computing time); alternative services that can
123 be used to run code on local or cloud-based computational infrastructure are available,
124 such as Drone (<http://try.drone.io/>) (Beaulieu-Jones & Greene, 2017). As detailed
125 below, we use CI to quality check data, test code using “unit tests” (Wilson et al., 2014),
126 build models, make forecasts, and publicly present and archive the results (Figure 2).

127 In addition to automatically running software pipelines, the other key component of
128 “continuous analysis” is making sure that the pipelines will continue to run even as
129 software dependencies change (Beaulieu-Jones & Greene, 2017). Many of us have
130 experienced the frustrations that can occur when software updates (e.g., changes in R
131 package versions) create errors in previously functional code. We experienced this issue
132 when a package one of our models relies on, `tscount` [liboschik2015], was
133 temporarily removed from CRAN (the R package repository) and, therefore, could not
134 be installed in the usual way. This broke our forecasting pipeline because we could no
135 longer run models that used that package. To minimize issues with changes in software
136 dependencies, we follow Beaulieu and Green’s (2017) recommendation to use software
137 containers. Software containers are standalone packages that contain copies of
138 everything you need to run some piece of software. Once created, a software container
139 is basically a time capsule, it contains all the software dependencies in the exact state
140 used to develop and run the software. If those dependencies change (or disappear) in the

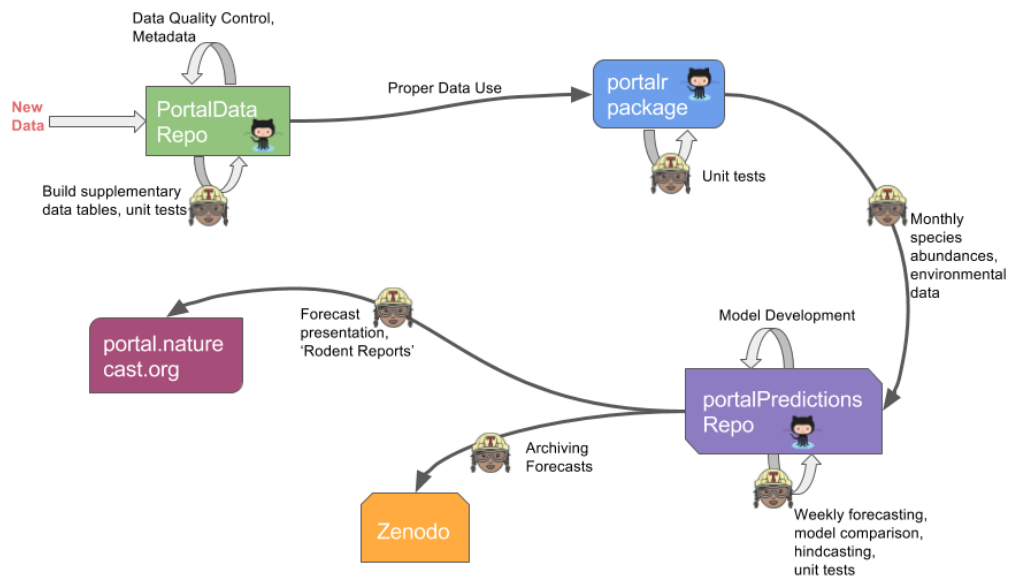



Figure 2: Figure 2. Continuous integration system. Each box denotes the core infrastructure used for each stage of the forecasting pipeline denoted in Figure 1. Continuous integration, denoted here with the Travis icon (a girl wearing safety glasses and hardhat), plugs into our pipeline at every stage. Travis triggers the code involved in major events that involve integration across stages of the pipeline, such as taking the output from the forecasting stage (purple box) to create an updated presentation (rose box). Travis also runs tasks within a stage, such as conducting tests to make sure code changes have not introduced errors (icons on arrows originating and ending on the same box)

141 wider world, they still exist, unchanged, in your container. We use an existing platform,
142 Docker, to store an exact image of the complete software environment for running the
143 forecasts. Docker also allows a specified set of packages to be used consistently across
144 different computer and server environments. Using containers allows us to update to
145 new package versions after testing and making any necessary changes to the data
146 processing and analysis code. We use a container created by the Rocker project which is
147 a Docker image with many important R packages (i.e. tidyverse) pre-installed (Boettiger
148 & Eddelbuettel, 2017). We use add our code and dependencies to this existing Rocker
149 image to create a software container for our forecasting pipeline. In combination, the
150 automated running of the pipeline (continuous integration) and the guarantee it will not
151 stop working unexpectedly due to software dependencies (via a software container)
152 allows continuous analysis to serve as the glue that connects all stages of the forecasting
153 pipeline.

154 **Data Collection, Entry, and Processing**

155 Iterative forecasting benefits from frequently updated ~~data on the state~~ of the system so
156 that state changes can be quickly incorporated into new forecasts (Dietze et al., 2016).
157 Frequent data collection and rapid ~~entry and~~ processing of data are both important for
158 providing ~~the most up-to-date data for forecasting~~. Since our data is collected monthly,
159 ensuring that the models have access to the newest data requires a data latency period of
160 less than 1 month from collection to availability for modeling. To accomplish this, we
161 automated components of the data processing and quality assurance/quality control
162 (QA/QC) process to reduce the time needed to add new data to the database (Figures 1
163 and 2).

164 New data is double-entered into Excel  using the “data validation” feature. The two
165 versions are then compared in an R script to control for errors in data entry. Quality
166 control (QC) checks written in R using the `testthat` R package (Wickham, 2011) are

167 run on the data to test for validity and consistency both within the new data (e.g. sexual
168 characteristics of an animal match M/F designation) and between the new and archived
169 data (e.g. species and sex are consistent for recaptures of the same animal based on tag
170 number). The local use of the QC scripts to flag problematic data greatly reduces the
171 time spent error-checking and ensures that the quality of data is consistent. The data is
172 then submitted to a GitHub-based data repository. Git is a software development
173 tool for managing computer code development, but we have also found it useful for data
174 management. Changes to data can be tracked through version control, and additions and
175 changes to the data can be monitored through pull requests (notifications that someone
176 has a modified version of the data to contribute). QA/QC checks are automatically rerun
177 on the submitted data using continuous integration to ensure that these checks have been
178 run and that no avoidable errors reach the official version of the dataset.

179 We also automated the updating of supplementary data tables, including information on
180 weather and trapping history, that were previously updated manually. As soon as new
181 field data is merged into the repository, continuous integration updates all
182 supplementary files. Weather data is automatically fetched from our cellular-connected
183 weather station, cleaned, and appended to the weather data table. Supplementary data
184 tables related to trapping history are updated based on the data added to the main data
185 tables. Using CI for this ensures that all supplementary data tables are always
186 up-to-date with the core data.

187 Data Sharing

188 The Portal Project has a long history of making its data publicly available, which means
189 that anyone can use it for forecasting or other projects. Historically the publication of
190 the data was conducted through data papers, the most common approach in ecology;
191 however, this approach caused years of data latency. Recently, the project has switched
192 to posting data directly to a public GitHub repository with a CC0 license (Figure 1).

193 This immediate posting reduces that data latency to less than one month and, therefore,
194 makes meaningful iterative near-term forecasting possible for not only our group but
195 other interested parties as well.

196 **Data Manipulation**

197 Once data is available, it needs to be processed into a form appropriate for modeling
198 (Figure 1). In many ecological datasets, this requires not only simple data manipulation
199 but also a good understanding of the dataset to allow data to be aggregated
200 appropriately. These data manipulation steps are often conducted using custom one-off
201 code to convert the data into the desired form (Morris & White, 2013), but this approach
202 has several limitations. First, each researcher must develop their own data manipulation
203 code, which is inefficient and can result in different decisions by different researchers
204 about the details of data cleaning and aggregation. Subtle differences in data processing
205 decisions have lead to confusion when reproducing results for the Portal data in the past.
206 Second, this kind of code is rarely robust to changes in data structure and location.
207 Based on our experience developing and maintaining the Data Retriever (Morris &
208 White, 2013; Senyondo et al., 2017), these kinds of changes are common. Finally, this
209 kind of code is generally poorly tested, which can lead to errors based on mistakes in
210 data manipulation. To avoid these issues for the Portal Project data, the Portal team has
211 been developing an R package (portalR; <http://github.com/weecology/portalr>) for
212 acquiring the data and handling common data cleaning and aggregation tasks. As a
213 result, our modeling and forecasting code only needs to install this package and run the
214 data manipulation and summary functions to get the appropriate data (Figure 2). The
215 package undergoes thorough automated unit tests to ensure that data manipulations are
216 achieving the desired results. Having data manipulation code maintained in a separate
217 package that focuses on consistently providing properly summarized forms of the most
218 recent data has made maintaining the forecasting code itself much more straightforward.

219 **Modeling and Forecasting**

220 Ideally, iterative near-term forecasting involves regularly refitting a variety of different
221 models (Figure 1). New models should be easy to incorporate to allow for iterative
222 improvements to the general modeling structure and approach. We use CI to refit the
223 models and make new forecasts each time the modeling code changes and when new
224 data becomes available (Figure 2). We use a plugin infrastructure to allow new models
225 to be easily added to the system. Each model is a script that takes in standardized inputs
226 and returns standardized outputs (Figure 3). All models are run by the main forecasting
227 code at each update, and the standardized outputs are combined to store the results of
228 the different models' forecasts. A weighted ensemble model is then added with weights
229 based on how well individual models fit the training data. This plugin infrastructure
230 allows flexibility in all aspects of the modeling process, making it easier to explore new
231 modeling approaches, and allows new models that fit the data well to immediately
232 improve the ensemble forecast.

233 To allow flexibility in what the models predict, we designed a forecast output format
234 that allows us to store relatively generic forecasts. Each forecast output file contains the
235 date being forecast, the collection date of the data used for fitting the models, the date
236 the forecast was made, the state variable being forecast (e.g., rodent biomass, the
237 abundance of a species), and the forecast value and associated uncertainty of that
238 forecast (Figure 3). This allows a variety of different forecasts to be stored in a common
239 format and may serve as a useful starting point for developing a standard for storing
240 ecological forecasts more generally.

241 Forecasts are currently evaluated using root mean square error (RMSE) to evaluate
242 point forecasts and coverage to evaluate uncertainty. We plan to add additional metrics
243 in the future. In addition to evaluating the actual forecasts, we also use hindcasting
244 (forecasting on already collected data) to gain additional insight into the methods that
245 work best for forecasting this system. For example, a model is fit using rodent

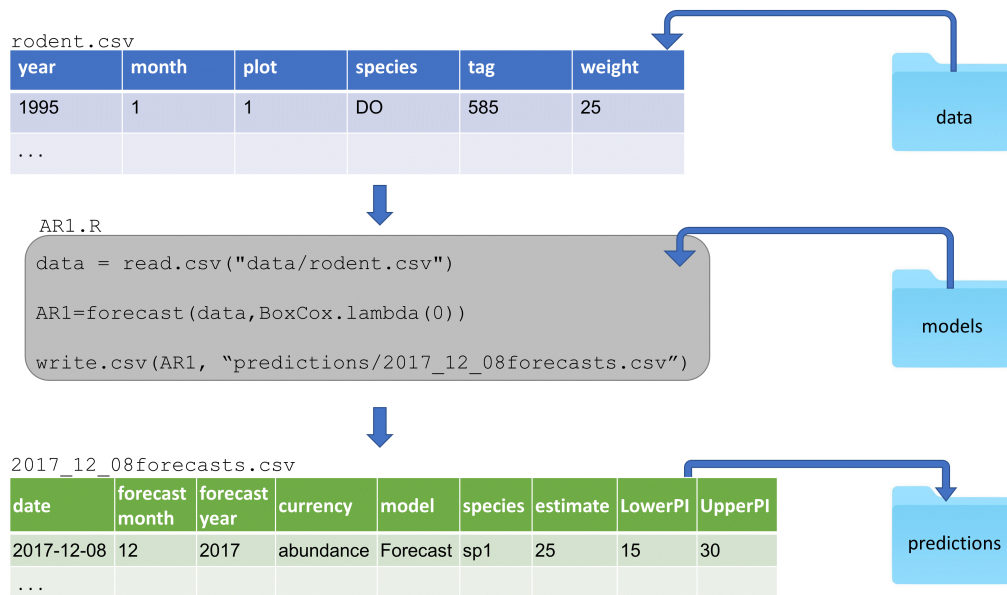


Figure 3: Figure 3. Demonstration of plugin infrastructure where each model script (represented here as AR1.R) uses data provided by the core forecasting code (represented here as rodent.csv) and returns its forecasts in a predefined structure that is consistent across models (represented here as 2017_12_08forecasts.csv) to allow cross-model evaluations.

246 observations up to June 2005, then used to make a forecast 12 months out to May 2006.
 247 The observations of that 12 month period can immediately be used to evaluate the
 248 model. Hindcasting can be conducted using all months from the beginning of the study,
 249 thus allowing model comparison of large numbers of hindcasts and giving insight into
 250 which models make the best forecasts. It can also be used to quickly evaluate new
 251 models instead of waiting for an adequate amount of data to accumulate.

252 Archiving

253 Publicly archiving forecasts before new data is collected allows the field to assess,
 254 compare, and build on forecasts made by different groups (McGill, 2012; Tredennick et
 255 al., 2016; Dietze, 2017; J. Harris David, Taylor, & White, 2018)(Figure 1). This
 256 archiving serves as a form of pre-registration for model predictions, helping facilitate
 257 unbiased interpretation of model performance. We explored three major repositories for

258 archiving our forecasts: FigShare, Zenodo, and Open Science Framework. While all
259 three repositories allowed for easy manual submissions (i.e., a human uploading files
260 after each forecast), automating this process was substantially more difficult. Various
261 combinations of repositories, APIs, and associated R packages had issues with: 1)
262 integrating authorization with continuous integration; 2) automatically making archived
263 files public; 3) adding new files to an existing location; or 4) automatically permanently
264 archiving the files. Our eventual solution was to leverage the GitHub-Zenodo
265 integration (<https://guides.github.com/activities/citable-code/>) and automatically push
266 forecasts to the GitHub repository from the CI server. The GitHub-Zenodo integration is
267 designed to automatically create versioned archives of GitHub repositories. There is an
268 existing one-time process for linking our forecasting repository on GitHub with Zenodo.
269 Once this link is created, each time a new forecast is created, our pipeline adds the new
270 forecasts to the GitHub repository and uses the GitHub API to create a new “release” of
271 our repository. This triggers the Zenodo-GitHub integration, which automatically
272 archives the resulting forecasts and the code that generated them under a top-level DOI
273 that refers to all archived forecasts (<https://doi.org/10.5281/zenodo.833438>). Through
274 this process, we automatically archive every forecast made with a documented history
275 of the archive. While this approach is functional because everything in the repository is
276 archived when a new forecast is made, these archives are complicated, making it more
277 complicated than necessary to find and access the forecasting results.

278 **Presentation**

279 In addition to archiving the results of each forecast, we present them on a website that
280 displays monthly rodent forecasts, model evaluation metrics, monthly reports, and
281 information about the study site (Figure 4; <http://portal.naturecast.org>). The website
282 includes a graphical presentation of the most recent month’s forecasts (including
283 uncertainty) and compares the latest data to the previous forecasts. Information on the

Portal Forecast

Total Abundance Forecast

This is the forecast for next month's sampling of rodents at Portal.





Figure 4: Figure 4. Screen capture of the homepage of the Portal Forecasting website; <http://portal.naturecast.org>. This site contains information on the most current forecasts, evaluation of forecast performance, and general information about the species being forecast.


284 the species and the field site targeted to a general audience are also included. The site is
 285 built using Rmarkdown (Allaire et al., 2017), which naturally integrates into the
 286 pipeline, and is automatically updated after each forecast. The `knitr` R package (Xie,
 287 2015) compiles the code into HTML, which is then published using Github Pages
 288 (<https://pages.github.com/>). The files for the website are stored in a subdirectory of the
 289 forecasting repository. As a result, the website is also archived automatically as part of
 290 the forecast archiving.



291 Discussion


292 Following the recommendations of Dietze et al [[dietze2018](#)], we developed an
293 automated iterative forecasting system (Figures 1 and 2) to support repeated forecasting
294 of an ecological system. Our forecasting system automatically acquires and processes
295 the newest data, refits the models, makes new forecasts, publicly archives those
296 forecasts, and presents both the current forecast and information on how previous
297 forecasts performed. Every week our forecasting system generates a new set of
298 forecasts with no human intervention, except for the [entry](#) of new field data. This
299 ensures that forecasts based on the most recent data are always available and allows us
300 to rapidly assess the performance of multiple forecasting models [for a number of](#) 
301 [different states of the system, including the abundances of individual species and](#)
302 [community-level variables such as total abundance](#). To create this iterative near-term
303 forecasting system, we used R to process data and conduct analyses, and we leveraged
304 already existing services (i.e. GitHub, Travis, Docker) for more complicated
305 cyberinfrastructure tasks. Thus, our approach to developing iterative near-term
306 forecasting infrastructure provides an example for how short-term ecological
307 forecasting systems can be initially developed.




308 We designed this forecasting system with the goal of making it relatively easy to build,
309 maintain, and extend. We used existing technology for both running the pipeline and
310 building individual components, which allowed [the system to be built](#) ~~to be built~~ relatively cheaply
311 in terms of both time and money. This included the use of tools like Docker for
312 reproducibility, the Travis CI continuous integration system for automatically running
313 the pipeline, Rmarkdown and `knitr` for generating the website, and the already
314 existing integration between Github and Zenodo to archive the forecasts. By using this
315 “continuous analysis” approach (Beaulieu-Jones & Greene, 2017), where analyses are
316 automatically rerun when changes are made to data, models, or associated code, we
317 have reduced the time required by scientists to run and maintain the forecasting pipeline.


318 To make the system extensible so that new models could be easily incorporated, we use
319 a plugin-based infrastructure so that adding a new model to the system is as easy as
320 adding a single file to the ‘models’ folder in our repository (Figure 3). This should
321 substantially lower the barriers to other scientists contributing models to this forecasting
322 effort. We also automatically archive the resulting forecasts publicly so that, as new data
323 is collected, the performance of these forecasts can be assessed by both us and other
324 researchers. This serves as a form of “egistration” by providing a quantitative
325 record of the forecast before the data being predicted were collected.

326 While building this system was facilitated by the use of existing technological solutions,
327 there were still a number of challenges in making existing tools work for automated
328 iterative forecasting. Continuous integration is designed primarily for running
329 automated tests on software, not for running a coordinated forecasting pipeline. As a
330 result, extra effort was sometimes necessary to figure out how to get these systems to
331 work properly in non-standard situations, like running code that was not part of a
332 software package. In addition, hosted continuous integration solutions, like Travis,
333 provide only limited computational resources. As the number and complexity of the
334 models we fit has grown, we have had to continually invest effort in reducing our total
335 compute time so we can stay within these limits. Finally, we found no satisfactory
336 existing solution for archiving our results. All approaches we tried had limitations when
337 it came to automatically generating publicly versioned archives of forecasts on a
338 repeated basis. Overall, we found existing technology to be sufficient to the task, but it
339 required greater expertise and a greater investment of time than is ideal. Tool
340 development to reduce the effort required for scientists to set up their own short-term
341 forecasting systems would clearly be useful, but our efforts show that it is still currently
342 possible for scientists using existing tools to develop initial iterative systems as a
343 method for both advancing scientific understanding and developing proof of concept
344 forecasting systems. By developing these systems so that they already produce

345 automated iterative forecasts, it will be easier to convert ~~these systems~~ into fully
346 operationalized forecasting systems that are relied on for decision making (Dietze et al.,
347 2016, other paper by operationalization co-authors).

348 Because of the breadth of expertise needed to set up our forecasting pipeline, our effort
349 required a team with diverse skills and perspectives, ranging from software
350 development to field site expertise. It is rare to find such breadth within a single
351 research group, and our system was developed as a collaboration between the lab
352 collecting the data and a computational ecology lab. When teams have a breadth of
353 expertise, communication can be challenging. We found a shared base of knowledge
354 related to both the field research and fundamental computational skills was important
355 for the success of the group.  Everyone on the team ~~had~~ received training in fundamental
356 data management and computing skills through a combination of university courses,
357 Software and Data Carpentry workshops (Teal et al., 2015), and informal lab training
358 efforts. In addition, everyone was broadly familiar with the study site and methods of
359 data collection, and most team members had participated in field work at the site on
360 multiple occasions. This provided a shared set of knowledge and vocabulary that
361 actively facilitated interdisciplinary interactions. Given the current state of ~~existing~~
362 tools for forecasting, forecasting teams will need people with significant experience in
363 working with continuous integration and APIs, which means interdisciplinary teams
364 will generally be required for creating these pipelines until tool development improves.

365 We developed this infrastructure for automatically making iterative forecasts with the
366 goals of making accurate forecasts for this well-studied system, learning what methods
367 work well for ecological forecasting more generally, and improving our understanding
368 of the processes driving ecological dynamics.  The most obvious applica of
369 automated iterative ecological forecasting is for speeding up development of forecasting
370 models by providing the most recent data available to models and by quickly iterating to
371 improve the models used for forecasting.  Learning what works best for forecasting in

372 this and other ecological systems, we will better understand what the best approaches
373 are for ecological forecasting more generally. By designing the pipeline so that it can
374 forecast many different aspects of the ecological community, we also hope to learn
375 about what aspects of ecology are more forecastable. lly, automated forecasting
376 infrastructures like this one also provide a core foundation for faster scientific inquiry
377 more broadly because new models can quickly be applied to data and compared to
378 existing models. The forecasting infrastructure does the time-consuming work of data
379 processing, data integration, and model assessment, allowing new research to focus on
380 the models being developed and the inferences about the system that can be drawn from
381 them (Dietze et al., 2016). We plan to use this pipeline to drive future research into
382 understanding the processes that govern the dynamics of individual populations and the
383 community as a whole. By regularly running different models for population and
384 community dynamics, a near-term iterative pipeline such as ours should also make it
385 possible to rapidly detect changes in how the system is operating, which should allow
386 the rapid identification of ecological transitions or even possibly allow them to be
387 prevented (??? example). By building an automated iterative near-term forecasting
388 infrastructure we can improve our ability to forecast natural systems, our understanding
389 of the biology driving ecological dynamics, and detect or even predict changes in
390 system state that are important for conservation and management.

391 Acknowledgements

392 This research was supported by the National Science Foundation through grant 1622425
393 to S.K.M. Ernest and by the Gordon and Betty Moore Foundation's Data-Driven
394 Discovery Initiative through grant GBMF4563 to E.P. White. We thank all of the
395 graduate students, postdocs, and volunteers who have collected the Portal Project over
396 the last 40 years and the developers of all of the software and tools that made this

397 project possible.

398 **Box 1. Key practices for automated iterative near-term** 399 **ecological forecasting**

400 A list of some of the key practices developed by Dietze et al [-dietze2018] for
401 facilitating iterative near-term ecological forecasting and discussion of why these
402 practices are important.

403 **Data**

404 **1. Frequent data collection**

405 Frequent data collection allows models to be regularly updated and forecasts to be
406 frequently evaluated (Dietze et al., 2016). Depending on the system being studied, this
407 frequency could range from sub-daily to annual, but typically the more frequently the
408 data is collected the better.

409 **2. Rapid data release under open licenses**

410 Data should be released as quickly as possible (low latency) under open licenses so that
411 forecasts can be made frequently and data can be accessed by a community of
412 forecasters (Dietze et al., 2016; Vargas et al., 2017).

413 **3. Best practices in data structure**


414 To reduce the time and effort needed to incorporate this data into models, best practices
415 in data structure need to be employed for managing and storing collected data to ensure
416 it is easy to integrate into other systems (interoperability) (Borer et al., 2009; Strasser et
417 al., 2011; White et al., 2013).

418 **Models**

419 **4. Focus on uncertainty**

420 Understanding the uncertainty of forecasts is crucial to interpreting the forecasts and
421 understanding their utility. Models used for forecasting should be probabilistic to
422 properly quantify uncertainty and to convey how this uncertainty increases through time.
423 Evaluation of forecast models should include assessment of how accurately they
424 quantify uncertainty as well as point estimates. This can be done using “proper and
425 local” scores (Hooten & Hobbs, 2015).

426 **5. Compare forecasts to simple baselines**

427 Understanding how much information is present in a forecast requires comparing its
428 accuracy to simple baselines to see if the models yield improvements over the naive
429 expectation that the system is  (J. Harris David et al., 2018).

430 **6. Compare and combine multiple modeling approaches**

431 To quickly learn about the best approaches to forecasting different aspects of ecology,
432 multiple modeling approaches should be compared for forecasting tasks (J. Harris
433 David et al., 2018). Different modeling approaches should also be combined into
434 ensemble models, which are known to outperform single models for many forecasting
435 and prediction tasks (Weigel, Liniger, & Appenzeller, 2008).

436 **Cyberinfrastructure**

437 In addition to improvements in data and models, iterative near-term forecasting requires
438 improved infrastructure and approaches to support continuous model development and
439 iterative forecasting (Dietze et al., 2016).

440 **7. Best practices in software development**

441 Best practices should be followed in the development of scientific software and
442 modeling to make it easier to maintain, integrate into pipelines, and build on by other
443 researchers. Key best practices include using open licenses, good documentation,
444 version control, and cross-platform support (Wilson et al., 2014; Hampton et al., 2015).

445 **8. Support easy inclusion of new models**

446 To facilitate the comparison and ensembling of different modeling approaches, code for
447 fitting models and making forecasts should be easily extensible, allowing models
448 developed by different groups to be easily integrated into a single framework (Dietze et
449 al., 2016).

450 **9. Automated end-to-end reproducibility**

451 Iteratively making forecasts requires that acquiring the newest data, refitting the models,
452 and making new forecasts is simple. Ideally, this should be done automatically without
453 requiring human intervention. Therefore, the process of making forecasts should
454 emphasize end-to-end reproducibility, including data, models, and evaluation (Stodden
455 & Miguez, 2014), to allow the forecasts to be easily rerun as new data becomes
456 available (Dietze et al., 2016). Ideally, the entire forecasting pipeline will be rerun
457 automatically as new data becomes available.

458 **10. Publicly archive forecasts**

459 Forecasts should be openly archived to demonstrate that the forecasts were made
460 without knowledge of the outcomes (i.e., as a form of pre-registration *sensu*) and to
461 allow the community to assess and compare the performance of different forecasting
462 approaches both now and in the future (McGill, 2012; Dietze et al., 2016; Tredennick et
463 al., 2016; J. Harris David et al., 2018). Ideally, the forecasts and evaluation of their
464 performance should be automatically posted publicly in a manner that is understandable
465 by both interested scientists and other stakeholders.

References

- Allaire, J., Cheng, J., Xie, Y., McPherson, J., Chang, W., Allen, J., . . . Arslan, R. (2017). *Rmarkdown: Dynamic documents for r*. Retrieved from <https://CRAN.R-project.org/package=rmarkdown>
- Beaulieu-Jones, B. K., & Greene, C. S. (2017). Reproducibility of computational workflows is automated using continuous analysis. *Nature Biotechnology*, 35(4), 342–346.
- Boettiger, C., & Eddelbuettel, D. (2017). An introduction to rocker: Docker containers for r. *arXiv Preprint arXiv:1710.03675*.
- Borer, E. T., Seabloom, E. W., Jones, M. B., & Schildhauer, M. (2009). Some simple guidelines for effective data management. *The Bulletin of the Ecological Society of America*, 90(2), 205–214.
- Clark, J. S., Carpenter, S. R., Barber, M., Collins, S., Dobson, A., Foley, J. A., . . . others. (2001). Ecological forecasts: An emerging imperative. *Science*, 293(5530), 657–660.
- Dietze, M. C. (2017). *Ecological forecasting*. Princeton University Press.
- Dietze, M. C., Fox, A., Betancourt, J. L., Hooten, M. B., Jarnevich, C. S., Keitt, T. H., . . . others. (2016). Iterative ecological forecasting: Needs, opportunities, and challenges. *Proceedings of the National Academy of Sciences*.
- Díaz, S., Demissew, S., Carabias, J., Joly, C., Lonsdale, M., Ash, N., . . . others. (2015). The ipbes conceptual framework—connecting nature and people. *Current Opinion in Environmental Sustainability*, 14, 1–16.
- Hampton, S. E., Anderson, S. S., Bagby, S. C., Gries, C., Han, X., Hart, E. M., . . . others. (2015). The tao of open science for ecology. *Ecosphere*, 6(7), 1–13.
- Harris, J., David, Taylor, S. D., & White, E. P. (2018). Forecasting biodiversity in

490 breeding birds using best practices. *PeerJ*.

491 Hooten, M. B., & Hobbs, N. (2015). A guide to bayesian model selection for ecologists.
 492 *Ecological Monographs*, 85(1), 3–28.

493 McGill, B. J. (2012). Ecologists need to do a better job of prediction – part ii – partly
 494 cloudy and a 20% chance of extinction (or the 6 p’s of good prediction). Retrieved from
 495 [https://dynamicecology.wordpress.com/2013/01/09/](https://dynamicecology.wordpress.com/2013/01/09/ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/)
 496 [ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/](https://dynamicecology.wordpress.com/2013/01/09/ecologists-need-to-do-a-better-job-of-prediction-part-ii-mechanism-vs-pattern/)

497 Morris, B. D., & White, E. P. (2013). The ecodata retriever: Improving access to
 498 existing ecological data. *PloS One*, 8(6), e65848.

499 Senyondo, H., Morris, B. D., Goel, A., Zhang, A., Narasimha, A., Negi, S., . . . White,
 500 E. P. (2017). Retriever: Data retrieval tool. *The Journal of Open Source Software*, 2(19),
 501 451. doi:10.21105/joss.00451

502 Stodden, V., & Miguez, S. (2014). Best practices for computational science: Software
 503 infrastructure and environments for reproducible and extensible research. *Journal of*
 504 *Open Research Software*, 2(1).

505 Strasser, C., Cook, R., Michener, W., Budden, A., & Koskela, R. (2011). Promoting data
 506 stewardship through best practices. In *Proceedings of the environmental information*
 507 *management conference 2011 (eim 2011)*. Oak Ridge National Laboratory (ORNL).

508 Tallis, H. M., & Kareiva, P. (2006). Shaping global environmental decisions using
 509 socio-ecological models. *Trends in Ecology & Evolution*, 21(10), 562–568.

510 Teal, T. K., Cranston, K. A., Lapp, H., White, E., Wilson, G., Ram, K., & Pawlik, A.
 511 (2015). Data carpentry: Workshops to increase data literacy for researchers.
 512 *International Journal of Digital Curation*, 10(1), 135–143.

513 Tredennick, A. T., Hooten, M. B., Aldridge, C. L., Homer, C. G., Kleinhesselink, A. R.,
 514 & Adler, P. B. (2016). Forecasting climate change impacts on plant populations over

515 large spatial extents. *Ecosphere*, 7(10).

516 Vargas, R., Alcaraz-Segura, D., Birdsey, R., Brunsell, N. A., Cruz-Gaistardo, C. O.,
 517 Jong, B. de, ... others. (2017). Enhancing interoperability to facilitate implementation
 518 of redd+: Case study of mexico. *Carbon Management*, 8(1), 57–65.

519 Weigel, A. P., Liniger, M., & Appenzeller, C. (2008). Can multi-model combination
 520 really enhance the prediction skill of probabilistic ensemble forecasts? *Quarterly*
 521 *Journal of the Royal Meteorological Society*, 134(630), 241–260.

522 White, E. P., Baldridge, E., Brym, Z. T., Locey, K. J., McGlinn, D. J., & Supp, S. R.
 523 (2013). Nine simple ways to make it easier to (re) use your data. *Ideas in Ecology and*
 524 *Evolution*, 6(2).

525 Wickham, H. (2011). Testthat: Get started with testing. *The R Journal*, 3, 5–10.

526 Retrieved from
 527 http://journal.r-project.org/archive/2011-1/RJournal_2011-1_Wickham.pdf

528 Wilson, G., Aruliah, D. A., Brown, C. T., Hong, N. P. C., Davis, M., Guy, R. T., ...
 529 others. (2014). Best practices for scientific computing. *PLoS Biology*, 12(1), e1001745.

530 Xie, Y. (2015). *Dynamic documents with r and knitr* (Vol. 29). CRC Press.